

Omnidirectional type inference for ML: principality any way

We propose a new concept of *omnidirectional* type inference: the ability to resolve ML-style typing constraints in disorder. In contrast, all known existing implementations which typically infer the types of let-bound expressions before typechecking their use sites. This relies on two technical devices: *suspended match constraints*, which suspend the resolution of some constraints until the context has more information about a type variable; and *partial type schemes*, which allow taking instances of a partially solved type scheme containing suspended constraints, with a mechanism to incrementally update instances as the scheme is refined.

The benefits of omnidirectional type inference are striking for several advanced ML extensions, typically those that rely on optional type annotations where principality is fragile. We illustrate them with OCaml's static overloading of record labels and datatype constructors, semi-explicit first-class polymorphism, and tuple projections *à la SML*.

1 Introduction

The Damas-Hindley-Milner (HM) [Damas and Milner 1982] type system has long occupied a sweet spot in the design space of strongly typed programming languages, as it enjoys the *principal types property*: every well-typed expression e has a most general type σ from which all other valid types for e are instances of σ . For example, the identity function $\lambda x. x$ has the principal type $\forall \alpha. \alpha \rightarrow \alpha$, generalizing types like $\text{int} \rightarrow \text{int}$ and $\text{bool} \rightarrow \text{bool}$.

This property ensures predictable and efficient inference. Local typing decisions are always optimal, yielding most general types without guessing or backtracking. As a result, inference of subexpressions can proceed in any order, and well-typedness is preserved under common program transformations such as let-contraction and argument reordering.

Over the years, many extensions of ML have been proposed. Some of them, such as extensible records with row-polymorphism, higher-kinded types, or dimensional types, fit perfectly into the ML framework. Others such as GADTs, higher-rank polymorphism, or static overloading, are *fragile*, as they sometimes require explicit type annotations. The return type of overloaded datatype constructors may be annotated; polymorphic expressions can be annotated with a type scheme; and for GADTs, the type of the match scrutinee and return type can be annotated to a rigid type which will be refined by type equalities in each branch. Those type annotations may sometimes—but not always—be omitted.

Consider impredicative higher-rank polymorphism for instance:

```
let self f = f f
```

With higher-rank types, one could guess the type of f to be either $\forall \alpha. \alpha \rightarrow \alpha$ or $\forall \alpha. \alpha \rightarrow \alpha \rightarrow \alpha$ in order to typecheck self —neither of which is more general than the other, violating principality.

To fix this, inference algorithms require a minimal amount of *known* type information to restore principality; in this example the binding of f should be annotated with a polymorphic type scheme. Yet specifying such requirements declaratively is difficult. As a result, the specifications are often twisted with some direct or indirect algorithmic flavor in order to preserve principality and completeness. Moreover, these (more or less) ad-hoc restrictions commonly reject examples whose type could easily be guessed. For instance, MLF [Le Botlan and Rémy 2009] accepts or rejects the following expression, depending on the position of the annotation (using a traffic-light scheme: green and red indicate typechecking success and failure; orange signals a warning):

```
let self' (f : ∀α. α → α) = if true then f f else f
let self' f = if true then f f else (f : ∀α. α → α)
```

MLF
MLF

50 1.1 Directional type inference

51 Each fragile construct admits a robust counterpart where the type annotation is mandatory. While
 52 robust constructs fit perfectly into the ML framework, they are significantly more cumbersome
 53 to use, as they always require explicit type annotations. Fragile constructs can be defined by
 54 elaboration into their robust counterpart. The difficulty lies in finding a specification that is
 55 sufficiently expressive, principled, intuitive for the user, and for which we have a complete and
 56 effective elaboration algorithm.

57 The solutions proposed so far all enforce some ordering in which type inference is performed,
 58 which can then be used to propagate both inferred types and user-provided type annotations as
 59 *known* types that can be used for disambiguation and enable the omission of some annotations.
 60

61 *π-directional type inference*. Most ML inference algorithms enforce a fixed order when typechecking
 62 let-bindings $\text{let } x = e_1 \text{ in } e_2$: first typecheck the definition e_1 , then the body e_2 . OCaml leverages
 63 this ordering to resolve overloaded or ambiguous constructs in a *principal* way: polymorphic types
 64 are treated as *known* and may guide disambiguation, whereas monomorphic types are considered
 65 not-yet-known and cannot be relied on for disambiguation.

66 We call this π -directional (read as “**pi**-directional”) type inference, to mean that polymorphic
 67 expressions must be typed before their instances. This strategy ensures principality for fragile
 68 constructs, but can lead to counter-intuitive behavior.

69 To illustrate the problem, consider the following two record types with overlapping field names:

```
70 type α one = {x : α; y : int}
 71 type two = {x : int; z : int}
```

72 In OCaml, both definitions are in scope, and the compiler must statically disambiguate field usage.
 73 The expression $\{x = 1; z = 1\}$ can only be two, and $r.y$ necessarily infers one for r . But field
 74 accesses such as $r.x$ are ambiguous unless the type of r is *known*. Consider:

```
76 let ex1 r = r.x
 77 let ex2 = let r = {x = 1; y = 1} in r.x
 78 let ex3 = (fun r → r.x) {x = 1; y = 1}
```

OCaml	OmniML
OCaml	OmniML
OCaml	OmniML

80 In ex_1 , the type of r is unconstrained, so disambiguation fails¹. At first glance, ex_2 and ex_3 appear
 81 equivalent: in both, the expression $r.x$ can only refer to the field from type int one . Yet OCaml
 82 accepts ex_2 and rejects ex_3 . This is because the **let**-binding in ex_2 allows r to be treated polymor-
 83 phically, and thus its type is considered *known*—enabling disambiguation. In ex_3 , by contrast, r is
 84 monomorphic at the point of projection, and disambiguation is therefore forbidden.

85 To emphasize that this behavior is specification-drive and not an artifact of OCaml’s inference
 86 algorithm, consider two equivalent versions of ex_3 , where @@ and $|>$ are the application and reverse
 87 application functions:

```
88 let ex32 = (fun r → r.x) @@ {x = 1; y = 1}
 89 let ex33 = {x = 1; y = 1} |> (fun r → r.x)
```

OCaml	OmniML
OCaml	OmniML

90 While these terms are semantically equivalent, they highlight a potential hazard: their typability
 91 would vary under a directionally biased inference algorithm, depending on whether the function
 92 or argument is typed first. To avoid such implementation-dependent behavior, OCaml chooses to
 93 infer all subexpressions *simultaneously*, until they are **let**-bound.
 94

95 ¹In fact, OCaml uses a default resolution strategy instead of failing when the type is ambiguous, which is to emit a warning
 96 and use the last definition in scope. To check these examples, you should use the options `-principal -w +41`, which
 97 enforce principality checks and enables the warning on default resolution.

Consequently, OCaml does not make any difference between ex_3 , ex_{32} , or ex_{33} ; in all cases, disambiguation is disallowed, and they are ill-typed. This criterion also warns on the following example, where r has a monomorphic type:

```
let ex4 p r = if p then r.x else (r : α one).x
```

OCaml OmniML

Warning here is preferable to silently accepting or rejecting the program based on the inference order between the `if` branches.

π -directional inference offers a way to specify and implement principal type inference for fragile features, aligning with the implicit inference order present in most ML-like typecheckers. This mechanism was originally proposed by [Garrigue and Rémy \[1999\]](#) for semi-explicit first-class polymorphism, and used in `MLF`. It has since been adopted in OCaml for features such as polymorphic object methods and the overloading of record fields and variant constructors. More generally, OCaml uses π -directionality whenever the typechecker disambiguates on type information.

Bidirectional type inference. Bidirectional type inference is a standard alternative to unification for propagating type information. It is typically formulated by splitting typing rules into two modes: *checking mode*, which typechecks a term e against a *known* type τ in a given context, and *inference mode* which infers e from the context alone.

For example, the type system designer can decide to typecheck function applications $e_1 e_2$ by first *inferring* that e_1 has some function type $\tau \rightarrow \tau'$, and then *checking* e_2 against τ . This is not the only possible choice: bidirectional type inference is a framework that must be instantiated by assigning modes—*checking* or *inference*—to each language construct. There is usually no optimal assignment of modes: for any choice of modes, some programs will typecheck successfully, while others will fail unnecessarily. Yet the typing rules must irrevocably commit to a fixed set of modes. After which, principal types often exist, but only with respect to a specification that made non-principal choices to begin with.

Bidirectional type inference has been largely used for languages with higher-rank polymorphism, dependent types, or subtyping. Still, both OCaml and Haskell only use a limited form of bidirectional type checking with an underlying first-order unification-based type inference engine, that limits the downsides of bidirectional type inference.

Limitations of directional type inference. Bidirectional type inference is lightweight, practical, and well-suited for complex language features. It supports the propagation of type information with minimal annotations. Its main downside lies in the need to fix an often arbitrary flow of type information—as in the case of function applications discussed above.

On the other hand, π -*directional* type inference appears better suited for ML, relying on polymorphism, the essence of ML. But it remains surprisingly weak in some cases: it does not even allow the propagation of user-provided type annotations from a function to its argument! For example, the following would be rejected as ambiguous with π -directional type inference alone:

```
let g (f : two → int) : int = f {x = 1; z = 1} in g (fun r → r.x)
```

OCaml OmniML

OCaml uses π -directional type inference as its primary mechanism, alongside a weak form of bidirectional propagation. In this example the type of the argument of g is known π -directionally, but OCaml then propagates this expected type within the function definition in bidirectional fashion, so that this example may be considered non-ambiguous.

Besides, the implementation of π -directional type inference has an algorithmic cost. For technical reasons, type annotations must unshare types (from acyclic graphs as naturally produced by unification to trees), which may increase the size of types and the cost of type inference. For that

148 reason, the implementation of OCaml cheats and is incomplete by default. The user must explicitly
 149 pass the `-principal` flag to require the more expensive computation when desired.

150
 151 **1.2 Omnidirectional type inference**

152 In absence of *implicit* polymorphism, type inference is solely based on unification constraints
 153 which can be solved in any order; omnidirectional inference is then natural and easy to implement.
 154 The difficulty originates from ML *implicit* let-polymorphism for which all known implementations
 155 follow the π -order: first typing the binding, generalizing it into a type scheme, and finally typing
 156 the body under the extended typing environment that binds the generalized scheme. The Hindley-
 157 Milner algorithm \mathcal{J} , one of its variant \mathcal{W} or \mathcal{M} [Lee and Yi 1998], or more flexible constraint-based
 158 type inference implementations [Odersky, Sulzmann and Wehr 1999; Pottier and Rémy 2005; Rémy
 159 1990, 1992] all follow this strategy, to the best of our knowledge. However, this state of affairs is
 160 not a necessity.

161 To efficiently achieve omnidirectional type inference for fragile ML extensions:

- 162 (1) We introduce *suspended match constraints* as a way to suspend ambiguity resolution until
 163 sufficient information has been found from context so that they can be discharged.
 164 (2) We work with *partial types schemes*, i.e., with the ability to instantiate type schemes that are
 165 not yet fully determined and consequently revisit their instances when they are being refined,
 166 incrementally. This allows inferring parts of a let-body to disambiguate its definition,
 167 without duplicating constraint-solving work.

168 These technical devices are introduced once and for all—in a general framework of constraint-
 169 based type inference. Each fragile ML construct can then be implemented by suspended constraints
 170 that expand to its robust counterpart once the annotation has been inferred. This generality comes
 171 at a cost, which is that everything is hard:

- 172 (1) Giving an adequate semantics for suspended constraints is hard, as we must capture declar-
 173 atively the intuition that some type information must be *known* rather than *guessed*.
 174 (2) Implementing partial types schemes is also hard.

175 In return, the techniques we developed for the semantics also help provide declarative typing rules
 176 for each fragile construct, for which the generated constraints are correct and complete.

177 *Illustrative examples.* Examples ex_2 to ex_4 are all typable with omnidirectional inference, as
 178 indicated by the green traffic light labeled OmniML—the calculus formalized in this paper.

179 In contrast, both bidirectional and π -directional inference rely on specifications that include
 180 choices that are subjective and somewhat arbitrary. As a result, they reject programs that have a
 181 unique well-typed solution. We now turn to further examples that illustrate such failures:

182
 183 let ex_6 $r = (\lambda x. r.x, r.y)$
 184 let ex_7 $r = \text{let } x = r.x \text{ in } x + r.y$
 185 let $\text{ex}_8 = \text{let } \text{getx } r = r.x \text{ in } \text{getx } \{x = 1; y = 1\}$
 186 let ex_9 $r = (\lambda x : \text{bool}.)$
 187 let ex_{10} $r = r.x.x$

OCaml	OmniML

188 All are arguably unambiguous; OCaml accepts none of them, OmniML accepts the first three.

189 In ex_6 , r can only be of type α one. Indeed, considering the second projection first, we should
 190 learn that r is of type α one and since it is λ -bound, this should then make the first projection
 191 unambiguous. Disambiguating this example is a matter of solving the typing constraints in the
 192 right order.

193 A similar failure occurs in ex_7 , where the type of the λ -bound variable r is initially ambiguous
 194 and unknown. It is only upon typing the projection $r.y$ that r is forced to have the type α one; this

197 requires inferring the let-body to disambiguate the let-definition. In ex_8 , disambiguation information
 198 flows from an instance to back the definition, opposite to the π -order; we call this *backpropagation*.

199 The example ex_9 can be disambiguated from the return type of the projection, rather than from
 200 the type of r . The typing rules for records that we present in this work restrict disambiguation to
 201 the record type alone, and thus rejects this example. Alternative typing rules using omnidirectional
 202 type inference could support this example as well.

203 Finally, ex_{10} is an example where none of the field projections has enough type information to be
 204 disambiguated on its own, but the constraints they impose can be combined to deduce that the type
 205 of r must be one, as the x field of two does not have a record type. This lies outside the framework
 206 of omnidirectional type inference, in which suspended constraints must be discharged one by one
 207 in some order, independently of other still-suspended constraints. We believe that this restriction is
 208 necessary for effective type inference, since the complexity of general overloading without this
 209 restriction is NP-hard, even in the absence of let-polymorphism, as shown by an encoding of 3-SAT
 210 problem by [Charguéraud, Bodin, Dunfield and Riboulet \[2025\]](#).

211 *Plan.* The paper is organized as follows: In §2, we give an overview of suspended constraints
 212 and their application to three extensions for ML of various kind. In §3, we describe suspended
 213 match constraints and their semantics. In §4, we define OmniML, an extension of ML featuring
 214 static overloading of record labels, overloaded tuple projections, and semi-explicit first-class poly-
 215 morphism. We sketch its typing rules and state the theorems of soundness and completeness for
 216 constraint generation, as well as principality. By lack of space, detailed typing rules and constraint
 217 generation are postponed to §A. In §5, we provide a formal definition of our constraint solver
 218 as a series of non-deterministic rewriting rules and state the main theorems for correctness. In
 219 §6, we describe an efficient implementation of suspended constraints and partial type schemes.
 220 In §7, we compare with related work, and in §8, we conclude with a discussion of future work,
 221 including prototyped extensions whose theory is less clear. Appendix §C contains a complete
 222 technical reference, collecting key definitions and figures for convenient lookup. All proofs are
 223 postponed to appendices.

224 *Our contributions.* Our contributions are: (1) A novel *omnidirectional* type inference framework for
 225 extensions of ML with advanced features, based on two new devices, suspended constraints and
 226 partial type schemes; (2) A declarative semantics of suspended constraints that captures the idea
 227 that they wait on information that must be propagated from the context, not *guessed*. This includes,
 228 in particular, a new declarative characterization of *known* type information. (3) A complete yet
 229 efficient constraint-solving type inference algorithm. (4) Three instantiation of our framework that
 230 give new declarative type systems and their implementation using suspended constraints for tuple
 231 projection in the style of SML, static overloading of record fields and datatype constructors, and
 232 for semi-explicit first-class polymorphism.

234 2 Suspended constraints: an overview

235 The syntax of types and constraints is given in Figure 1. Monotypes (or just types) include, as usual,
 236 type variables α , the unit type 1 , arrow types, but also² structural tuples $\Pi_{i=1}^n \tau_i$, nominal types³ $t \bar{\tau}$,
 237 and polytypes $[\sigma]$. Type schemes σ are of the form $\forall \bar{\alpha}. \tau$, they are equal up to the reordering of
 238 binders and removal of useless variables. We write \mathcal{V} the set of type variables.

239 Building atop the constraint-based type inference framework of [Pottier and Rémy \[2005\]](#), we adopt
 240 a constraint language that includes both term and type variables. The language (in Figure 1) contains

241 ²These are grayed, as they will be introduced in the following subsections.

242 ³Type constructors are prefixed, except in OCaml code, where they are postfixized.

246	$\alpha, \beta, \gamma \in \mathcal{V}$	Type variables
247	$\tau ::= \alpha \mid 1 \mid \tau_1 \rightarrow \tau_2 \mid \prod_{i=1}^n \tau_i \mid t \bar{\tau} \mid [\sigma]$	Types
248	$\sigma ::= \tau \mid \forall \alpha. \sigma$	Type schemes
249	$C ::= \text{true} \mid \text{false} \mid C_1 \wedge C_2 \mid \exists \alpha. C \mid \forall \alpha. C \mid \tau_1 = \tau_2$	Constraints
250	$\mid \text{let } x = \lambda \alpha. C_1 \text{ in } C_2 \mid x \tau$	
251	$\mid \text{match } \tau \text{ with } \bar{\chi}$	
252	$\chi ::= \rho \rightarrow C$	Branches
253	ρ	Patterns
254	$\mathcal{C} ::= \square \mid \mathcal{C} \wedge C \mid C \wedge \mathcal{C} \mid \exists \alpha. \mathcal{C} \mid \forall \alpha. \mathcal{C}$	Constraint contexts
255	$\mid \text{let } x = \lambda \alpha. \mathcal{C} \text{ in } C \mid \text{let } x = \lambda \alpha. C \text{ in } \mathcal{C}$	
256	ζ	Shapes
257	ς	Canonical principal shapes
258		
259		

Fig. 1. Syntax of types and constraints.

tautological (true) and unsatisfiable (false) constraints, conjunctions ($C_1 \wedge C_2$). The constraint form ($\exists \alpha. C$) binds an existentially quantified type variable α in C , while the constraint ($\forall \alpha. C$) binds α universally. The constraint form ($\tau_1 = \tau_2$) asserts that the types τ_1 and τ_2 are equal. When σ is a polymorphic type scheme $\forall \bar{\alpha}. \tau'$, we use the notation ($\sigma \leq \tau$) as syntactic sugar for the instantiation constraint $\exists \bar{\alpha}. \tau' = \tau$.

Two constructs deal with the introduction and elimination of constraint abstractions. A constraint abstraction $\lambda \alpha. C$ can simply be seen as a function which when applied to some type τ returns $C[\alpha := \tau]$. Constraint abstractions are introduced by a let construct (let $x = \lambda \alpha. C_1$ in C_2) which binds the constraint abstraction to the term variable x in C_2 —additionally ensuring the abstraction is satisfiable. They are eliminated using the application constraint ($x \tau$) which applies the type τ to the abstraction constraint bound to x .

Finally, we introduce *suspended match constraints* (match τ with $\bar{\chi}$). These constraints are suspended until the *shape* of τ , such as its top-level constructor, is known. Then they are *discharged*: a unique branch is selected and its associated constraint has to be solved. A match constraint that is never discharged is considered unsatisfiable.

More precisely:

- (1) The matchee τ is a type. The constraint remains suspended while τ is a type variable, that is, until the shape of τ is determined.
- (2) $\bar{\chi}$ is a list of branches of the form $\rho \rightarrow C$, where ρ is a shape pattern. For example, the pattern $\alpha \rightarrow \beta$ matches function types, binding its domain and codomain to α and β , respectively. The constraint C is then solved in the extended context. To ensure determinism, the set of patterns $\bar{\rho}$ must be *disjoint*—that is, no shape may be matched by more than one pattern in the list.

We keep the grammar of shapes and patterns abstract in this section, to explain the general framework of suspended constraints. For now, it suffices to think of shapes as top-level constructors like function arrows $\cdot \rightarrow \cdot$. This will be made more precise in §3.

Throughout this paper, we will find it convenient to work with *constraint contexts*. A constraint context is simply a constraint with a *hole*, analogous to evaluation contexts \mathcal{C} used extensively in operational semantics. We write $\mathcal{C}[C]$ to denote filling the hole of the context \mathcal{C} with the constraint C . Hole filling may capture variables, so we frequently require explicit side conditions

when variable capture must be avoided. We write $\text{bv}(\mathcal{C})$ for the set of variables bound at the hole in \mathcal{C} .

Suspended constraints in action. The remainder of this section illustrates the role of suspended constraints in supporting *fragile* language features as defined above. In particular:

- (§2.1) Semi-explicit first-class polymorphism;
- (§2.2) Constructor and record label overloading for nominal algebraic datatypes;
- (§2.3) Overloaded tuple projection in the style of SML.

We demonstrate how the typability of each of these features can be elaborated into constraints, formalized using a constraint generation function of the form $\llbracket e : \alpha \rrbracket$, which, given a term e and expected type α , produces a constraint C which is satisfiable if and only if e is well-typed.

As we will see, once we adopt the suspended constraint machinery developed in this paper, much of the complexity of these typing fragile constructs vanishes—suspended constraints do most of the heavy lifting.

2.1 Semi-explicit first-class polymorphism

Semi-explicit first-class polymorphism [Garrigue and Rémy 1999] brings some System F-like expressiveness to ML by allowing impredicative, first-class polymorphism while preserving principal type inference. It has since been adopted by OCaml, notably for polymorphic object methods.

The type constructor $[\sigma]^{\varepsilon}$ boxes a polymorphic type scheme σ , turning it into a *polytype* annotated with the *annotation variable* ε . Once boxed, the polytype $[\sigma]^{\varepsilon}$ is considered a monotype, thereby enabling impredicative polymorphism. Annotation variables record the origins of polytypes and may themselves be generalized, yielding type schemes such as $\forall \varepsilon. [\sigma]^{\varepsilon}$. When ε is generalized, the polytype is considered *known*, rather than still being inferred—this distinction is precisely the purpose of annotation variables, and it captures π -directionality explicitly.

The introduction form for polytypes is a boxing operator $\llbracket e : \exists \bar{\alpha}. \sigma \rrbracket$ with an explicit polytype annotation $\exists \bar{\alpha}. \sigma$ where the $\bar{\alpha}$ are all the type variables that are free in σ . The resulting expression has type $[\sigma[\bar{\alpha} := \bar{\tau}]]^{\varepsilon}$ where ε is an arbitrary (typically fresh) annotation variable and $\bar{\tau}$ are arbitrary types that replace the free variables $\bar{\alpha}$. The annotation variable ε can thus be generalized. That is $\llbracket e : \exists \bar{\alpha}. \sigma \rrbracket$ can also be assigned the type scheme $\forall \varepsilon. [\sigma[\bar{\alpha} := \bar{\tau}]]^{\varepsilon}$.

Conversely, to instantiate a polytype expression, one must use an explicit unboxing operator $\langle e \rangle$, which requires no accompanying type annotation. However, the operator requires e to have a polytype scheme of the form $\forall \varepsilon. [\sigma]^{\varepsilon}$ and then assigns $\langle e \rangle$ a type τ that is an instance of σ . If, by contrast, e has the type $[\sigma]^{\varepsilon}$ for some non-generalizable annotation variable ε , then e is considered of a not-yet-known polytype, and therefore $\langle e \rangle$ is ill-typed.

For example, the expression $\lambda x. \langle x \rangle$ is not typable. Indeed, the λ -bound variable x is assigned a monotype. The only admissible type for x is $x : [\sigma]^{\varepsilon}$ for some σ and ε . Since ε is bound in the surrounding context at the point of typing $\langle x \rangle$, it cannot be generalized prior to unboxing, rendering the term ill-typed.

However, type annotations can be used to freshen annotation variables. We usually omit annotation variables in annotations, since we can implicitly introduce fresh ones in their place. For example, $\lambda x : [\sigma]. \langle x \rangle$ —which is syntactic sugar for $\lambda x. \text{let } x = (x : [\sigma]) \text{ in } \langle x \rangle$ —is well-typed because the explicit annotation introduces a fresh variable annotation ε_1 , which can then be generalized, yielding $\forall \varepsilon_1. [\sigma]^{\varepsilon_1}$.

This behavior can be counter-intuitive: type information that has just been inferred must still be considered as yet-unknown until its generalization. It also makes the system sensitive to the placement of type annotations, an artifact of the fixed directionality of generalization in π -directional inference. For instance, the following two terms differ only in the position of the

344 annotation, yet only the one on the left-hand side is well-typed.

$$\lambda f. \langle(f : [\forall\alpha. \alpha \rightarrow \alpha])\rangle f \quad \lambda f. \langle f \rangle (f : [\forall\alpha. \alpha \rightarrow \alpha])$$

345
346
347 The difference lies in how generalization and annotation variables interact. In the first term, the
348 annotation occurs in an unboxing operator introducing fresh annotation variables and may therefore
349 be generalized to the type scheme $\forall\epsilon. [\forall\alpha. \alpha \rightarrow \alpha]^{\epsilon}$, enabling unboxing to proceed. Whereas the
350 second term applies the annotation to the argument f , which fixes f 's type to the monotype
351 $[\forall\alpha. \alpha \rightarrow \alpha]^{\epsilon_1}$ for some fresh annotation variable ϵ_1 . Because this type is assigned to f at its
352 binding site, ϵ_1 is bound in the context when typing $\langle f \rangle$ and cannot be generalized, so the second
353 term is ill-typed despite the annotation.

354 Suspended match constraints eliminate this sensitivity to directionality when typechecking $\langle e \rangle$.
355 If e is already known to have the type $[\sigma]$, then we can simply instantiate it. However, if the type
356 of e is not yet known—i.e., it is a (possibly constrained) type variable α —then we must defer until
357 more information is available. We capture this behavior with a suspended match constraint:

$$\llbracket \langle e \rangle : \alpha \rrbracket \triangleq \exists \beta. \llbracket e : \beta \rrbracket \wedge \text{match } \beta \text{ with } ([s] \rightarrow s \leq \alpha)$$

358 The match constraint is suspended until β is resolved to a polytype $[\sigma]$ matching the pattern $[s]$,
359 which binds the type scheme σ to the scheme variable s . The selected branch then performs the
360 instantiation $s \leq \alpha$, that is $\sigma \leq \alpha$. By waiting for the type of e to be *known*, we ensure principal
361 types without annotation variables.
362

363 2.2 Static overloading of constructors and record labels

364 *Static overloading* denotes a form of overloading in which resolution is performed entirely at
365 compile time, enabling the compiler to select a unique implementation without relying on runtime
366 information—in contrast to *dynamic overloading*, which defers resolution to runtime via mechanisms
367 such as dictionary-passing or dynamic dispatch.
368

369 Many languages offer statically resolved overloading to avoid the overhead of dynamic dispatch.
370 C++ and Java resolve overloaded functions through compile-time specialization based on argument
371 types. Conversely, languages like Rust and Haskell primarily employ dynamic overloading via
372 traits and type classes, respectively, which can incur runtime overhead unless optimized away by
373 monomorphization and aggressive inlining.
374

375 As noted in the introduction, OCaml supports a limited yet useful form of static overloading
376 for record labels and datatype constructors. When encountering overloaded labels or constructors,
377 OCaml resolves ambiguity using local type information, guided by π -directional inference. Nominal
378 types $t \bar{\tau}$ carry annotation variables ϵ , written $t^{\epsilon} \bar{\tau}$. As discussed in §2.1, this mechanism allows one
379 to deduce that types polymorphic over their annotation variable $\forall\epsilon. t^{\epsilon} \bar{\tau}$ are *known*.

380 Because static overloading involves more intricate flows of information than polytype inference,
381 OCaml supplements π -directionality with a limited, ad-hoc form of bidirectional type inference.
382 This mechanism is folklore; no formal account has been given.
383

384 Beyond propagation, OCaml also exploits *closed-world reasoning* to resolve ambiguities in record
385 types. For instance:

```
386 let ex11 = {x = 42; z = 1337}
```

OCaml OmniML

387 Here, x and y appear together only in the type two , allowing the type checker to unambiguously
388 infer the type of e_{11} as two . If local type information and closed-world reasoning are insufficient,
389 OCaml falls back to a syntactic default: it selects the most recently defined compatible type. For
390 example:
391

```
let getx r = r.x
```

OCaml OmniML

393 The expression is compatible with both one and two, since each defines a field x. But two is chosen
 394 simply because it appears later in the source. We do not treat this behavior as principal; accordingly,
 395 we provide no formalization of such “default” rules, though their implementation is discussed
 396 further in §8. This fallback mechanism highlights the directionality of OCaml inference. Once the
 397 compiler selects a type, it commits to it—even if that choice causes errors downstream. Consider
 398 ex₇ from §1:

399 `let ex7 r = let x (* infers [two] *) = r.x in x + r.y`

400 OCaml OmniML

401 Here, OCaml defaults to two for r when typing r.x, but then fails to type r.y, as this default choice
 402 is fixed—even though one would have satisfied both projections.

403 We assume a global typing environment Ω mapping labels to type schemes, written $\ell : \forall \bar{\alpha}. \tau \rightarrow$
 404 $t \in \Omega$. A given label ℓ may be defined several times in Ω , but at most once at a given record type t.
 405 We write $\Omega(\ell/t)$ for the type scheme of ℓ in t when it exists.

406 We propose an alternative account of static overloading using suspended match constraints. For
 407 example, in the case of an ambiguous record projection e.ℓ, we generate the typing constraint:

$$408 \quad \llbracket e.\ell : \alpha \rrbracket \triangleq \exists \beta. \llbracket e : \beta \rrbracket \wedge \text{match } \beta \text{ with } t _ \rightarrow (\Omega(\ell/t) \leq \alpha \rightarrow \beta)$$

409 This constraint suspends resolution of the return type α until the record type β of e is known. Its
 410 branch matches against the nominal type pattern $t _$, binding the type constructor name to t. Using
 411 this, the appropriate type scheme for ℓ is retrieved from $\Omega(\ell/t)$, instantiated, and the resulting
 412 constraints are imposed on the domain and codomain of the field-access type.

413 OCaml programs that do not use the default rule are accepted by this approach. Certain ex-
 414 pressions, such as e₁₂ are well-typed under our account but rejected by OCaml’s current type
 415 checker.

416 Our approach also applies to overloaded datatype constructors. Since the formal treatment is
 417 analogous to that of record fields, we focus only on fields in this work. However, our prototype
 418 implementation of OmniML supports both.

421 2.3 Tuple projections à la SML

422 SML supports positional projections from tuples using expressions of the form #_j e to extract the j-
 423 th component of the tuple e. Internally, tuples in SML are treated as structural records with numeric
 424 labels, so (#_j e) desugars into a structural record field access e.j: if e has the type {j = τ_j ; ϱ },
 425 where ϱ is a row describing the remaining tuple fields, then e.j has type τ_j .

426 SML enforces an additional restriction: the tail ϱ must be fully determined (*i.e.*, it cannot be a
 427 polymorphic row variable). This ensures that the arity of the tuple is *known* statically from the
 428 surrounding context, thereby avoiding the need for row polymorphism. However, this restriction
 429 is not expressed in the typing rules themselves, but is specified operationally as part of the type
 430 inference process.

431 From a typing perspective, tuple projection in SML behaves like a form of static overloading: the
 432 expression e.j is valid only when e is known to be an n-ary tuple for some fixed $n \geq j$.

433 We can capture the typing of tuple projections precisely using suspended constraints. For the
 434 projection e.j, we generate the following constraint:

$$435 \quad \llbracket e.j : \alpha \rrbracket \triangleq \exists \beta. \llbracket e : \beta \rrbracket \wedge \text{match } \beta \text{ with } \Pi \gamma_j \rightarrow \alpha = \gamma$$

436 The suspended constraint (match β with $\Pi \gamma_j \rightarrow \alpha = \gamma$) blocks until the shape of e (β) is known
 437 to be a tuple of sufficient arity. The pattern $\Pi \gamma_j$ matches only tuple types $\Pi_{i=1}^n \tau_i$, where $n \geq j$,
 438 binding the j-th component to γ , which is then unified with the expected result type α .

442 *Comparison to SML.* Our understanding is that SML typecheckers implement row-polymorphic
 443 records under the hood, but they never generalize row variables, rejecting any declaration that
 444 leaves a row variable undetermined. This is a neat approach, and we conjecture that it accepts
 445 the same programs as our implementation of overloaded tuples using suspended constraints. On
 446 the other hand, it only works for structural types. It cannot be applied to disambiguate between
 447 nominal records or variants that share field or constructor names, particularly when some field
 448 projections or constructors need non-uniform typing rules, such polymorphic fields or GADT
 449 constructors.

$$\phi ::= \emptyset \mid \phi[\alpha := g] \mid \phi[x := \mathfrak{G}]$$

Semantic environments

453 TRUE	CONJ	EXISTS	FORALL	UNIF
454	$\frac{\phi \vdash C_1 \quad \phi \vdash C_2}{\phi \vdash C_1 \wedge C_2}$	$\frac{\phi[\alpha := g] \vdash C}{\phi \vdash \exists \alpha. C}$	$\frac{\forall g, \phi[\alpha := g] \vdash C}{\phi \vdash \forall \alpha. C}$	$\frac{\phi(\tau_1) = \phi(\tau_2)}{\phi \vdash \tau_1 = \tau_2}$
455				
456 LET		APP		
457	$\frac{\phi \vdash \exists \alpha. C_1 \quad \phi[x := \phi(\lambda \alpha. C_1)] \vdash C_2}{\phi \vdash \text{let } x = \lambda \alpha. C_1 \text{ in } C_2}$	$\frac{\phi(\tau) \in \phi(x)}{\phi \vdash x : \tau}$	$\phi(\lambda \alpha. C) \triangleq \{g \in \mathcal{G} : \phi[\alpha := g] \vdash C\}$	$C_1 \models C_2 \triangleq \forall \phi, \phi \vdash C_1 \implies \phi \vdash C_2$
458				
459			$C_1 \equiv C_2 \triangleq (C_1 \models C_2) \wedge (C_2 \models C_1)$	
460				

Fig. 2. Semantics of constraints (without suspended constraints).

3 Semantics of constraints

466 To implement a typechecker using constraint-based type inference, it suffices to generate constraints
 467 from terms and to solve them. To study the meta-theory of this approach, we follow the standard
 468 approach of assigning a *semantics* for our constraints—as declaratively as possible. The existence
 469 of well-defined declarative semantics provides a foundation for reasoning about correctness and
 470 validates the design of the constraint language.

471 In our work on suspended constraints, defining a satisfying semantics was the most challenging
 472 aspect. The key difficulty lies in capturing what it means for type information to be *known*. Our
 473 semantics is declarative, but not syntax-directed unlike the standard constraint semantics of Pottier
 474 and Rémy [2005]. This lack of syntax-directness complicates reasoning and proofs. On the upside,
 475 the semantics directly suggest declarative typing rules for the surface language.

476 The semantics of constraints follows the standard form of a satisfiability judgment $\phi \vdash C$. The
 477 semantic environment ϕ contains a ground assignment for each free variable of C (type and term
 478 variable), and $\phi \vdash C$ states that these assignments indeed satisfy C . Let us write \mathcal{G} for the set of
 479 *ground* types, types without free variables⁴. ϕ maps each type variable α to a ground types $g \in \mathcal{G}$,
 480 and each term variable x to sets of ground types $\mathfrak{G} \subseteq \mathcal{G}$ (the set of ground instances of a type
 481 scheme for x). We write $\phi[\alpha := g]$ and $\phi[x := \mathfrak{G}]$ for the extension of ϕ with a new binding. For a
 482 type τ , we write $\phi(\tau)$ for the ground type obtained by substitution.

483 The judgment is defined in Figure 2 for all constraint formers except suspended constraints; its
 484 definition on this fragment is standard and somewhat tautological. The constraint true is satisfied
 485 by any environment, and false by none. An environment ϕ satisfies $C_1 \wedge C_2$ if it satisfies both C_1 and
 486 C_2 . Satisfying $\exists \alpha. C$ requires finding a witness g for α . The universal constraint $\forall \alpha. C$ is satisfiable

488 ⁴Ground types are thus finite trees, assuming the existence of some base types such as int. In §8, we discuss the alternative
 489 choice of regular trees for the set of ground types that models equirecursive types.

if C is satisfiable for any binding of α . The unification constraint $\tau_1 = \tau_2$ is satisfied when $\phi(\tau_1)$ and $\phi(\tau_2)$ are equal.

The rule for let $x = \lambda\alpha.C_1$ in C_2 states that C_1 must be satisfied under *some* instantiation of its bound variable, and that C_2 must be satisfiable when x is bound to $\lambda\alpha.C_1$, or rather to its semantic interpretation as a set of ground types.

An application constraint $x : \tau$ is interpreted by checking that τ belongs to the set of types mapped to x in ϕ , that is, $\phi(\tau) \in \phi(x)$. Note that when $\phi(x)$ is of the form $\phi'(\lambda\alpha.C)$, where ϕ' is the environment at the binding site of x , then $\phi(\tau) \in \phi(x)$ holds iff $\phi'[\alpha := \phi(\tau)] \vdash C$, which corresponds to the intuition that the application $(\lambda\alpha.C) \tau$ should be equivalent to $C[\alpha := \tau]$.

Closed constraints are either satisfiable in any semantic environment (i.e., they are tautologies) or unsatisfiable. For example, the satisfiability of the constraint $\exists\alpha. \alpha = \text{int}$ is established by the derivation on the right-hand side.

$$\frac{\text{int} = \text{int}}{\phi[\alpha := \text{int}] \vdash \alpha = \text{int}} \text{UNIF}$$

$$\frac{}{\phi \vdash \exists\alpha. \alpha = \text{int}} \text{EXISTS}$$

We write $C_1 \models C_2$ to express that C_1 *entails* C_2 , meaning every solution ϕ to C_1 is also a solution to C_2 . We write $C_1 \equiv C_2$ to indicate that C_1 and C_2 are equivalent, that is, they have exactly the same set of solutions.

3.1 Shapes

We introduce *shapes* as a generalization of type constructors for suspended match constraints. They provide a uniform treatment of both constructors and polytypes, and are useful in defining polytype unification (§6).

A shape ζ is a type with holes, written $v\bar{y}. \tau$, where \bar{y} denotes the set of type variables representing the holes. By construction, we require \bar{y} to be *exactly* the free variables of τ . Hence, shapes are closed and do not contain useless binders. We consider shapes up to α -conversion. When τ is a ground type, we omit the binder and write simply τ . We write \perp for the shape $v\gamma. \gamma$, which we call the *trivial* shape. We write \mathcal{S} the set of non-trivial shapes.

Shapes are equipped with the standard instantiation ordering, defined by **INST-SHAPE**. When writing $\zeta \leq \zeta'$, we say that ζ is more general than ζ' . When ζ and ζ' are more general than one another, they are actually equal. The trivial shape \perp is the most general shape. If ζ is $v\bar{y}. \tau$, the shape application $\zeta \bar{t}$ is defined as $\tau[\bar{y} := \bar{t}]$. We say that ζ is a shape of τ when there exists \bar{t} such that $\tau = \zeta \bar{t}$; in this case we write that the pair (ζ, \bar{t}) is a decomposition of τ .

Definition 3.1. A non-trivial shape $\zeta \in \mathcal{S}$ is the principal shape of the type τ iff:

- (1) $\exists \bar{t}', \tau = \zeta \bar{t}'$
- (2) $\forall \zeta' \in \mathcal{S}, \forall \bar{t}', \tau = \zeta' \bar{t}' \implies \zeta \leq \zeta'$

THEOREM 3.2 (PRINCIPAL SHAPES). Any non-variable type τ has a non-trivial principal shape ζ .

There is an equivalent direct description of principal shapes ζ . They are precisely the shapes $v\bar{y}. \tau$ satisfying two conditions: (1) \bar{y} must be linear in τ i.e., each variable y in \bar{y} occurs exactly once in τ . (2) The type τ must be shallow, meaning that its structure is limited in the following way. When τ is not a polytype, all of its subterms must be variables. Shapes of this form are 1 , $\gamma_1 \rightarrow \gamma_2$, and $\Pi_{i=1}^n \gamma_i$, or $t \bar{y}$. When τ is a polytype $[\forall \bar{x}. \sigma']$, the only subterms of σ' that do not contain one of the polymorphic variables \bar{x} must be variables in \bar{y} .

A principal shape $v\bar{y}. \tau$ is *canonical* if its free variables appear in the sequence \bar{y} in the order in which they occur in τ . We write ζ for canonical principal shapes. Each non-variable type τ has a unique canonical principal shape, which we write $\text{shape}(\tau)$. For example, $\text{shape}(t \bar{t})$ is $(v\bar{y}. t \bar{y})$.

Polytypes are particularly interesting in this setting because they can be decomposed into shapes and treated analogously to type constructors. For instance, the polytype $[\forall\alpha. ([\forall\beta. (\beta \rightarrow \text{int list}) * \beta]) \rightarrow \alpha \rightarrow \alpha]$ has the principal shape $\varsigma := v\gamma. [\forall\alpha. ([\forall\beta. (\beta \rightarrow \gamma) * \beta]) \rightarrow \alpha \rightarrow \alpha]$. The original polytype can thus be represented as the shape application ς (`int list`).

3.2 Suspended constraints

We have left the syntax of shape patterns deliberately abstract. We also assume a matching relation:

$$\rho \text{ matches } \varsigma \bar{\gamma} = \theta$$

This partial function matches a pattern ρ against a principal shape ς opened with shape names $\bar{\gamma}$ (which must have the same arity as ς), yielding a substitution θ . The substitution binds the pattern variables to shape components, that may contain occurrences of the shape variables $\bar{\gamma}$. For our examples we define the trivial pattern `_` which matches any shape and binds nothing:

$$_ \text{ matches } \varsigma \bar{\gamma} \triangleq \emptyset$$

Definition 3.3 (Discharged match constraint). Given a suspended constraint (match τ with $\bar{\chi}$) and a canonical shape ς , we introduce the syntactic sugar (match $\tau := \varsigma$ with $\bar{\chi}$) for the *discharged match constraint* that selects the branch in $\bar{\chi}$ that matches ς :

$$\text{match } \tau := \varsigma \text{ with } \bar{\rho} \rightarrow \bar{C} \triangleq \begin{cases} \exists \bar{\alpha}. \tau = \varsigma \bar{\alpha} \wedge \theta(C_i) & \text{if } \rho_i \text{ matches } \varsigma \bar{\alpha} = \theta \\ \text{false} & \text{otherwise} \end{cases}$$

The first conjunct ($\tau = \varsigma \bar{\alpha}$) ensures that ς is indeed the canonical shape of τ , and the second conjunct is the selected branch constraint C_i under the appropriate substitution. Since the syntax of suspended match constraints requires that branch patterns are non-overlapping, the matching branch $\rho_i \rightarrow C_i$ is uniquely determined; but it may not exist as branches need not be exhaustive, in which case the discharged constraint is false.

A natural attempt. To provide semantics for our suspended constraints, a first idea is to propose the following rule—henceforth referred to as the *natural semantics* of suspended constraints.

$$\frac{\text{SUSP-NAT}}{\varsigma = \text{shape } (\phi(\tau)) \quad \phi \vdash \text{match } \tau := \varsigma \text{ with } \bar{\chi}} \phi \vdash \text{match } \tau \text{ with } \bar{\chi}$$

This rule states that a suspended constraint is satisfied by ϕ whenever the corresponding discharged constraint holds for the canonical shape ς of τ in the semantic environment ϕ . If ς matches no branch in $\bar{\chi}$, then the discharged constraint is not defined, so this rule cannot be applied, and the suspended constraint is unsatisfiable.

This semantics rule is nicely declarative, but unfortunately accepts too many constraints. For example, $\exists\alpha. \text{match } \alpha \text{ with } _ \rightarrow \alpha = \text{int}$ is satisfiable under this natural semantics:

$$\frac{_ \text{ matches int } \emptyset = \emptyset \quad \frac{\text{int} = \text{int}}{\phi[\alpha := \text{int}] \vdash \alpha = \text{int}} \text{ UNIF}}{\phi[\alpha := \text{int}] \vdash \text{match } \alpha \text{ with } _ \rightarrow \alpha = \text{int}} \text{ SUSP-NAT}$$

$$\frac{\phi[\alpha := \text{int}] \vdash \text{match } \alpha \text{ with } _ \rightarrow \alpha = \text{int}}{\phi \vdash \exists\alpha. \text{match } \alpha \text{ with } _ \rightarrow \alpha = \text{int}} \text{ EXISTS}$$

The semantics can *guess* the type of α and use it to unlock the match constraint, rather than requiring it to be *known* from the surrounding context. One could call the guess of $\alpha = \text{int}$ an “out of thin air” behavior. This does not match the intended meaning of suspended match constraints, and

589 raises several problems: (1) a reasonable solver—one that avoids guessing or backtracking—cannot
 590 be complete with respect to this semantics; (2) this breaks the existence of principal solutions.
 591 Consider the function $\lambda x. (x.2)$, which projects the second component of a tuple. The natural
 592 semantics lets us guess for x any tuple type of arity at least 2; so there is no principal type for x .
 593

594 *Contextual semantics.* To rule out guessing, we instead adopt a *contextual* semantics: a match
 595 constraint is satisfiable only if the shape of the type is determined by the surrounding context. The
 596 corresponding rule for suspended constraints, *SUSP-CTX* in Figure 3, is the only non-syntax-directed
 597 rule in our semantics. In this rule, the shape ς is not guessed from ϕ , but it must be recovered from
 598 the constraint context \mathcal{C} . The *uniqueness* condition $\mathcal{C}[\tau ! \varsigma]$ (defined below) ensures that ς is uniquely
 599 determined by \mathcal{C} .

600 *Definition 3.4 (Erasure).* The erasure $[\mathcal{C}]$ of a constraint C is defined as the constraint obtained
 601 by replacing suspended match constraints in C with true.

602 *Definition 3.5 (Simple constraints).* We say that C is *simple* if it contains no suspended match
 603 constraints. We write $\phi \vdash_{\text{simple}} C$ for a derivation of $\phi \vdash C$ that only uses the rules listed in Figure 2,
 604 without using *SUSP-CTX*. This judgment coincides with $\phi \vdash C$ on simple constraints.

605 *Definition 3.6 (Uniqueness).* We define the uniqueness condition $\mathcal{C}[\tau ! \varsigma]$, which states that τ has a unique
 606 canonical shape ς within the context \mathcal{C} as: $\forall \phi, g. \phi \vdash_{\text{simple}} [\mathcal{C}[\tau = g]] \implies \text{shape}(g) = \varsigma$.

$$\begin{array}{c} \text{SUSP-CTX} \\ \mathcal{C}[\tau ! \varsigma] \quad \phi \vdash \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{x}] \\ \hline \phi \vdash \mathcal{C}[\text{match } \tau \text{ with } \bar{x}] \end{array} \qquad \mathcal{C}[\tau ! \varsigma] \triangleq \forall \phi, g. \phi \vdash_{\text{simple}} [\mathcal{C}[\tau = g]] \implies \text{shape}(g) = \varsigma$$

614 Fig. 3. Semantics of suspended constraints.
 615

616 The use of erasure $[\mathcal{C}[\tau = g]]$ in the definition of $\mathcal{C}[\tau ! \varsigma]$ ensures that the unicity of ς is deter-
 617 mined only by the constraints that have already been discharged in \mathcal{C} ; it excludes suspended match
 618 constraints, which may be discharged in the future. Implicitly, this induces a linear partial order
 619 between the suspended match constraints within a constraint, reflecting a *temporal* dependency: a
 620 match constraint may only be discharged once all of its dependencies have been discharged.
 621

622 The erasure $[\mathcal{C}[\tau = g]]$ is simple, so the use of \vdash_{simple} avoids well-foundedness issues that would
 623 arise from a negative occurrence of (\vdash) in a premise of *SUSP-CTX*. Note that, when τ is not a variable,
 624 then $\square[\tau ! \varsigma]$ holds trivially for $\varsigma = \text{shape}(\tau)$. Likewise, when \mathcal{C} is unsatisfiable, then $\mathcal{C}[\alpha ! \varsigma]$
 625 holds vacuously for any ς . The interesting cases arise when τ is a type variable and \mathcal{C} is satisfiable.

626 We summarize the definition of the unicity condition and *SUSP-CTX* in Figure 3. Together with
 627 the rules of Figure 2, this forms the complete semantics of our constraint language.

628 *Example 3.7.* Consider the two examples from above:

$$\exists \alpha. \alpha = \text{int} \wedge \text{match } \alpha \text{ with } _- \rightarrow \text{true} \qquad \exists \alpha. \text{match } \alpha \text{ with } _- \rightarrow \alpha = \text{int}$$

629 In the first example, we apply the contextual rule with the context $\mathcal{C} := \exists \alpha. \alpha = \text{int} \wedge \square$. Any
 630 solution ϕ of this context necessarily satisfies $\alpha = \text{int}$, so we have $\mathcal{C}[\alpha ! \text{int}]$ and the suspended
 631 constraint can be discharged. By contrast, the second example has no contextual information
 632 around the suspended constraint: $\mathcal{C} := \square$. So any solution ϕ satisfies it, allowing $\phi(\alpha)$ to have an
 633 arbitrary shape (e.g. `int`, `bool`, etc.). As a result, the uniqueness condition $\mathcal{C}[\alpha ! \varsigma]$ never holds and
 634 the constraint is unsatisfiable as intended.
 635

638 *Example 3.8.* Consider the more intricate example:

$$639 \quad \exists\alpha\beta. (\text{match } \alpha \text{ with } _ \rightarrow \beta = \text{bool}) \wedge (\text{match } \beta \text{ with } _ \rightarrow \text{true}) \wedge (\alpha = \text{int})$$

640
 Suppose we attempt to apply **SUSP-CTX** to the match on β first. We want to show $\mathcal{C}[\beta ! \text{bool}]$ for the context \mathcal{C} equal to $\text{match } \alpha \text{ with } (_ \rightarrow \beta = \text{bool}) \wedge \square \wedge \alpha = \text{int}$. Its erasure is $[\mathcal{C}] = \text{true} \wedge \square \wedge \alpha = \text{int}$. In this constraint β is unconstrained, so for example $[\mathcal{C}[\beta = \text{int}]]$ and $[\mathcal{C}[\beta = \text{bool}]]$ are both satisfiable: unicity does not hold and **SUSP-CTX** cannot be applied.

641
 Now consider instead applying **SUSP-CTX** to the match on α first. To do so, we must show that α has a uniquely determined shape in the context \mathcal{C} equal to $\square \wedge \text{match } \beta \text{ with } _ \rightarrow \text{true} \wedge \alpha = \text{int}$. Its erasure $[\mathcal{C}]$ is $\square \wedge \text{true} \wedge \alpha = \text{int}$. Since α is unified with int in the erasure, we have $\mathcal{C}[\alpha ! \text{int}]$. We may now discharge the match on α , rewriting it as $(\text{match } \alpha := \text{int} \text{ with } _ \rightarrow \beta = \text{bool})$, that is, $(\alpha = \text{int} \wedge \beta = \text{bool})$. Substituting back, we are left to satisfy the constraint $\mathcal{C}[\alpha = \text{int} \wedge \beta = \text{bool}]$, that is, $(\alpha = \text{int} \wedge \beta = \text{bool} \wedge \text{match } \beta \text{ with } _ \rightarrow \text{true} \wedge \alpha = \text{int})$.

642
 At this point, we can safely apply **SUSP-CTX** to the remaining match constraint on β . The unicity condition now holds, as the erasure of the context includes the discharged constraint $\beta = \text{bool}$, allowing us to eliminate the final match constraint.

643
 This demonstrates that suspended match constraints must be resolved in a dependency-respecting order: attempting to resolve a match constraint too early may result in unsatisfiability.

644 *Example 3.9.* Let us consider a constraint with a cyclic dependency between match constraints:

$$645 \quad \exists\alpha\beta. (\text{match } \alpha \text{ with } _ \rightarrow \beta = \text{bool}) \wedge (\text{match } \beta \text{ with } _ \rightarrow \alpha = \text{int})$$

646
 This constraint can be proved satisfiable under the “natural semantics” introduced earlier: by guessing the assignment $\alpha := \text{int}, \beta := \text{bool}$, the two match constraints succeed. However, our solver and the contextual semantics reject it.

647
 Without loss of generality, suppose we attempt to apply **SUSP-CTX** on α first. We must show $\mathcal{C}[\alpha ! \text{int}]$ where \mathcal{C} is $\square \wedge \text{match } \beta \text{ with } _ \rightarrow \alpha = \text{int}$. But the erasure $[\mathcal{C}]$ is $\square \wedge \text{true}$ imposes no constraint on α , so unicity fails, and **SUSP-CTX** cannot be applied.

648 *Example 3.10.* Considering the example `ex7` from §1:

649 `let ex7 r = let x = r.x in x + r.y`

OCaml OmniML

650
 The typing constraint generated for `ex7` contains the following, where α stands for the type of r :

$$651 \quad \exists\alpha\gamma. \text{let } x = \lambda\beta. (\text{match } \alpha \text{ with } \dots) \text{ in } x \text{ int} \wedge \alpha = \text{one } \gamma$$

652
 The suspended constraint can be discharged under our semantics, as intended. We apply the **SUSP-CTX** rule with context \mathcal{C} equal to $\text{let } x = \lambda\beta. \square \text{ in } x \text{ int} \wedge \alpha = \text{one } \gamma$. Although the context includes a **let**-binding—which in practice involves **let**-generalization—we can still deduce $\mathcal{C}[\alpha ! v\gamma'. \text{one } \gamma']$, since the erased context $[\mathcal{C}]$ contains the unification $\alpha = \text{one } \gamma$.

653
 This example illustrates that our formulation of suspended constraints interacts nicely with **let**-polymorphism. Although the two features are specified in a modular fashion, they are carefully crafted to work together, as we will further show in our next example.

654 *Example 3.11.* A subtle yet crucial feature of our semantics is its support for *backpropagation*:

655 `let ex8 = let getx r = r.x in getx {x = 1; y = 1}`

OCaml OmniML

656
 As in the previous example, the type of r cannot be disambiguated in the **let**-definition alone. In the previous example, this type was unified to a known shape in the **let**-body. Here, this is more subtle: an *instance* of the type scheme is taken, which is only well-typed if r has a variable type or a type of the form $\alpha \text{ one}$. The projection $r.x$ would be forbidden if r had a variable type, so

687 $e ::= x \mid () \mid \lambda x. e \mid e_1 e_2 \mid \text{let } x = e_1 \text{ in } e_2 \mid (e : \exists \bar{\alpha}. \tau) \mid \{\overline{l = e}\} \mid e.l$	Terms
688 $\mid (e_1, \dots, e_n) \mid e.j \mid e.j/n \mid [e] \mid [e : \exists \bar{\alpha}. \sigma] \mid \langle e \rangle \mid \langle e : \exists \bar{\alpha}. \sigma \rangle$	
689 $\tau ::= \alpha \mid 1 \mid \tau_1 \rightarrow \tau_2 \mid t \bar{\tau} \mid \prod_{i=1}^n \tau_i \mid [\sigma]$	Types
690 $\sigma ::= \tau \mid \forall \alpha. \sigma$	Type schemes
691 $\Gamma ::= \emptyset \mid \Gamma, x : \sigma$	Contexts
692	

Fig. 4. Syntax of OmniML.

696 α one is the unique solution. We call this flow of information from instances back to definitions
697 *backpropagation*.

698 The constraint generated when typing ex_8 is:

$$699 \quad \exists \alpha. \text{let } getx = \lambda \delta. \exists \beta, \gamma. (\delta = \beta \rightarrow \gamma \wedge \text{match } \beta \text{ with } \dots) \text{ in } getx \text{ (int one} \rightarrow \alpha)$$

700 With the context \mathcal{C} equal to $\text{let } getx = \lambda \delta. \exists \beta, \gamma. \delta = \beta \rightarrow \gamma \wedge \square$ in $getx \text{ (int one} \rightarrow \alpha)$, we
702 can show the unicity predicate $\mathcal{C}[\beta! \varsigma]$ for the shape $\varsigma = (v\gamma. \gamma \text{ one})$. For any ϕ, g , the erasure
703 $[\mathcal{C}[\beta = g]]$ is $\text{let } getx = \lambda \delta. \exists \beta, \gamma. \delta = \beta \rightarrow \gamma \wedge \beta = g$ in $getx \text{ (int one} \rightarrow \alpha)$. Since $getx$ is bound
704 to the constraint abstraction $\lambda \delta. \exists \gamma. \delta = (g \rightarrow \gamma)$, the instantiation $getx \text{ (int one} \rightarrow \alpha)$ can only
705 be satisfied when $g = \text{int one}$. This proves unicity, hence ex_7 is accepted by our semantics.

4 The OmniML calculus

708 To prove correctness of constraint generation, we must first define a surface language and its type
709 system. Surprisingly, identifying an appropriate declarative type system to use as a specification is
710 itself an interesting problem! In particular, naïve specifications often fail to ensure principal types.

711 Take overloaded tuple projections *à la SML*. We can ask the user to provide the length of the
712 tuple explicitly, via an annotated syntax $e.j/n$, which has a simple typing rule (**PROJ-X**).

$$\frac{\begin{array}{c} \text{PROJ-X} \\ \Gamma \vdash e : \prod_{i=1}^n \tau_i \quad 1 \leq j \leq n \end{array}}{\Gamma \vdash e.j/n : \tau_j} \quad \frac{\begin{array}{c} \text{PROJ-I-NAT} \\ \Gamma \vdash e : \prod_{i=1}^n \tau_i \quad 1 \leq j \leq n \end{array}}{\Gamma \vdash e.j : \tau_j}$$

717 On the other hand, the natural typing rule for the fragile construct $e.j$ breaks principality (**PROJ-I-NAT**).
718 The term $e.j$ admits infinitely many typings for e , provided the tuple is of sufficient length. This
719 is the exact same issue we had with the naïve semantics of suspended constraints, and in fact we
720 solve it in the same way, with a unicity condition and a contextual rule (**PROJ-I**) that transforms the
721 fragile, implicit construct into the robust, explicit counterpart:

$$\frac{\begin{array}{c} \text{PROJ-I} \\ \mathcal{E}[e \triangleright v\bar{\gamma}. \prod_{i=1}^n \bar{\gamma}] \quad \Gamma \vdash \mathcal{E}[e.j/n] : \tau \end{array}}{\Gamma \vdash \mathcal{E}[e.j] : \tau}$$

4.1 Syntax

726 In Figure 4, we give the grammar for our calculus. Terms include all of the ML calculus: variables x ,
727 the unit literal $()$, lambda-abstractions $\lambda x. e$, applications $e_1 e_2$, annotations $(e : \exists \bar{\alpha}. \tau)$ and let-
728 bindings $\text{let } x = e_1 \text{ in } e_2$. Our extensions include:

- 731 (1) Overloaded variant constructors and record labels, modeled using record literals $\{l_1 =$
732 $e_1; \dots; l_n = e_n\}$ and field projections $e.l$. Variant constructors are not treated formally in
733 OmniML, but behave analogously in practice.
- 734 (2) Tuples (e_1, \dots, e_n) with implicit projections $e.j$ and explicit projections $e.j/n$.

- 736 (3) For semi-explicit first-class polymorphism, we have implicit and explicit introduction and
 737 elimination forms: boxing $[e]$ and $[e : \exists\bar{\alpha}.\sigma]$, and unboxing $\langle e \rangle$ and $\langle e : \exists\bar{\alpha}.\sigma \rangle$.

738 We use the metavariable e^i to range over the fragile/implicit constructions, and e^x to range over
 739 their explicit counterpart.

740 4.2 Typing rules and unicity

741 We have detailed typing rules for the full OmniML calculus, but unfortunately they do not fit in
 742 the margins of the 25 pages of this document. We moved them all, along with detailed examples, in
 743 Appendix §A.

744 Our typing rules $\Gamma \vdash e : \sigma$ are mostly standard, except for the rules governing implicit (or
 745 fragile) constructs e^i . These rules are inspired by our contextual constraint semantics (§3): each is a
 746 contextual typing rule paired with a unicity condition and an elaboration into an explicit form.

747 The unicity condition requires that the shape ς is fully determined by the surrounding term
 748 context \mathcal{E} , including any subexpressions (e.g. e in $e.j$). They are analogous to the unicity condition
 749 $\mathcal{C}[\tau ! \varsigma]$ for constraints, though the analogy is not exact. Different fragile features require slightly
 750 different formulations, depending on whether they infer a unique shape for a subexpression $\mathcal{E}[e \triangleright \varsigma]$
 751 or for the expected type of the context $\mathcal{E}[e \triangleleft \varsigma]$.

752 In order to define the unicity conditions, we introduce *typed holes* $\{e\}$, which
 753 allow any well-typed term e to be treated as if it had any type (via MAGIC). Types
 754 holes are forbidden in the source language—they are a device solely used to define
 755 unicity conditions. We also introduce an erasure function $[e]$, the term counterpart
 756 of constraint erasure $[C]$, which erases all not-yet-elaborated implicit constructs e^i in e with typed
 757 holes around their subterms. This ensures the subterms—such as type annotations—remain present,
 758 so that any constraints they introduce can still contribute to unicity. For example, $[e.j]$ is $\{[e]\}$.
 759 The full definition is given in Appendix §C.

760 We can now formalize the two unicity conditions:

$$\mathcal{E}[e \triangleright \varsigma] \triangleq \forall \Gamma, \tau, g, \Gamma \vdash [\mathcal{E}[\{(e : g)\}]] : \tau \implies \text{shape}(g) = \varsigma$$

$$\mathcal{E}[e \triangleleft \varsigma] \triangleq \forall \Gamma, \tau, g, \Gamma \vdash [\mathcal{E}[\{(\{e\} : g)\}]] : \tau \implies \text{shape}(g) = \varsigma$$

$$\frac{\text{MAGIC}}{\Gamma \vdash e : \tau} \quad \frac{\Gamma \vdash e : \tau}{\Gamma \vdash \{e\} : \tau'}$$

761 We use the unicity condition $\mathcal{E}[e \triangleright \varsigma]$ when we disambiguate using the type of a subterm, as
 762 in overloaded tuple projections, record projections, and polytype unboxing. Conversely, we use
 763 $\mathcal{E}[e \triangleleft \varsigma]$ for polytype boxing and overloaded records, where we disambiguate them using the
 764 expected type of the context.

765 *Example 4.1.* Let e be let $f = \lambda x. x.1$ in $f(1, 2)$. e is well-typed using *backpropagation*. e is of
 766 the form $\mathcal{E}[x]$ where \mathcal{E} is the context let $f = \lambda x. \square$ in $f(1, 2)$. We have $\emptyset \vdash \mathcal{E}[x.1/2] : \text{int}$. Let us
 767 show that $\mathcal{E}[x \triangleright \gamma_1, \gamma_2. \gamma_1 * \gamma_2]$. Assume $\emptyset \vdash \mathcal{E}[\{(x : g)\}] : \tau$. As g is a ground type, the type g of
 768 x is not a variable. Then, g cannot be that of an arbitrary sized tuple, since there is no such type
 769 for a tuple of arbitrary size. Hence, g must be a tuple $\Pi_{i=1}^n \bar{\tau}$ for some size n . Since the codomain of
 770 f must be a tuple of size 2 (for $f(1, 2)$ to be well-typed), then n must also be 2. This shows that
 771 $\mathcal{E}[x \triangleright \gamma_1, \gamma_2. \gamma_1 * \gamma_2]$.

772 4.3 Metatheory

773 Constraint generation is sound and complete with respect to the typing judgment. That is to say,
 774 the term e is typable with τ if and only if $\llbracket e : \alpha \rrbracket$ is satisfiable when α is τ .

775 THEOREM 4.2 (CONSTRAINT GENERATION IS SOUND AND COMPLETE). *Given a closed term e and
 776 type τ . Then for any $\alpha \# \tau$, $\vdash e : \tau$ iff $\alpha = \tau \models \llbracket e : \alpha \rrbracket$.*

THEOREM 4.3 (PRINCIPAL TYPES). *For any well-typed closed term e , there exists a type τ , which we call principal, such that: (i) $\vdash e : \tau$. (ii) For any other typing $\vdash e : \tau'$, then $\tau' = \theta(\tau)$ for some substitution θ .*

It is also interesting to discuss the stability of typing by common program transformations.

Application equi-typability does hold. The expressions $f\ e_1\ e_2$ and $\text{swap}\ f\ e_1\ e_2$ are equitypable where swap is $\lambda f. \lambda x_1. \lambda x_2. f\ x_2\ e_1$. We also have that $f\ e$ and $\text{app}\ f\ e$ and $\text{rev_app}\ e\ f$ are equitypable, where app and rev_app are the application function $\lambda f. \lambda x. f\ x$ and the reverse application function $\lambda x. \lambda f. f\ x$, respectively. It is well-known that bidirectional types inference breaks application equi-typability. Both π -directional and omnidirectional type inference preserve it.

Factorization does not hold. If $\Gamma \vdash e[x := e_0] : \tau$ with x appearing in e , we do not necessarily have $\Gamma \vdash \text{let } x = e_0 \text{ in } e : \tau$. This is not a defect of our system, but a general property of all systems that support static overloading: the expanded term $e[x := e_0]$ can pick a different overloading choice for each occurrence of e_0 , and if they are incompatible the factored form may not typecheck.

Inlining does not hold. If $\Gamma \vdash \text{let } x = e_0 \text{ in } e : \tau$, we do not necessarily have $\Gamma \vdash e[x := e_0] : \tau$. This is specific to our support for *backpropagation*: the let-form will use information from all occurrences of x in e to resolve fragile constructs in e_0 , but in the inlined form each copy of e must resolve its implicit constructs independently, and it has access to less information to establish unicity. As a result, the implicit OmniML calculus does not preserve typability in its operational semantics.

5 Solving constraints

We now present a machine for solving constraints in our language. The solver operates as a rewriting system on constraints $C \longrightarrow C'$. Once no further transitions are applicable, *i.e.*, $C \longrightarrow$, the constraint C is either in solved form—from which we can read off a most general solution—or the solver becomes stuck, indicating that C is unsatisfiable.

Definition 5.1 (Solved form \hat{U}). A solved form is a constraint \hat{U} of the form $\exists \bar{\alpha}. \bigwedge_{i=1}^n \epsilon_i$, where: (1) each ϵ_i contains at most one non-variable type; (2) head variables do not occur in multiple equations; (3) the constraint is acyclic.

5.1 Unification

Our constraints ultimately reduce to equations between types, which we solve using first-order unification. Like our solver, we specify unification as a non-deterministic rewriting relation between *unification problems* $U_1 \longrightarrow U_2$, that eventually reduces to a solved form \hat{U} or to false.

$$\begin{array}{lll} U ::= \text{true} \mid \text{false} \mid U_1 \wedge U_2 \mid \exists \alpha. U \mid \epsilon & \text{Unification problems} & \text{MULTI-UNIF} \\ \epsilon ::= \emptyset \mid \tau = \epsilon & \text{Multi-equations} & \frac{\forall \tau \in \epsilon. \phi(\tau) = g}{\phi \vdash \epsilon} \\ C ::= \dots \mid \epsilon & \text{Constraints} & \end{array}$$

Fig. 5. Syntax and semantics of unification problems.

Unification problems U (Figure 5) are a restricted subset of constraints, extended with *multi-equations* [Pottier and Rémy 2005]—a multi-set of types considered equal. These generalize binary equalities: ϕ satisfies a multi-equation ϵ if all of its members are mapped to a single ground type g (**MULTI-UNIF**). Multi-equations are considered equal modulo permutation of their members.

Our algorithm is largely standard, with its main novelty being the use of *canonical principal shapes* in place of type constructors. This uniform treatment of monotypes and polytypes simplifies

unification and improves on the previous treatment of polytype unification [Garrigue and Rémy 1999]. For a detailed discussion of the unification rules, see §B (appendix).

5.2 Solving rules

We now gradually introduce the rules of the constraint solver itself (Figures 6, 8 and 10). These rules define a non-deterministic rewriting system, operating modulo α -equivalence, and the associativity and commutativity of conjunction. Rewriting takes place under an arbitrary one-hole constraint context \mathcal{C} . A constraint C is satisfiable if it rewrites to a solved form \hat{U} (Definition 5.1); otherwise it gets stuck.

$$\begin{array}{c}
 \frac{\text{S-UNIF} \quad U_1 \quad U_1 \longrightarrow U_2}{U_2} \qquad \frac{\text{S-FALSE} \quad \mathcal{C}[\text{false}] \quad \mathcal{C} \neq \square}{\text{false}} \qquad \frac{\text{S-LET} \quad \text{let } x = \lambda\alpha. C_1 \text{ in } C_2}{\text{let } x \alpha [\emptyset] = C_1 \text{ in } C_2} \qquad \frac{\text{S-EXISTS-CONJ} \quad (\exists\alpha. C_1) \wedge C_2 \quad \alpha \# C_2}{\exists\alpha. C_1 \wedge C_2} \\
 \\
 \frac{\text{S-LET-EXISTSLEFT} \quad \text{let } x \alpha [\bar{\alpha}] = \exists\beta. C_1 \text{ in } C_2 \quad \beta \# \alpha, \bar{\alpha}, C_2}{\text{let } x \alpha [\bar{\alpha}, \beta] = C_1 \text{ in } C_2} \qquad \frac{\text{S-LET-EXISTSRIGHT} \quad \text{let } x \alpha [\bar{\alpha}] = C_1 \text{ in } \exists\beta. C_2 \quad \beta \# \alpha, \bar{\alpha}, C_1}{\exists\beta. \text{let } x = \lambda\bar{\alpha}. C_1 \text{ in } C_2} \\
 \\
 \frac{\text{S-LET-CONJLEFT} \quad \text{let } x \alpha [\bar{\alpha}] = C_1 \wedge C_2 \text{ in } C_3 \quad C_1 \# \alpha, \bar{\alpha}}{C_1 \wedge \text{let } x \alpha [\bar{\alpha}] = C_2 \text{ in } C_3} \qquad \frac{\text{S-LET-CONJRIGHT} \quad \text{let } x \alpha [\bar{\alpha}] = C_1 \text{ in } (C_2 \wedge C_3) \quad x \# C_3}{C_3 \wedge \text{let } x \alpha = C_1 \text{ in } C_2}
 \end{array}$$

Fig. 6. Basic rewriting rules $C_1 \longrightarrow C_2$

Basic rules. S-UNIF invokes the unification algorithm to the current unification problem. The unification algorithm itself is treated as a black box by the solver, so the system could be extended with any equational theory of types implemented by the unification algorithm.

$$C ::= \dots | \text{let } x \alpha [\bar{\alpha}] = C_1 \text{ in } C_2 \qquad \text{Constraints}$$

$$\phi(\lambda\alpha[\bar{\alpha}]. C) \triangleq \{\alpha[\phi[\alpha := g, \bar{\alpha} := \bar{g}]] \in \mathcal{R} : \phi[\alpha := g, \bar{\alpha} := \bar{g}] \vdash C\}$$

$$\frac{\text{LET R} \quad \phi \vdash \exists\alpha, \bar{\alpha}. C_1 \quad \phi[x := \phi(\lambda\alpha[\bar{\alpha}]. C_1)] \vdash C_2}{\phi \vdash \text{let } x \alpha [\bar{\alpha}] = C_1 \text{ in } C_2} \qquad \frac{\text{APP R} \quad \alpha[\phi'] \in \phi(x) \quad \phi(\tau) = \phi'(\alpha)}{\phi \vdash x \tau}$$

Fig. 7. Syntax and semantics of region-based let constraints.

In general, existential quantifiers $\exists\alpha. C$ are lifted to the nearest enclosing let, if one exists, or otherwise to the top of the constraint. The resulting existential prefix $\exists\bar{\alpha}$ is called a *region*. To make regions explicit, we introduce the syntax $\text{let } x \alpha [\bar{\alpha}] = C_1 \text{ in } C_2$, where α is the *root* of the region and $\bar{\alpha}$ are auxiliary existential variables. The order of $\bar{\alpha}$ is immaterial; regions are considered equal up to permutation of these variables.

Satisfiability of regional let-constraints is defined in Figure 7. The semantics of an abstraction with a region, written $\phi(\lambda\alpha[\bar{\alpha}]. C)$, is a set of *ground regions* that satisfy C . A ground region is a

satisfying interpretation for the region ϕ' with a designated *root* variable α , written $\alpha[\phi']$. Regional let-constraints strictly generalize ordinary constraint abstractions, as captured by the equivalence:

$$\text{let } x = \lambda\alpha. C_1 \text{ in } C_2 \equiv \text{let } x \alpha [\emptyset] = C_1 \text{ in } C_2$$

Rule S-LET rewrites let constraints into regional form. S-EXISTS-CONJ lifts existentials across conjunctions; S-LET-EXISTSLEFT and S-LET-EXISTSRIGHT lift existentials across let-binders; S-LET-CONJLEFT, S-LET-CONJRIGHT hoist constraints out of let-binders when they are independent of the local variables. Collectively, these lifting rules normalize the structure of each region into a block of existentially bound variables, whose body consists of a conjunction of solved multi-equations followed by a residual constraint—typically an application, let-binding, or suspended constraint.

OmniML-specific constraints, such as the label and polytype instantiations from §2.1 and §2.2, require no special treatment in our solver. Once their pattern variables are substituted—after solving a match constraint—they are desugared into constraints already handled by the solver.

$$\frac{\begin{array}{c} \text{S-MATCH-TYPE} \\ \text{match } \tau \text{ with } \bar{\chi} \quad \tau \notin \mathcal{V} \end{array}}{\text{match } \tau := \text{shape } (\tau) \text{ with } \bar{\chi}}$$

$$\frac{\begin{array}{c} \text{S-MATCH-VAR} \\ \mathcal{C}[\text{match } \alpha \text{ with } \bar{\chi}] \quad \alpha = \tau = \epsilon \in \mathcal{C} \end{array}}{\mathcal{C}[\text{match } \alpha := \text{shape } (\tau) \text{ with } \bar{\chi}]}$$

Fig. 8. Rewriting rules for suspended match constraints.

Suspended match constraints. S-MATCH-TYPE solves suspended match constraints whose scrutinee is a non-variable type τ by rewriting them using the sugar (match $\tau := \text{shape } (\tau)$ with $\bar{\chi}$), introduced in §2.

S-MATCH-VAR applies when the scrutinee is a variable α and the context \mathcal{C} proves that α is equal to some non-variable type τ , which establishes the unicity property $\mathcal{C}[\tau! \text{shape } (\tau)]$. To check whether a context \mathcal{C} proves an equality—or more generally, a multi-equation ϵ —we search for a decomposition $\mathcal{C} = \mathcal{C}_1[\epsilon \wedge \mathcal{C}_2]$ where $\text{fv}(\epsilon)$ is disjoint from the binders of \mathcal{C}_2 .

Let constraints. Application constraints can be solved by copying constraints:

$$\frac{\begin{array}{c} \text{S-LET-APP-BETA} \\ \text{let } x \alpha [\bar{\alpha}] = C_1 \text{ in } \mathcal{C}[x \tau] \quad \alpha, \bar{\alpha} \# \tau \quad x \# \text{bv}(\mathcal{C}) \end{array}}{\text{let } x = \lambda\alpha. C_1 \text{ in } \mathcal{C}[\exists \alpha, \bar{\alpha}. \alpha = \tau \wedge C_1]}$$

This resembles β -reduction, except that the original abstraction is retained. While correct for *simple* constraints, it may duplicate solving work across applications of the same abstraction. A more efficient approach first solves the abstraction once—e.g. reducing it to $\lambda\alpha[\bar{\alpha}]. \bar{\epsilon}$, where $\bar{\alpha}$ are generalizable variables—and then reuses the result at each application site by only copying the solved constraint ϵ . This mirrors ML generalization and instantiation, a connection formalized by Pottier and Rémy [2005], where $\lambda\alpha[\bar{\alpha}]. \bar{\epsilon}$ corresponds to the type scheme $\forall \bar{\alpha}. \vartheta(\alpha)$ and ϑ is the *mgu* of $\bar{\epsilon}$. This optimization underlies efficient implementations of HM inference, such as OCaml’s.

However, this approach *does not* extend to suspended constraints. To illustrate this, let us examine ex₇ (from §1):

```
let ex7 r = let x = r.x in x + r.y
```

OCaml OmniML

The generated typing constraint contains:

$$\exists \alpha, \gamma. \text{let } x = \lambda\beta. \text{match } \beta \text{ with } (t _) \rightarrow \mathcal{C}[t, \alpha, \beta] \text{ in } x \text{ int} \wedge \alpha = \text{one } \gamma$$

where $\mathcal{C}[t, \alpha, \beta]$ is $\Omega(\ell/t) \leq \alpha \rightarrow \beta$. Here, α stands for r’s type. The constraint remains suspended until r.y forces r’s type to be one. Crucially, the variable β (introduced inside the abstraction for

932 $C ::= \dots \exists i^x. C i[\alpha \rightsquigarrow \tau]$ 933 $\phi ::= \dots \phi[i := \phi']$	934 Constraints 935 Semantic environments	936 EXISTS-INST $\alpha[\phi'] \in \phi(x)$ $\frac{\phi[i := \phi'] \vdash C}{\phi \vdash \exists i^x. C}$	937 PARTIAL-INST $\phi(i)(\alpha) = \phi(\tau)$ $\phi \vdash i[\alpha \rightsquigarrow \tau]$
---	--	---	--

Fig. 9. The syntax and semantics of partial instantiations.

the type of y) is captured by the suspended match constraint that is not yet resolved at the point of solving the let constraint that binds x .

Nonetheless, to continue solving the let-body, we must assign a scheme to x . We naively pick $\forall \beta. \beta$. This appears unsound, since β will later unify with int once the match constraint is discharged. But it would be incomplete to lower β as a monomorphic variable. This motivates *partial type schemes*, our second novel mechanism for omnidirectional inference. Partial type schemes are type schemes that delay commitment to certain quantifications (e.g. β). Such *partially generalized* variables are treated as generalized, but can be incrementally refined in future as suspended constraints are discharged.

To support this, we extend the constraint language with *partial instantiation constraints*. Instead of duplicating an abstraction at each application site, we introduce: (1) $\exists i^x. C$, which binds a fresh instantiation i of x 's region within C , and (2) $i[\alpha \rightsquigarrow \tau]$, which asserts that the copy of α in i equals τ . The instantiation variable i is required to ensure all partial instantiations $i[\alpha \rightsquigarrow \tau]$ are solved uniformly. Within the solver, we view partial instantiations as markers indicating which parts of the abstraction still need to be copied.

Partial instantiations enables efficient incremental instantiation of constraint abstractions: solved parts are reused immediately, while suspended constraints can be solved later, further refining the abstraction and propagation additional equations to the application sites.

The semantics of the existential constraint $\exists i^x. C$ (**EXISTS-INST**) introduces the fresh instantiation i by “guessing” a region ϕ' that satisfies the regional constraint abstraction bound to x . Partial instantiations (**PARTIAL-INST**) equate the copy of α in i with τ . The domain of partial instantiation constraints must lie within the closure of the abstraction or among the regional variables of x . Consequently, the variables $\alpha, \bar{\alpha}$ bound by the let-constraint $\text{let } x \alpha [\bar{\alpha}] = C_1 \text{ in } C_2$ are bound not only in the body of the abstraction C_1 , but also in the constraint C_2 , where they may appear in partial instantiations of x via renamings—and only there. Hence, they cannot appear in C_2 when the corresponding variable x does not itself appear in C_2 .

Partial instantiation constraints are reduced using the following rules:

- (1) **S-INST-COPY** copies the shape of a type to the instantiation site, introducing fresh variables for each subcomponents and marking them with corresponding instantiation constraints. We write $i^x[\alpha \rightsquigarrow \tau]$ as a shorthand for $i[\alpha \rightsquigarrow \tau]$ when i is bound with $\exists i^x$ in the context. To ensure termination, the abstraction must contain acyclic types.
- (2) **S-INST-UNIFY** unifies two instantiations if they refer to the same source variable.

There are three cases in which an instantiation constraint is eliminated:

- (1) A nullary shape is copied and no further instantiations are needed (**S-INST-COPY**).
- (2) The copied variable β is polymorphic, and thus the instantiation constraint imposes no restriction (**S-INST-POLY**), provided no other instantiations of β remain (if not, then apply **S-INST-UNIFY**).
- (3) The copy is monomorphic and in scope, so we unify it directly (**S-INST-MONO**).

S-LET-SOLVE remove a let constraint when the bound term variable is unused and the abstraction is satisfiable. **S-COMPRESS** determines that a regional variable β is an alias for γ . We replace every

<p>981 S-EXISTS-LOWER</p> $\frac{\begin{array}{l} \text{let } x \alpha [\bar{\alpha}, \bar{\beta}] = C_1 \text{ in } C_2 \\ \exists \alpha, \bar{\alpha}. C_1 \text{ determines } \bar{\beta} \end{array}}{\exists \bar{\beta}. \text{let } x \alpha [\bar{\alpha}] = C_1 \text{ in } C_2}$	<p>982 S-LET-APPR</p> $\frac{\begin{array}{l} \text{let } x \alpha [\bar{\alpha}] = C \text{ in } \mathcal{C}[x \tau] \\ \gamma \# \tau \quad x \# \text{bv}(\mathcal{C}) \end{array}}{\text{let } x \alpha [\bar{\alpha}] = C \text{ in } \mathcal{C}[\exists \gamma, i^x. \gamma = \tau \wedge i[\alpha \rightsquigarrow \gamma]]}$
<p>983</p>	
<p>984</p>	
<p>985</p>	
<p>986 S-INST-COPY</p> $\frac{\begin{array}{l} \text{let } x \alpha [\bar{\alpha}] = C \text{ in } \mathcal{C}[i^x[\alpha' \rightsquigarrow \gamma]] \\ \alpha' \in \alpha, \bar{\alpha} \quad \neg \text{cyclic } (\mathcal{C}) \end{array}}{\text{let } x \alpha [\bar{\alpha}] = C \text{ in } \mathcal{C}[\exists \bar{\beta}' . \gamma = \varsigma \bar{\beta}' \wedge i^x[\bar{\beta} \rightsquigarrow \bar{\beta}']]}$	<p>987 S-INST-UNIFY</p> $\frac{i[\alpha \rightsquigarrow \gamma_1] \wedge i[\alpha \rightsquigarrow \gamma_2]}{i[\alpha \rightsquigarrow \gamma_1] \wedge \gamma_1 = \gamma_2}$
<p>988</p>	
<p>989</p>	
<p>990</p>	
<p>991 S-INST-POLY</p> $\frac{\begin{array}{l} \text{let } x \alpha [\bar{\alpha}] = \bar{\epsilon} \wedge C \text{ in } \mathcal{C}[i^x[\alpha' \rightsquigarrow \gamma]] \\ \forall \alpha'. \exists \alpha, \bar{\alpha} \setminus \alpha'. \bar{\epsilon} \equiv \text{true} \quad \alpha' \in \alpha, \bar{\alpha} \\ \alpha' \# C \quad i.\alpha' \# \text{insts}(\mathcal{C}) \quad x \# \text{bv}(\mathcal{C}) \end{array}}{\text{let } x \alpha [\bar{\alpha}] = \bar{\epsilon} \wedge C \text{ in } \mathcal{C}[\text{true}]}$	<p>992 S-INST-MONO</p> $\frac{\begin{array}{l} \text{let } x \alpha [\bar{\alpha}] = C \text{ in } \mathcal{C}[i^x[\beta \rightsquigarrow \gamma]] \\ \beta \notin \alpha, \bar{\alpha} \quad x, \beta \# \text{bv}(\mathcal{C}) \end{array}}{\text{let } x \alpha [\bar{\alpha}] = C \text{ in } \mathcal{C}[\beta = \gamma]}$
<p>993</p>	
<p>994</p>	
<p>995</p>	
<p>996</p>	
<p>997 S-LET-SOLVE</p> $\frac{\begin{array}{l} \text{let } x \alpha [\bar{\alpha}] = \bar{\epsilon} \text{ in } C \\ \exists \alpha, \bar{\alpha}, \bar{\epsilon} \equiv \text{true} \quad x \# C \end{array}}{C}$	<p>998 S-COMPRESS</p> $\frac{\begin{array}{l} \text{let } x \alpha [\bar{\alpha}, \bar{\beta}] = C_1 \wedge \beta = \gamma = \bar{\epsilon} \text{ in } C_2 \\ \beta \neq \gamma \end{array}}{\text{let } x \alpha [\bar{\alpha}] = C_1[\beta := \gamma] \wedge \gamma = \bar{\epsilon}[\beta := \gamma] \text{ in } C_2[x, \beta := \gamma]}$
<p>999</p>	
<p>1000</p>	
<p>1001</p>	
<p>1002 S-BACKPROP</p> $\frac{\begin{array}{l} \mathcal{C}[\text{let } x \alpha [\bar{\alpha}] = \mathcal{C}_1[\text{match } \alpha' \text{ with } \bar{\chi}] \text{ in } \mathcal{C}_2[i^x[\alpha' \rightsquigarrow \gamma]]] \\ \alpha' \in \alpha, \bar{\alpha} \quad \alpha' = \tau = \bar{\epsilon} \in \mathcal{C}[\mathcal{C}_2] \quad x \# \text{bv}(\mathcal{C}_2) \end{array}}{\mathcal{C}[\text{let } x \alpha [\bar{\alpha}] = \mathcal{C}_1[\text{match } \alpha' := \text{shape } (\tau) \text{ with } \bar{\chi}] \text{ in } \mathcal{C}_2[i^x[\alpha' \rightsquigarrow \gamma]]]}$	
<p>1003</p>	
<p>1004</p>	
<p>1005</p>	
<p>1006</p>	
<p>1007</p>	
<p>1008</p>	

Fig. 10. Select rewriting rules for let-bindings and applications.

1010 free occurrence of β with γ —*including* the domains of any partial instantiation constraints, written
 1011 as the substitution $[x.\beta := \gamma]$. We view this as an analogous copy rule for variables.

1012 **S-EXISTS-LOWER** implements the non-trivial case of lowering existentials across let-binders. It
 1013 identifies a subset of variables in the region of a let constraint that are unified with variables from
 1014 outside the region. Such variables are considered monomorphic and thus cannot be generalized;
 1015 they can instead be safely lowered to the outer scope.

1016 This is the case when the types of $\bar{\beta}$ are *determined* in a unique way. In short, C determines $\bar{\beta}$ if
 1017 and only if the solutions for $\bar{\beta}$ are uniquely fixed by the solutions to other variables in C .

1018 Definition 5.2. C determines $\bar{\beta}$ if and only if every ground assignments ϕ and ϕ' that satisfy (the
 1019 erasure of) C and coincide outside of $\bar{\beta}$ coincide on $\bar{\beta}$ as well.

$$C \text{ determines } \bar{\beta} \triangleq \forall \phi, \phi'. \phi \vdash [C] \wedge \phi' \vdash [C] \wedge \phi =_{\bar{\beta}} \phi' \implies \phi = \phi'$$

1021 Conceptually, this corresponds to the negation of the generalization condition in ML: a type
 1022 variable cannot be generalized if it appears in the typing context. In the constraint setting, it cannot
 1023 be generalized if it depends on variables from outside the region.

1024 To decide when $\exists \bar{\alpha}. C$ determines $\bar{\beta}$ holds for $\bar{\beta} \# \bar{\alpha}$, we search for a multi-equation ϵ in C of
 1025 the form: (1) $\gamma = \epsilon'$ where $\gamma \# \bar{\alpha}, \bar{\beta}$ and $\bar{\beta} \subseteq \text{fv}(\epsilon')$, or (2) $\bar{\beta} = \tau = \epsilon'$ where $\text{fv}(\tau) \# \bar{\alpha}, \bar{\beta}$. For instance,
 1026 $\exists \beta_1. \alpha = \beta_1 \rightarrow \beta_2$ determines β_2 , as β_2 is free.

1029

1030 Lowering such variables improves solver efficiency. It avoids unnecessary duplication of work
 1031 that would otherwise occur via **S-INST-COPY**. By reducing the number of variables that need to be
 1032 copied, lowering directly reduces instantiation overhead.

1033 *Backpropagation.* Finally, **S-BACKPROP** expresses *backpropagation*, previously illustrated in [Example 3.11](#). In particular, the shape of a regional variable can sometimes be determined from its
 1034 instantiations. If an abstraction contains a suspended match constraint on a regional variable α' , and
 1035 the constraint context includes a partial instantiation $i^x[\alpha' \rightsquigarrow \gamma]$ together with a multi-equation
 1036 constraining the copy of α' (γ) to a non-variable type τ , then $\text{shape } (\tau)$ must be the unique shape
 1037 of α' . Any other shape would render the instantiation unsatisfiable.
 1038

1039 **THEOREM 5.3 (PROGRESS).** *If constraint C cannot take a step $C \longrightarrow C'$, then either:*

- 1040 (i) C is solved.
- 1041 (ii) C is stuck, it is either: (a) false; (b) $\hat{\mathcal{C}}[x \tau]$ where $x \# \hat{\mathcal{C}}$; (c) $\hat{\mathcal{C}}[i^x[\alpha \rightsquigarrow \gamma]]$ where $x \# \hat{\mathcal{C}}$ and
 1042 $i.\alpha \# \text{insts}(\hat{\mathcal{C}})$; (d) for every match constraint $\hat{\mathcal{C}}[\text{match } \alpha \text{ with } \bar{\chi}]$ in C , $\hat{\mathcal{C}}[\alpha !\varsigma]$ does not
 1043 hold for any ς . Here, $\hat{\mathcal{C}}$ is a normal context i.e., such that no other rewrites can be applied.

1044 **THEOREM 5.4 (TERMINATION).** *The constraint solver terminates on all inputs.*

1045 **THEOREM 5.5 (PRESERVATION).** *If $C_1 \longrightarrow C_2$, then $C_1 \equiv C_2$.*

1046 6 Implementation

1047 We have two working prototypes implementing the OmniML language with suspended match
 1048 constraints and partial type schemes, in which we have reproduced the various type-system fea-
 1049 tures and examples presented in this work. One closely follows the constraint-based presentation
 1050 described here⁵. It is public and open-source (link omitted for anonymity). Its implementation is
 1051 inspired by previous work such as Inferno [[Pottier 2014, 2018](#)]. It uses state-of-the-art implemen-
 1052 tation techniques for efficiency, such as a Tarjan’s union-find data structure for unification [[Tarjan 1053 1975](#)] and *ranks* (or *levels*) for efficient generalization [[Rémy 1992](#)]. Let us discuss a few salient
 1054 points.

1055 **Unification and scheduling.** Each unsolved unification variable maintains a *wait list* of suspended
 1056 constraints that are blocked until the variable is unified with a concrete type. When such a unification
 1057 occurs, the wait list is flushed: the suspended constraints are scheduled on the global constraint
 1058 scheduler, which is responsible for eventually solving them.

1059 **From a stack to a tree.** Many standard ML implementations, for example Inferno, represent
 1060 the solver state as a linear *stack* of inference regions, from the outermost variable scope to the
 1061 current region. Unification associates an integer *rank* (or *level*) for each variable, that indexes the
 1062 region in the stack to which it belongs. This approach does not work for partial generalization. If
 1063 generalization at some region is suspended by a match constraint, the region must remain alive
 1064 while we continue inference in other regions. However, later parts of the constraint may introduce
 1065 a new let-region at the same rank that is unrelated to the suspended one—neither its ancestor nor
 1066 its descendant—breaking the linear assumption of ranks.

1067 Our implementation must instead use a *tree* of nested let-regions. Under this scheme, ranks
 1068 no longer uniquely determine a variable’s region. Instead, we interpret a rank relative to a path
 1069 in the region tree from the root. When two variables are unified, they must always lie on some
 1070 shared path—by scoping invariants—so computing their minimum rank (along this path) suffices to
 1071 determine the lowered region: we keep the efficient integer comparisons of generalization.

1072 ⁵The other prototype is a direct implementation of type inference based on semi-unification. We mention it here only to
 1073 indicate that we have explored multiple implementation strategies leading to the same results.

1079 *Partial generalization.* Partial generalization arises when a region cannot be fully generalized
 1080 due to suspended constraints that may still update its variables. To manage this, we classify type
 1081 variables into four categories:

1082 (I) Variables are yet to be generalized.

1083 *Introduced by instantiations or source types in constraints*

1084 (G) Variables that are generalized.

1085 *Not accessible from any instance type. Treated polymorphically.*

1086 (PG) Variables that are partially generalizable.

1087 *Generalizable variables mentioned by suspended match constraint or partial instantiations.*

1088 (PI) Variables that were previously partially generalized but have since been updated.

1089 *Awaiting re-generalization. Introduced by the unification of partial generics.*

1090 At generalization time, we conservatively approximate whether a variable may be updated in the
 1091 future using *guards*. A guard is a mark on a variable that indicates the variable is captured by some
 1092 suspended constraint that has not yet been solved. Guarded variables are generalized as partial
 1093 generics (PG); unguarded ones are fully generalized (G).

1094 When an instance is taken from a partial generic, we retain a forward reference from the partial
 1095 generic (PG) to the instance. This enables the generic to notify the instance that it has been updated,
 1096 propagating the updated type structure to all instances. This mirrors, in reverse, the way our
 1097 formalized solver uses partial instantiation constraints to track copies. In addition, the instance
 1098 remains guarded by the partial generic until the latter is either lowered or fully generalized.

1099 Once a suspended match constraint is solved, it removes the guards it introduced. This may enable
 1100 previously partial generics to become fully generalizable. Conversely, if a partially generalized
 1101 variable is lowered (e.g. by [S-LOWER-EXISTS](#)), it must be unified with all its instances.

1102 *Lazy generalization.* Repeatedly generalizing a region after every update is expensive. Instead
 1103 we generalize on demand. We mark regions as “stale” when they may require re-generalization.
 1104 When an instance is taken, we re-generalize the stale descendants of the region in the region tree.

1106 7 Related work

1107 *Principality tracking in OCaml.* [Garrigue and Rémy \[1999\]](#) introduced an approach to principality
 1108 tracking for polytypes—what we now call π -directional inference—in which generalization and
 1109 instantiation govern the flow of known type information. This approach has since been extended to
 1110 other features of the OCaml language: whenever the typechecker need to know if a type is known
 1111 in a robust way, it checks whether the type is generalizable. We compared their approach to ours
 1112 in §1 and §2.1.

1113 *Bidirectional type inference.* At the level of simply-typed terms or ML, we believe that omnidirectional
 1114 inference works better than bidirectional type inference: it can type more programs than
 1115 a given fixed bidirectional system, and has a more declarative specification—we would say that
 1116 it is “more principal”. In fact, a direct inspiration for the present work was a user complaint in
 1117 [Rossberg \[2016\]](#) on the type-based disambiguation of OCaml: its bidirectional logic propagates
 1118 type information from patterns to definitions in `let`-bindings, when the WebAssembly reference
 1119 implementation would sometimes prefer the other direction. On the other hand, bidirectional typing
 1120 is known to scale to powerful systems such as fully-implicit predicative polymorphism [[Dunfield](#)
 1121 and [Krishnaswami 2013](#)] and we have not considered scaling our approach to those systems yet.

1122 *Qualified types.* Qualified types [[Jones 1995](#)], most well-known via their usage in Haskell type-
 1123 classes, are related to our suspended match constraints as they represent constraints on types or
 1124 type variables. At generalization time, the constraints on generalizable variables are kept as part

1128 of the generalized type scheme, and they get copied during instantiation. This is much simpler to
 1129 implement than our partial type schemes, but it provides a different behavior where each instance
 1130 can choose independently how to resolve the constraint. Qualified types are an excellent choice
 1131 when this is the desired behavior, typically for dynamic overloading. To handle cases that require a
 1132 unique resolution of the constraint across all instances—such as static overloading—we require the
 1133 more complex mechanism of partial generalization.

1134 *Suspended constraints in dependent-type systems.* Suspending the constraints that cannot be
 1135 solved yet is not a novel idea: it is a standard approach to implement unification dependently-typed
 1136 systems. This goes back to Huet’s algorithm for higher-order unification [Huet 1975] and pattern
 1137 unification [Miller 1991] where flexible-flexible pairs are delayed until at least one side becomes
 1138 rigid. The novelty of our work lies in combining constraint suspension with ML-style implicit
 1139 polymorphism—absent from most dependently-typed systems—and in the design of a declarative
 1140 constraint semantics used to establish principality.
 1141

1142 *OutsideIn.* OutsideIn [Schrijvers, Jones, Sulzmann and Vytiniotis 2009] is a type system for
 1143 GADTs that introduces *delayed implications* of the form $[\bar{\alpha}](\forall \bar{\beta}. C_1 \Rightarrow C_2)$. Constraint solving for
 1144 delayed implications proceeds in two steps; solving simple constraints first and then solving delayed
 1145 implications. The deferral ensures that inference for GADT match branches occurs when more
 1146 is known about the scrutinee and expected return type from the context. To ensure principality,
 1147 OutsideIn enforces an algorithmic restriction: the variables $\bar{\alpha}$ must already be instantiated to
 1148 concrete type constructors before they may be unified by the implication’s conclusion C_2 . This
 1149 ensures information only flows from the outside into the implication’s conclusion. They do not give
 1150 a declarative semantics for delayed implication that their solver preserves. Moreover, later work
 1151 on OutsideIn argues [Vytiniotis, Jones, Schrijvers and Sulzmann 2011] that delayed implication
 1152 constraints make local let-generalization all but unmanageable, both in theory and implementation.
 1153 Their proposed fix is to abandon local let-generalization altogether. We believe that we have solved
 1154 the troubling interactions between let-generalization and suspended constraints in this work, and
 1155 would be interested in studying applications to GADT typing, which was also one of our original
 1156 motivations.
 1157

1158 8 Conclusions and future work

1159 In this work, we developed a constraint-based framework for omnidirectional type inference,
 1160 capable of supporting fragile features that would otherwise break principality. Central to our
 1161 approach is a new declarative account of when type inference is *known* from the context, rather
 1162 than *guessed*.

1163 Our constraint solver is omnidirectional: constraints may be solved in any order, made possible
 1164 by our introduction of *partial type schemes*. We formalized the solver as a non-deterministic,
 1165 terminating rewrite system, and implemented an efficient prototype to demonstrate its practicality.

1166 Through three instantiations of our framework—static overloading of tuples, nominal records and
 1167 variant constructors, and semi-explicit first-class polymorphism—we showed that our framework
 1168 yields a sound and complete inference algorithm and a principal type system. In short, it appears
 1169 that principality holds *anyway* from our approach. Naturally, all this begs the question: what else can
 1170 be done with omnidirectional inference beyond the features of OmniML?

1171 *Static overloading.* Our nominal records use a restricted form of overloading. We have also
 1172 experimented with a more general overloading mechanism in which several definitions may be
 1173 bound to the same identifier $M.x$, but prefixed with a namespace⁶ M used for disambiguation: an
 1174

1175 ⁶Reusing the notation of Leijen and Ye [2025].
 1176

1177 implicit form x in the source is elaborated to an explicit form $M.x$. Although implemented in a
 1178 prototype, we have not yet formalized this feature. Nevertheless, we conjecture that it should be
 1179 typable with our framework.

1180 Modular implicits [White, Bour and Yallop 2014] are a proposed extension to OCaml’s module
 1181 system, intended to support ad-hoc polymorphism through type-directed implicit parameters. We
 1182 believe omnidirectional type inference could serve as a principled, constraint-based approach
 1183 foundation for resolving implicits in the presence of let-generalization. As future work, we aim to
 1184 extend our constraint language to typecheck an implicit-parameters calculus, similar to COCHIS
 1185 [Schrijvers, d. S. Oliveira, Wadler and Marntirosian 2019], but with ML polymorphism.
 1186

1187 *Higher-rank polymorphism.* In §1 and §7, we compared omnidirectional and bidirectional type
 1188 inference in the context of static overloading. While overloading is non-trivial, it poses little
 1189 challenge for the bidirectional framework, making the comparison somewhat limited. Bidirectional
 1190 typing is best known for its scalability to more complex settings, such as higher-rank polymorphism.
 1191 We are therefore interested in extending our framework to support higher-rank polymorphism, in
 1192 the style of Dunfield and Krishnaswami [2013]. This would provide a more meaningful basis for
 1193 understanding the trade-offs of omnidirectional and bidirectional inference.

1194 MLF is an extension of ML that support first-class polymorphism that goes beyond the power of
 1195 System F, while retaining type inference. It is a generalization of OCaml’s polytypes, relying on
 1196 π -directionality. It would therefore be worth exploring whether omnidirectional type inference
 1197 could further empower MLF.
 1198

1199 *Default rules.* Some type systems disambiguate fragile constructs using known type information,
 1200 but fall back on default, non-principal choices when none is available. OCaml selects the most
 1201 recent matching record type in scope for ambiguous field names; SML assigns default types to
 1202 overloaded numeric literals [Rossberg 2008, Section 5.8].
 1203

1204 We explored adding such *default rules* to suspended match constraints, allowing unresolved
 1205 constraints to discharge with a default shape rather than fail. While pragmatic, such rules are
 1206 inherently non-principal and difficult to reconcile with our framework.
 1207

1208 In particular, they introduce subtle semantic complexities: if two suspended constraints could
 1209 unlock each other, then defaulting one over the other may force an unsatisfiable branch to be taken.
 1210 The optimal or principled strategy for applying defaults in such cases remains unclear. Should
 1211 we fire all default rules of all suspended constraints that remain after the solver terminates, or in
 1212 batches, restricted to connected components of mutually dependent suspended constraints? Our
 1213 prototype opts for the latter, but this warrants further study.

1214 *Equi-recursive types.* OCaml allows equi-recursive types to express recursive polymorphic variants
 1215 and objects types. Supporting such types is a necessary step towards integrating our approach
 1216 into OCaml’s typechecker.
 1217

1218 In this work, for the sake of simplicity, we treat ground types as finite trees. Supporting equi-
 1219 recursive types amounts to using regular trees instead [Pottier and Rémy 2005]. Our prototype
 1220 already supports them, but the formalization of our solver relies on acyclicity to ensure termination.
 1221 Extending the formalization to accommodate cycles would require some changes. Following the
 1222 implementation, incremental instantiation might require to instantiate cycles atomically.
 1223

1224 Shapes may also be equi-recursive, though only minimal shapes of polytypes can be recursive.
 1225 In the acyclic setting, shape equality is syntactic; with cycles, this no longer holds—but we do not
 1226 anticipate any fundamental issues.
 1227

1226 References

- 1227 Arthur Charguéraud, Martin Bodin, Jana Dunfield, and Louis Riboulet. 2025. Typechecking of Overloading. In *Journées
1228 Francophones des Langages Applicatifs*.
- 1229 Luis Damas and Robin Milner. 1982. Principal type-schemes for functional programs. In *Proceedings of the 9th ACM SIGPLAN-
1230 SIGACT Symposium on Principles of Programming Languages* (Albuquerque, New Mexico) (*POPL '82*). Association for
1231 Computing Machinery, New York, NY, USA, 207–212. <https://doi.org/10.1145/582153.582176>
- 1232 Jana Dunfield and Neelakantan R. Krishnaswami. 2013. Complete and easy bidirectional typechecking for higher-rank
1233 polymorphism. In *ICFP*.
- 1234 Jacques Garrigue and Didier Rémy. 1999. Extending ML with Semi-Explicit Higher-Order Polymorphism. *Information and
1235 Computation* 155, 1/2 (1999), 134–169. <http://www.springerlink.com/content/m303472288241339/> A preliminary version
appeared in TACS'97.
- 1236 Gérard Huet. 1975. A unification algorithm for typed λ -calculus. *Theoretical Computer Science* 1, 1 (1975).
- 1237 Mark P. Jones. 1995. *Qualified types: theory and practice*. Cambridge University Press.
- 1238 Didier Le Botlan and Didier Rémy. 2009. Recasting MLF. *Information and Computation* 207, 6 (2009), 726–785. <https://doi.org/10.1016/j.ic.2008.12.006>
- 1239 Oukséh Lee and Kwangkeun Yi. 1998. Proofs about a folklore let-polymorphic type inference algorithm. *ACM Transactions
1240 on Programming Languages and Systems* 20, 4 (July 1998), 707–723. <https://doi.org/10.1145/291891.291892>
- 1241 Daan Leijen and Wenjia Ye. 2025. Principal Type Inference under a Prefix. In *PLDI'25*. ACM, 1–24. <https://www.microsoft.com/en-us/research/publication/principal-type-inference-under-a-prefix/> Backup Publisher: ACM SIGPLAN.
- 1242 Dale Miller. 1991. Unification of Simply Typed Lambda-Terms as Logic Programming. In *Logic Programming, Proceedings of
1243 the Eighth International Conference, Paris, France, June 24-28, 1991*, Koichi Furukawa (Ed.). MIT Press, 255–269.
- 1244 Martin Odersky, Martin Sulzmann, and Martin Wehr. 1999. Type inference with constrained types. *Theory and Practice of
1245 Object Systems* 5, 1 (1999), 35–55. [https://doi.org/10.1002/\(SICI\)1096-9942\(199901/03\)5:1<35::AID-TAPO4>3.0.CO;2-4](https://doi.org/10.1002/(SICI)1096-9942(199901/03)5:1<35::AID-TAPO4>3.0.CO;2-4)
- 1246 FranCois Pottier. 2014. Hindley-Milner elaboration in applicative style. [http://cambium.inria.fr/~fpottier/publis/fpottier-elaboration.pdf](http://cambium.inria.fr/~fpottier/publis/fpottier-
elaboration.pdf)
- 1247 FranCois Pottier. 2018. Inferno. A library for constraint-based Hindley-Milner type inference. [https://gitlab.inria.fr/fpottier/inferno](https://gitlab.inria.fr/fpottier/
1248 inferno) Available on opam <https://opam.ocaml.org/>.
- 1249 FranCcois Pottier and Didier Rémy. 2005. The Essence of ML Type Inference. In *Advanced Topics in Types and Programming
1250 Languages*, Benjamin C. Pierce (Ed.). MIT Press, Chapter 10, 389–489. <https://pauillac.inria.fr/~remy/attapl/>
- 1251 Didier Rémy. 1990. *Algèbres Touffues. Application au Typage Polymorphe des Objets Enregistrements dans les Langages
Fonctionnels*. Thèse de doctorat. Université de Paris 7.
- 1252 Didier Rémy. 1992. *Extending ML Type System with a Sorted Equational Theory*. Research Report 1766. Institut National de
1253 Recherche en Informatique et Automatisme, Rocquencourt, BP 105, 78 153 Le Chesnay Cedex, France.
- 1254 Andreas Rossberg. 2008. HaMLet: To Be Or Not To Be Standard ML. [https://people.mpi-sws.org/~rossberg/hamlet/hamlet-1.3.1.pdf](https://people.mpi-sws.org/~rossberg/hamlet/hamlet-
1255 1.3.1.pdf) hamlet manual, version 1.3.1.
- 1256 Andreas Rossberg. 2016. WebAssembly: high speed at low cost for everyone. presented at the ML Family Workshop.
- 1257 Tom Schrijvers, Bruno C. d. S. Oliveira, Philip Wadler, and Koar Marntirosian. 2019. COCHIS: Stable and coherent implicits.
J. Funct. Program. 29 (2019), e3. <https://doi.org/10.1017/S0956796818000242>
- 1258 Tom Schrijvers, Simon L. Peyton Jones, Martin Sulzmann, and Dimitrios Vytiniotis. 2009. Complete and decidable type
1259 inference for GADTs. In *Proceeding of the 14th ACM SIGPLAN international conference on Functional programming, ICFP
1260 2009, Edinburgh, Scotland, UK, August 31 - September 2, 2009*, Graham Hutton and Andrew P. Tolmach (Eds.). ACM,
341–352. <https://doi.org/10.1145/1596550.1596599>
- 1261 Robert Endre Tarjan. 1975. Efficiency of a Good But Not Linear Set Union Algorithm. *J. ACM* 22, 2 (1975), 215–225.
<https://doi.org/10.1145/321879.321884>
- 1262 Dimitrios Vytiniotis, Simon L. Peyton Jones, Tom Schrijvers, and Martin Sulzmann. 2011. OutsideIn(X) Modular type
1263 inference with local assumptions. *J. Funct. Program.* 21, 4-5 (2011), 333–412. <https://doi.org/10.1017/S0956796811000098>
- 1264 Leo White, Frédéric Bour, and Jeremy Yallop. 2014. Modular implicits. In *Proceedings ML Family/OCaml Users and Developers
1265 workshops, ML/OCaml 2014, Gothenburg, Sweden, September 4-5, 2014. (EPTCS, Vol. 198)*, Oleg Kiselyov and Jacques
Garrigue (Eds.), 22–63. <https://doi.org/10.4204/EPTCS.198.2>
- 1266
- 1267
- 1268
- 1269
- 1270
- 1271
- 1272
- 1273
- 1274

1275 Organization of our appendices

1276 *Content appendices.* These appendices are intended to be readable prose in the style of the rest of
 1277 the paper.

- 1278 • §A presents and explains the typing rules for OmniML; this is the long version of §4.
- 1279 • §B presents a relatively standard part of our constraint solver (§5), namely the unification
 1280 rules.

1281 *Reference appendix.*

- 1282 • §C gives a full reference for all definitions, grammars and figures in the paper, including all
 1283 cases (even those omitted from the main paper for reasons of space).

1284 *Proof appendices.* These appendices contain proofs for the formal claims in the article. They are
 1285 typically written tersely.

- 1286 • §D proves properties of the constraint language and its semantics. The main result is
 1287 canonicalization, which morally establishes that uses of the contextual rule **SUSP-CTX** can
 1288 be “permuted down” in the proof until they are all at the bottom of the derivation, followed
 1289 by a proof on a simple constraint.
- 1290 • §E proves the correctness of the constraint solver with respect to the semantics.
- 1291 • §F proves the properties about the OmniML type system, in particular the correctness of
 1292 constraint generation.

1293 A The OmniML calculus: typing rules and constraint generation

1294 A.1 Typing rules

1295 As usual, the main typing judgment $\Gamma \vdash e : \sigma$ states that in context Γ , expression e has type scheme
 1296 σ . Typing rules are given on Figure 11. They use auxiliary typing judgments $\Gamma \vdash \ell = e : \tau$ and
 1297 $l : \tau \rightarrow \tau'$ for the typing of record assignments and label instantiations respectively.

1298 **VAR** retrieves the type scheme $x : \sigma$ from the context Γ . Function types are introduced via lambda
 1299 abstractions: in **FUN**, the system guesses a well-formed type τ_1 for the type of x , typechecks the body
 1300 e is under the extended context $\Gamma, x : \tau_1$ producing the return type τ_2 , and assigns the abstraction
 1301 the function type $\tau_1 \rightarrow \tau_2$. Conversely, function types are eliminated by applications; in Rule **APP**,
 1302 the type of the argument must match the function’s parameter type τ_1 and application returns the
 1303 type τ_2 . **UNIT** asserts that () has the unit type 1.

1304 **GEN** and **INST** correspond to implicit *generalization* and *instantiation* respectively. Generalization
 1305 universally quantifies a type variable α , introducing it as a fresh polymorphic variable in the typing
 1306 context. In **INST**, we specialize a type scheme $\forall \alpha. \sigma$ to $\sigma[\alpha := \tau]$, substituting α for an arbitrary
 1307 monotype τ .

1308 Let-polymorphism is handled by the **LET** rule, where a *polymorphic* term can be bound. This
 1309 allows a single definition to be instantiated differently at each use site—an essential feature of
 1310 ML. In this rule, the term e_1 has a polymorphic type scheme σ , adds $x : \sigma$ into the context Γ to
 1311 typecheck e_2 .

1312 Annotations ($e : \exists \bar{\alpha}. \tau$) ensures that the type of e is (an instance of) the type τ . The type variables
 1313 $\bar{\alpha}$ are *flexibly* (or existentially) bound in τ , meaning that $\bar{\alpha}$ may be unified with some types $\bar{\tau}$ to
 1314 produce a well-typed term. For instance, the term $(\lambda x. x+1 : \exists \alpha. \alpha \rightarrow \alpha)$ is well-typed with $\alpha := \text{int}$
 1315 in **ANNOT**.

1316 *Polytypes and overloaded tuples.* The typing rules for fully annotated terms (e^x) are unsurprising.
 1317 However, typing rules for terms with omitted type annotations are non-compositional as they
 1318 depend on a surrounding one-hole context \mathcal{E} . Hence, they assert that the typability of the expression
 1319

1324	VAR	$x : \sigma \in \Gamma$	FUN	$\frac{}{\Gamma, x : \tau_1 \vdash e : \tau_2}$	APP	$\frac{\Gamma \vdash e_1 : \tau_1 \rightarrow \tau_2 \quad \Gamma \vdash e_2 : \tau_1}{\Gamma \vdash e_1 e_2 : \tau_2}$	UNIT	$\frac{}{\Gamma \vdash () : 1}$
1325								
1326								
1327								
1328	ANNOT	$\frac{}{\Gamma \vdash e : \tau[\bar{\alpha} := \bar{\tau}]}$	GEN	$\frac{\Gamma \vdash e : \sigma \quad \alpha \# \Gamma}{\Gamma \vdash e : \forall \alpha. \sigma}$	INST	$\frac{\Gamma \vdash e : \forall \alpha. \sigma}{\Gamma \vdash e : \alpha[\alpha := \tau]}$		
1329								
1330		$\frac{}{\Gamma \vdash (e : \exists \bar{\alpha}. \tau) : \tau[\bar{\alpha} := \bar{\tau}]}$						
1331								
1332	LET	$\frac{\Gamma \vdash e_1 : \sigma \quad \Gamma, x : \sigma \vdash e_2 : \tau}{\Gamma \vdash \text{let } x = e_1 \text{ in } e_2 : \tau}$	TUPLE	$\frac{(\Gamma \vdash e_i : \tau_i)_{i=1}^n}{\Gamma \vdash (e_1, \dots, e_n) : \prod_{i=1}^n \tau_i}$	PROJ-X	$\frac{\Gamma \vdash e : \prod_{i=1}^n \tau_i \quad 1 \leq j \leq n}{\Gamma \vdash e.j/n : \tau_j}$		
1333								
1334								
1335								
1336	PROJ-I	$\frac{\mathcal{E}[e \triangleright v \bar{\gamma}. \Pi_{i=1}^n \bar{\tau}_i] \quad \Gamma \vdash \mathcal{E}[e.j/n] : \tau}{\Gamma \vdash \mathcal{E}[e.j] : \tau}$	POLY-X	$\frac{\Gamma \vdash e : \sigma[\bar{\alpha} := \bar{\tau}]}{\Gamma \vdash [e : \exists \bar{\alpha}. \sigma] : [\sigma[\bar{\alpha} := \bar{\tau}]]}$				
1337								
1338								
1339								
1340	POLY-I	$\frac{\mathcal{E}[e \triangleleft v \bar{\gamma}. [\sigma]] \quad \Gamma \vdash \mathcal{E}[[e : \exists \bar{\gamma}. \sigma]] : \tau}{\Gamma \vdash \mathcal{E}[[e]] : \tau}$	USE-X	$\frac{\Gamma \vdash e : [\sigma][\bar{\alpha} := \bar{\tau}]}{\Gamma \vdash \langle e : \exists \bar{\alpha}. \sigma \rangle : \sigma[\bar{\alpha} := \bar{\tau}]}$				
1341								
1342								
1343								
1344	USE-I	$\frac{\mathcal{E}[e \triangleright v \bar{\gamma}. [\sigma]] \quad \Gamma \vdash \mathcal{E}[\langle e : \exists \bar{\gamma}. \sigma \rangle] : \tau}{\Gamma \vdash \mathcal{E}[\langle e \rangle] : \tau}$	RCD-ASSN	$\frac{\Gamma \vdash e : \tau \quad l : \tau \rightarrow \tau'}{\Gamma \vdash l = e : \tau'}$				
1345								
1346								
1347								
1348	RCD	$\frac{(\Gamma \vdash l_i = e_i : \tau)_{i=1}^n \quad \bar{l} ! \tau}{\Gamma \vdash \{l_1 = e_1; \dots; l_n = e_n\} : \tau}$	RCD-PROJ	$\frac{\Gamma \vdash e : \tau' \quad l : \tau \rightarrow \tau' \quad \bar{l} ! \tau}{\Gamma \vdash e.l : \tau}$	LAB-I	$\frac{\mathcal{L}[\ell ! t] \quad \Gamma \vdash \mathcal{L}[\ell / t] : \tau}{\Gamma \vdash \mathcal{L}[\ell] : \tau}$		
1349								
1350								
1351								
1352	LAB-X	$\Omega(\ell / t) = \forall \bar{\alpha}. \tau \rightarrow t \bar{\alpha}$	LAB-I	$\frac{\mathcal{L}[\ell ! t] \quad \Gamma \vdash \mathcal{L}[\ell / t] : \tau}{\Gamma \vdash \mathcal{L}[\ell] : \tau}$	LAB-!	$\frac{\bar{l} ! t \in \Omega}{\bar{l} ! t \bar{\tau}}$	LAB-?	$\bar{l} ! \tau$
1353								
1354		$\ell / t : \bar{\tau}[\bar{\alpha} := \bar{\tau}] \rightarrow t \bar{\tau}$						
1355								
1356								
1357								
1358								
1359								
1360								
1361								
1362								
1363								
1364								
1365								
1366								
1367								
1368								
1369								
1370								
1371								
1372								

Fig. 11. Typing rules of OmniML.

1358 $\Gamma \vdash \mathcal{E}[e^i] : \tau$ where e^i is an expression with an implicit type annotation. We first request a typing
 1359 for the expression with an explicit annotation $\Gamma \vdash \mathcal{E}[e^x] : \tau$ where e^x is a fully annotated variant
 1360 of e^i . We then request that (the shape of) the annotation is fully determined from context, either
 1361 from the type of the expression, which we write $\mathcal{E}[e \triangleright \varsigma]$, or from the type of the hole, which we
 1362 write $\mathcal{E}[e \triangleleft \varsigma]$.

1363 In order to describe the judgments $\mathcal{E}[e \triangleright \varsigma]$ and $\mathcal{E}[e \triangleleft \varsigma]$, we introduce a *typed hole* construct $\{e\}$
 1364 that allows any well-typed expression e to be treated as if it had any type. That is the typing rule
 1365 for holes is:

$$\frac{\text{MAGIC}}{\Gamma \vdash \{e\} : \tau'}$$

1370 Typed holes are not allowed on source terms and are just a device for the definition of non-
 1371 ambiguous shapes. Finally, we define what it means for a shape to be determined from the type of

1373 a context or an expression:

$$\begin{aligned} \mathcal{E}[e \triangleright \varsigma] &\triangleq \forall \Gamma, \tau, g, \Gamma \vdash [\mathcal{E}[\{(e : g)\}]] : \tau \implies \text{shape}(g) = \varsigma \\ \mathcal{E}[e \triangleleft \varsigma] &\triangleq \forall \Gamma, \tau, g, \Gamma \vdash [\mathcal{E}[(\{e\} : g)]] : \tau \implies \text{shape}(g) = \varsigma \end{aligned}$$

1378 These states that the shape ς of expression e in context \mathcal{E} is determined by the expression e , in
 1379 the former case, or by the context \mathcal{E} in the latter case. Just like constraints, we must erase implicit
 1380 constructs in the term that have not yet been elaborated, written $[e]$ (defined in §C).

1381 The implicit rule **PROJ-I** types the projection $e.j$ provided the context \mathcal{E} *infers* that the shape of e
 1382 must be a tuple with arity n . Similarly, **USE-I** permits instantiating a polytype in $\langle e \rangle$ if the context \mathcal{E}
 1383 infers that the type of e must be a polytype with shape $v\bar{y}. [\sigma]$. The rule **POLY-I** types the implicit
 1384 boxing construct $[e]$ by *checking* the expected type of $[e]$ in the context \mathcal{E} is a polytype with the
 1385 shape $v\bar{y}. [\sigma]$. This rule differs from the previous two as the shape is determined by the expected
 1386 type within the context as opposed to the inferred type of e .

1387 *Overloaded record labels.* We adopt a similar non-compositional approach for elaborating over-
 1388 loaded labels, whether in record projection $(e.\ell)$ or record construction $(\{\ell = e\})$, although the
 1389 definitions are slightly more involved. Here, a one-hole label context \mathcal{L} provides the surrounding
 1390 context in which a label ℓ may appear:

$$\mathcal{L} ::= \mathcal{E}[e.\square] \mid \mathcal{E}[\{l_1 = e_1; \dots; \square = e_i; \dots; l_n = e_n\}] \quad \text{Label contexts}$$

1394 As with our contextual rules for expressions, we define two rules for labels. **LAB-X** handles
 1395 explicitly annotated labels ℓ/t by instantiating the type scheme $\forall \bar{\alpha}. \tau \rightarrow t \bar{\alpha}$ associated with ℓ in
 1396 label environment Ω . **LAB-I** handles unannotated labels by elaborating ℓ to ℓ/t if the context \mathcal{L}
 1397 uniquely infers the record type t for ℓ and the resulting elaboration is well-typed.

1398 The unicity of the inferred record type is captured by the judgment $\mathcal{L}[\ell ! t]$. The definition fits
 1399 into the framework we established for expressions above by introducing analogous annotation and
 1400 hole constructs for labels.

$$\frac{\text{LAB-MAGIC}}{\Gamma \vdash \{\ell\} : \tau' \rightarrow \tau} \quad \frac{\text{LAB-ANNOT}}{\Gamma \vdash l : \tau' \rightarrow \tau[\bar{\alpha} := \bar{t}]} \quad \frac{\Gamma \vdash l : \exists \bar{\alpha}. \tau \rightarrow \tau[\bar{\alpha} := \bar{t}]}{\Gamma \vdash (l : \exists \bar{\alpha}. \tau) : \tau' \rightarrow \tau[\bar{\alpha} := \bar{t}]}$$

$$\mathcal{L}[\ell ! t] \triangleq \forall \Gamma, \tau, g, \Gamma \vdash [\mathcal{L}[(\{\ell\} : g)]] : \tau \implies \text{shape}(g) = v\bar{y}. \bar{y}t$$

1408 **RCD** types a record $\{\bar{l} = e\}$ as a record type τ provided that each field assignment $l = e$ can be
 1409 assigned the record type τ . **RCD-Assn** checks that e has the appropriate field type in $l = e$ and returns
 1410 the instantiated record type τ' for the label l . **RCD-Proj** types the projection $e.l$ by checking that the
 1411 type of e matches the record type associated with label l , returning the field type τ .

1412 Both **RCD** and **RCD-Proj** impose additional constraints on their record types to support *closed-*
 1413 *world* reasoning. These constraints exploit the uniqueness of type definitions in the global label
 1414 environment Ω to resolve overloaded labels: (1) in a record projection $e.\ell$, if the label ℓ is not
 1415 overloaded, then the global record typing context Ω assigns a unique record type t to ℓ ; (2) in a
 1416 record expression $\{l_1 = e_1; \dots; l_n = e_n\}$, if the set of labels l_1, \dots, l_n uniquely identifies a record
 1417 type t in the typing context Ω , then we can assign this type to the record expression.

1418 We formalize this with the judgment $\bar{l} ! \tau$, which either: (1) enforces τ to be of the form $t \bar{t}$ if the
 1419 labels \bar{l} uniquely identify a nominal record type t in Ω (**LAB-!**), or (2) imposes no constraint on τ in
 1420 the ambiguous case (**LAB-?**).

1422 Label declarations in Ω have the form $\ell : \forall \bar{\alpha}. \tau' \rightarrow t \bar{\alpha}$, assigning labels to field types and record
 1423 types⁷. We write $\ell/t \in \Omega$ if such a declaration of ℓ exists for the record type t . This membership
 1424 relation extends to explicitly annotated and casted labels:

$$\frac{\text{LAB-INX} \quad \ell/t \in \Omega}{(\ell/t)/t \in \Omega} \quad \frac{\text{LAB-INMAGIC} \quad \ell/t \in \Omega}{\{\ell\}/t \in \Omega} \quad \frac{\text{LAB-INANNOT} \quad l/t \in \Omega}{(l : \exists \bar{\alpha}. \tau)/\tau \in \Omega}$$

1429 We then define the uniqueness predicate $\bar{l}!t \in \Omega$ as:

$$\frac{\text{LAB-U} \quad \bar{l}/t \in \Omega \quad \forall t'. \bar{l}/t' \in \Omega \implies t = t'}{\bar{l}!t \in \Omega}$$

1433 This states that the set of labels \bar{l} determines a unique nominal type t in Ω if no other type t' can
 1434 be associated with the same label set.
 1435

1436 A.2 Examples of typings

1437 The following lemma shows that we can always take a larger context \mathcal{E} or \mathcal{L} for implicit rules
 1438 **Proj-I**, **Use-I**, **Poly-I** and **Lab-I**. That is, there is always a derivation using only toplevel contexts.

1440 LEMMA A.1. *If $\mathcal{E}_2[e \triangleright \varsigma]$, then $(\mathcal{E}_1[\mathcal{E}_2])[e \triangleright \varsigma]$. Similarly, for label contexts, if $\mathcal{L}[\ell ! t]$, then
 1441 $(\mathcal{E}[\mathcal{L}])[\ell ! t]$.*

1443 We now illustrate the typing of implicit constructs with a few examples.

1444 Example A.2. To illustrate a simple case of non-typability, we show that the expression e equal
 1445 to $\lambda x. x.k$ is ambiguous, i.e., that it does not typecheck. If there is a derivation of $\lambda x. x.k$ then there
 1446 must be one of the form:

$$\frac{\mathcal{E}[x \triangleright v\bar{\gamma}. \prod_{i=1}^n \bar{\gamma}_i] \quad \emptyset \vdash \mathcal{E}[x.k/n] : \tau}{\emptyset \vdash \mathcal{E}[x.k] : \tau} \text{ PROJ-I}$$

1450 where E is the term $\lambda x. \square$, which is the largest possible context, thanks to Lemma A.1. Let τ be
 1451 $\prod_{i=1}^n \tau_i \rightarrow \tau_k$ for some $n \geq k$. We have the following derivation:

$$\frac{\begin{array}{c} x : \prod_{i=1}^n \tau_i \vdash x : \prod_{i=1}^n \tau_i \\ \hline x : \prod_{i=1}^n \tau_i \vdash x.k/n : \tau_k \end{array}}{\emptyset \vdash \mathcal{E}[x.k/n] : \tau} \begin{array}{l} \text{PROJ-X} \\ \text{FUN} \end{array}$$

1457 Unfortunately, $\mathcal{E}[x \triangleright v\bar{\gamma}. \prod_{i=1}^n \bar{\gamma}_i]$ does not hold. Indeed, we have $\emptyset \vdash \mathcal{E}[\{(x : g)\}] : \tau$ for any g
 1458 assuming τ is of the form $g \rightarrow \tau'$. Hence, $v\bar{\gamma}. \prod_{i=1}^n \bar{\gamma}_i$ and $v\bar{\gamma}. \prod_{i=1}^{n+1} \bar{\gamma}_i$ are two possible shapes for the
 1459 type of x .

1460 Example A.3. We now illustrate a non-ambiguous example, showing that the expression e equal
 1461 to $\vdash (\lambda x. x.1) (1, 2) : \text{int}$. Let \mathcal{E} be the context $(\lambda x. \square) (1, 2)$. We have the derivation:

$$\frac{\mathcal{E}[x \triangleright v\gamma_1, \gamma_2. \gamma_1 * \gamma_2] \quad \emptyset \vdash \mathcal{E}[x.1/2] : \text{int}}{\emptyset \vdash \mathcal{E}[x.1] : \text{int}} \text{ PROJ-I}$$

1465 We have $\emptyset \vdash \mathcal{E}[x.1/2] : \text{int}$, indeed. Therefore, it just remains to show $\mathcal{E}[x \triangleright v\gamma_1, \gamma_2. \gamma_1 * \gamma_2] (1)$.
 1466 Assume $\emptyset \vdash \mathcal{E}[\{(x : g)\}] : \tau$. Since $x : \text{int} * \text{int}$ is bound in the context at the hole in \mathcal{E} , there is no
 1467 other choice but take g equal to $\text{int} * \text{int}$, hence shape $(g) = v\gamma_1, \gamma_2. \gamma_1 * \gamma_2$, which proves (1).

1468
 1469 ⁷For a given record type t , we assume each label associated with it is unique.

The following example of non-typability illustrates how the typing rules still forces to reject typing of some expressions whose elaboration would be unambiguous. This is intended, to prevent us from having to focus at several terms simultaneously. Our typing rules enforce the resolution of shape inference, locally, one at a time.

Example A.4. Let τ_{id} be $[\forall \alpha. \alpha \rightarrow \alpha]$. We show that the expression e equal to let $x = [\lambda z. z]$ in $(\langle x \rangle 1, \langle x \rangle ())$ is rejected as ambiguous. Let τ_{id} be $[\forall \alpha. \alpha \rightarrow \alpha]$. Clearly, we have let $x = [\lambda z. z : \tau_{\text{id}}]$ in $(\langle x : \tau_{\text{id}} \rangle 1, \langle x : \tau_{\text{id}} \rangle ())$. This is actually the only possible fully annotated derivation. To show that e is typable, we must be able to make all annotations optional, sequentially. Therefore, the final step, which will eliminate the last annotation has a single point of focus of the form $\mathcal{E}[e^i]$, where e^i can be any of the three positions with a missing annotation. We consider each case independently, and show that it is actually not typable.

Case \mathcal{E} is let $x = \square$ in $(\langle x \rangle 1, \langle x \rangle ())$. If this holds, we should have a derivation that ends with

$$\frac{\mathcal{E}[\lambda z. z \triangleleft [\tau_{\text{id}}]] \quad \emptyset \vdash \mathcal{E}[[\lambda z. z : \tau_{\text{id}}]] : \tau}{\emptyset \vdash \mathcal{E}[[\lambda z. z]] : \tau} \text{ POLY-I}$$

However, $\mathcal{E}[\lambda z. z \triangleleft [\tau_{\text{id}}]]$ does not hold. Indeed, the following judgment $\emptyset \vdash \mathcal{E}[(\{\lambda z. z\} : [\sigma])] : \tau$ holds, where σ is either $\forall \alpha. \alpha \rightarrow \alpha$ or $\forall \alpha. \alpha \rightarrow \alpha \rightarrow \alpha$. Hence, the shape of the type of $\lambda z. z$ is not uniquely determined and this case cannot occur.

Case \mathcal{E} is let $x = [\lambda z. z]$ in $\langle \square \rangle 1, \langle x \rangle ()$. The derivation must end with:

$$\frac{\mathcal{E}[x \triangleright [\tau_{\text{id}}]] \quad \emptyset \vdash \mathcal{E}[\langle x : \tau_{\text{id}} \rangle] : \tau}{\emptyset \vdash \mathcal{E}[\langle x \rangle] : \tau} \text{ PROJ-X}$$

However, $\mathcal{E}[x \triangleright \tau_{\text{id}}]$ does not hold (the proof is similar to the previous case).

Case \mathcal{E} is let $x = [\lambda z. z]$ in $(\langle x \rangle 1, \langle \square \rangle ())$. This is symmetric to the previous case, which cannot hold either.

Example A.5. Let e be let $f = \lambda x. x.1$ in $f(1, 2)$. e is well-typed using *backpropagation*. e is of the form $\mathcal{E}[x]$ where \mathcal{E} is the context let $f = \lambda x. \square$ in $f(1, 2)$. We have $\emptyset \vdash \mathcal{E}[x.1/2] : \text{int}$. Let us show that $\mathcal{E}[x \triangleright v\gamma_1, \gamma_2. \gamma_1 * \gamma_2]$. Assume $\emptyset \vdash \mathcal{E}[\{(x : g)\}] : \tau$. As g is a ground type, the type g of x is not a variable. Then, g cannot be that of an arbitrary sized tuple, since there is no such type for a tuple of arbitrary size. Hence, g must be a tuple $\prod_{i=1}^n \bar{\tau}$ for some size n . Since the codomain of f must be a tuple of size 2 (for $f(1, 2)$ to be well-typed), then n must also be 2. This shows that $\mathcal{E}[x \triangleright v\gamma_1, \gamma_2. \gamma_1 * \gamma_2]$.

A.3 Constraint generation

We now present the formal translation from terms e to constraints C , such that the resulting constraint is satisfiable if and only if the term is well typed. The translation is defined as a function $\llbracket e : \alpha \rrbracket$, where e is the term to be translated and α is the expected type of e .

Pattern constraints. Thus far, our formal presentation of constraint patterns has remained abstract, deliberately leaving the syntax and semantics of patterns unspecified to accommodate a range of language features. We now concertize this by specifying the patterns used in OmniML (in Figure 12), and introducing the corresponding constraints for the variables they bind. Patterns include: (1) Tuple patterns $\Pi \alpha_j$, matching a tuple type $\prod_{i=1}^n \bar{\tau}$ of arity $n \geq j$, and binding the j -th component to α . (2) Nominal patterns $t _$, binding the name of a nominal type $t \bar{\tau}$ to the nominal variable t . (3) Polytype patterns $[s]$ matching a polytype $[\sigma]$ and binding the resulting scheme to the variable s .

1520 Each new constraint has an unsubstituted form ($s \leq \tau, x \leq s$ etc.), whose semantics is defined
 1521 via substitution into a sugared form ($\sigma \leq \tau, x \leq \sigma$, etc.). Semantic environments ϕ are extended to
 1522 interpret nominal variables t as names t and scheme variables s as ground type schemes s , that is
 1523 type schemes with no unbound variables (i.e., $\forall \text{fv}(\tau). \tau$).
 1524

	Patterns
	Constraints
1525 $\rho ::= \Pi \alpha_j \mid t _ \mid [s]$	
1526 $C ::= \dots \mid \Omega(\ell/t) \leq \tau_1 \rightarrow \tau_2 \mid \Omega(\ell/t) \leq \tau_1 \rightarrow \tau_2$	
1527 $s \leq \tau \mid \sigma \leq \tau$	
1528 $x \leq s \mid x \leq \sigma$	
1529	
1530 $\Pi \alpha_j \text{ matches } (\nu \bar{\gamma}. \Pi_{i=1}^n \bar{\gamma}) \bar{\beta} \triangleq [\alpha := \beta_j] \quad \text{if } n \geq j$	
1531	
1532 $t _ \text{ matches } (\nu \bar{\gamma}. t) \bar{\beta} \triangleq [t := t]$	
1533 $[s] \text{ matches } (\nu \bar{\gamma}. [\sigma]) \bar{\beta} \triangleq [s := \sigma[\bar{\gamma} := \bar{\beta}]]$	
1534	
1535 LAB-INST	SCM-INST
1536 $\frac{\phi \vdash \Omega(\ell/\phi(t)) \leq \tau_1 \rightarrow \tau_2}{\phi \vdash \Omega(\ell/t) \leq \tau_1 \rightarrow \tau_2}$	$\frac{\phi \vdash \phi(s) \leq \tau}{\phi \vdash s \leq \tau}$
1537	
1538	
1539 $\Omega(\ell/t) \leq \tau_1 \rightarrow \tau_2 \triangleq \exists \bar{\alpha}. \tau_1 = \tau \wedge \tau_2 = t \bar{\alpha} \quad \text{if } \Omega(\ell/t) = \forall \bar{\alpha}. \tau \rightarrow t \bar{\alpha}$	
1540	
1541 $(\forall \bar{\alpha}. \tau') \leq \tau \triangleq \exists \bar{\alpha}. \tau' = \tau$	
1542 $x \leq (\forall \bar{\alpha}. \tau) \triangleq \forall \bar{\alpha}. x \tau$	
1543	

Fig. 12. Patterns for OmniML.

1544 *Constraint generation.* The function $\llbracket - : = \rrbracket$ is defined in Figure 13. All generated type variables
 1545 are fresh with respect to the expected type α , ensuring capture-avoidance. Unsurprisingly, variables
 1546 generate an instantiation constraint. Unit () requires the type α to be 1. A function generates a
 1547 constraint that binds two fresh flexible type variables for the argument and return types. We use
 1548 a let constraint to bind the argument in the constraint generated for the body of the function.
 1549 The let constraint is monomorphic since β' is fully constrained by type variables defined outside
 1550 the abstraction’s scope and therefore cannot be generalized. Applications introduce two fresh
 1551 flexible, one for the argument type and one for the type of the function, typing each subterm with
 1552 these, ensuring α is the expected return type. Let-bindings generates a polymorphic let constraint;
 1553 $\lambda \alpha. \llbracket e : \alpha \rrbracket$ is a principal constraint abstraction for e : its intended interpretation is the set of all
 1554 types that e admits.
 1555

1556 Annotations bind their flexible variables and enforce the equality of the annotated type τ and the
 1557 expected type α . Tuples introduce fresh variables for each component and unify their product with
 1558 α . Explicit projections ensure e has a tuple type $\Pi_{i=1}^n \bar{\beta}$ and extract the j -th component β_j , unifying
 1559 it with α . Implicit projections defer this via a suspended match constraint, until the shape of e ’s
 1560 expected type is known to be a tuple, extracting the j -th component with the pattern $\Pi \beta_j$.

1561 For polytypes, boxing asserts that e has the polymorphic type σ (using universal quantification)
 1562 and that the expected type is the polytype $[\sigma]$. Unboxing suspends until the inferred type of e is
 1563 known to be a polytype, captured by the pattern $[s]$, at which point we require α to be an instance
 1564 of s . Explicit unboxing is analogous, but uses an explicit scheme σ and therefore does not require a
 1565 suspended match constraint. Implicit boxing infers the principal type for e using a let constraint
 1566

1569	$\llbracket x : \alpha \rrbracket$	$\triangleq x \alpha$
1570	$\llbracket () : \alpha \rrbracket$	$\triangleq \alpha = 1$
1571	$\llbracket \lambda x. e : \alpha \rrbracket$	$\triangleq \exists \beta. \gamma. \alpha = \beta \rightarrow \gamma \wedge \text{let } x = \lambda \beta'. \beta' = \beta \text{ in } \llbracket e : \gamma \rrbracket$
1572	$\llbracket e_1 e_2 : \alpha \rrbracket$	$\triangleq \exists \beta \gamma. \gamma = \beta \rightarrow \alpha \wedge \llbracket e_1 : \gamma \rrbracket \wedge \llbracket e_2 : \beta \rrbracket$
1573	$\llbracket \text{let } x = e_1 \text{ in } e_2 : \alpha \rrbracket$	$\triangleq \text{let } x = \lambda \beta. \llbracket e_1 : \beta \rrbracket \text{ in } \llbracket e_2 : \alpha \rrbracket$
1574	$\llbracket (e : \exists \bar{\alpha}. \tau) : \alpha \rrbracket$	$\triangleq \exists \bar{\alpha}. \alpha = \tau \wedge \llbracket e : \alpha \rrbracket$
1575	$\llbracket (e_1, \dots, e_n) : \alpha \rrbracket$	$\triangleq \exists \bar{\alpha}. \alpha = \prod_{i=1}^n \bar{\alpha} \wedge \bigwedge_{i=1}^n \llbracket e_i : \alpha_i \rrbracket$
1576	$\llbracket [e.j/n] : \alpha \rrbracket$	$\triangleq \exists \bar{\beta}, \bar{\beta}. \llbracket e : \beta \rrbracket \wedge \beta = \prod_{i=1}^n \bar{\beta} \wedge \alpha = \beta_j$
1577	$\llbracket [e.j] : \alpha \rrbracket$	$\triangleq \exists \beta. \llbracket e : \beta \rrbracket \wedge \text{match } \beta \text{ with } \Pi \gamma_j \rightarrow \alpha = \gamma$
1578	$\llbracket [e : \exists \bar{\alpha}. \sigma] : \alpha \rrbracket$	$\triangleq \exists \bar{\alpha}. \llbracket e : \sigma \rrbracket \wedge \alpha = [\sigma]$
1579	$\llbracket \langle e : \exists \bar{\alpha}. \sigma \rangle : \alpha \rrbracket$	$\triangleq \exists \bar{\alpha}, \beta. \llbracket e : \beta \rrbracket \wedge \beta = [\sigma] \wedge \sigma \leq \alpha$
1580	$\llbracket \langle e \rangle : \alpha \rrbracket$	$\triangleq \exists \alpha. \llbracket e : \alpha \rrbracket \wedge \text{match } \alpha \text{ with } [s] \rightarrow s \leq \alpha$
1581	$\llbracket [e] : \alpha \rrbracket$	$\triangleq \text{let } x = \lambda \beta. \llbracket e : \beta \rrbracket \text{ in match } \alpha \text{ with } [s] \rightarrow x \leq s$
1582	$\llbracket [e.l] : \alpha \rrbracket$	$\triangleq \exists \beta. \llbracket e : \beta \rrbracket \wedge \llbracket l ! \beta \rrbracket \wedge \llbracket l : \alpha \rightarrow \beta \rrbracket$
1583	$\llbracket \overline{\{l = e\}} : \alpha \rrbracket$	$\triangleq \llbracket \bar{l} ! \alpha \rrbracket \wedge \bigwedge_{i=1}^n \llbracket l_i = e_i : \alpha \rrbracket$
1584	$\llbracket \{e\} : \alpha \rrbracket$	$\triangleq \exists \beta. \llbracket e : \beta \rrbracket$
1585		
1586		
1587		
1588		
1589	$\llbracket e : \tau \rrbracket$	$\triangleq \exists \alpha. \alpha = \tau \wedge \llbracket e : \alpha \rrbracket$
1590	$\llbracket e : \forall \bar{\alpha}. \tau \rrbracket$	$\triangleq \forall \bar{\alpha}. \llbracket e : \tau \rrbracket$
1591		
1592		
1593	$\llbracket l = e : \alpha \rrbracket$	$\triangleq \exists \beta. \llbracket e : \beta \rrbracket \wedge \llbracket l : \beta \rightarrow \alpha \rrbracket$
1594		
1595	$\llbracket \ell : \alpha \rightarrow \beta \rrbracket$	$\triangleq \text{match } \tau_2 \text{ with } t _ \rightarrow \Omega(\ell/t) \leq \alpha \rightarrow \beta$
1596	$\llbracket \ell/t : \alpha \rightarrow \beta \rrbracket$	$\triangleq \Omega(\ell/t) \leq \alpha \rightarrow \beta$
1597	$\llbracket \{\ell\} : \alpha \rightarrow \beta \rrbracket$	$\triangleq \text{true}$
1598	$\llbracket (l : \exists \bar{\alpha}. \tau) : \alpha \rightarrow \beta \rrbracket$	$\triangleq \exists \bar{\alpha}. \beta = \tau \wedge \llbracket l : \alpha \rightarrow \beta \rrbracket$
1599	$\llbracket \bar{l} ! \alpha \rrbracket$	$\triangleq \begin{cases} \exists \bar{\alpha}. \alpha = t \bar{\alpha} & \text{if } \bar{l} ! t \in \Omega \\ \text{true} & \text{otherwise} \end{cases}$
1600		
1601		
1602		
1603		
1604	Fig. 13. The constraint generation translation for OmniML.	
1605		
1606		

and suspends until the expected type of the entire term is known to be a polytype, bound to s . We then assert that the principal type of e is at least as general as s , via the constraint $x \leq s$.

Record projections generate a fresh variable for the nominal record type, constraining e to this type, and use the auxiliary function $\llbracket l : \alpha \rightarrow \beta \rrbracket$ to instantiate the label. The function $\llbracket \bar{l} ! \alpha \rrbracket$ checks whether a label sequence \bar{l} uniquely determines a record type, unifying α with $t \bar{\alpha}$ if so, or leaving it unconstrained if ambiguous. This function enables closed-world reasoning for both projections and constructions, and corresponds to the judgment $\bar{l} ! \tau$ defined in §A.1.

Record construction checks label uniqueness and generates a per-field constraint $l_i = e_i$, introducing a fresh variable β for each field's type and ensuring that e has this type and the label l instantiates to $\beta \rightarrow \alpha$. Label instantiation constraints $\llbracket \ell : \alpha \rightarrow \beta \rrbracket$ suspend until β is known to be a

1618 record type; once resolved, the label type is looked up in Ω and instantiated. Explicit instantiations
 1619 bypass suspension and directly instantiate the label's type.

1620 B Unification

1622 The unification rules are listed in Figure 14. Rewriting proceeds under an arbitrary context \mathcal{U} ,
 1623 modulo α -equivalence and associativity/commutativity of conjunctions.

1624 Our algorithm is largely standard [Pottier and Rémy 2005] but replaces type constructors with
 1625 canonical *principal shapes*, enabling a uniform treatment of monotypes and polytypes within
 1626 unification compared to prior formulations [Garrigue and Rémy 1999].

$$\begin{array}{c}
 \text{U-EXISTS} \quad \text{U-CYCLE} \quad \text{U-TRUE} \quad \text{U-FALSE} \\
 \frac{(\exists \alpha. U_1) \wedge U_2 \quad \alpha \# U_2}{\exists \alpha. U_1 \wedge U_2} \quad \frac{U \quad \text{cyclic } (U)}{\text{false}} \quad \frac{U \wedge \text{true}}{U} \quad \frac{\mathcal{U}[\text{false}] \quad \mathcal{U} \neq \square}{\text{false}} \\
 \text{U-MERGE} \quad \text{U-STUTTER} \quad \text{U-NAME} \quad \text{U-DECOMP} \\
 \frac{\alpha = \epsilon_1 \wedge \alpha = \epsilon_2}{\alpha = \epsilon_1 = \epsilon_2} \quad \frac{\alpha = \alpha = \epsilon}{\alpha = \epsilon} \quad \frac{\zeta(\bar{\tau}, \tau_i, \bar{\tau}') = \epsilon \quad \alpha \# \bar{\tau}, \bar{\tau}', \epsilon \quad \tau_i \notin \mathcal{V}}{\exists \alpha. \alpha = \tau_i \wedge \zeta(\bar{\tau}, \alpha, \bar{\tau}') = \epsilon} \quad \frac{\zeta \bar{\alpha} = \zeta' \bar{\beta} = \epsilon}{\zeta \bar{\alpha} = \epsilon \wedge \bar{\alpha} = \bar{\beta}} \\
 \text{U-CLASH} \quad \text{U-TRIVIAL} \\
 \frac{\zeta \bar{\alpha} = \zeta' \bar{\beta} = \epsilon \quad \zeta \neq \zeta'}{\text{false}} \quad \frac{\epsilon \quad |\epsilon| \leq 1}{\text{true}}
 \end{array}$$

1640 Fig. 14. Unification algorithm as a series of rewriting rules $U_1 \longrightarrow U_2$. All shapes are principal.
 1641

1642 We briefly summarize the role of each rule. **U-EXISTS** lifts existential quantifiers, enabling ap-
 1643 plications of **U-MERGE** and **U-CYCLE** since all multi-equations eventually become part of a single
 1644 conjunction. **U-MERGE** combines multi-equations sharing a common variable and **U-STUTTER** removes
 1645 duplicate variables. **U-DECOMP** decomposes equal types with matching shapes into equalities be-
 1646 tween their subcomponents, while **U-CLASH** detects shape mismatches that result in failure. **U-NAME**
 1647 introduces fresh variable for subcomponents, ensuring unification operates on *shallow terms*, mak-
 1648 ing sharing of type variables explicit and avoiding copying types in rules such as **U-DECOMP**. **U-TRUE**
 1649 and **U-TRIVIAL** eliminate trivial constraints, and **U-FALSE** propagates failure. Finally, **U-CYCLE** imple-
 1650 ments the *occurs check*, ensuring that a type variable does not occur in the type it is being unified
 1651 with. This is a necessary condition for unification, as it would otherwise lead to infinite types⁸.
 1652 This is formalized by the relation $\alpha <_{\mathcal{U}} \beta$ indicating that α occurs in a type assigned to β in \mathcal{U} . A
 1653 unification problem is cyclic, written *cyclic* (U), if $\alpha <_{\mathcal{U}}^* \alpha$ for some α .

1655
 1656
 1657
 1658
 1659
 1660
 1661
 1662
 1663
 1664
 1665 ⁸We discuss relaxing this constraint in §8.
 1666

1667 **C Full technical reference**

1668 This section repeats all the technical definitions mentioned in the paper, including the cases, rules,
 1669 and definitions that were omitted from the main paper to save space. It can serve as a useful
 1670 cheatsheet to understand a definition in full, or when studying the meta-theory of the system.
 1671

1672 $\alpha, \beta, \gamma \in \mathcal{V}$	Type variables
1673 $\tau ::= \alpha \mid 1 \mid \tau_1 \rightarrow \tau_2 \mid \prod_{i=1}^n \tau_i \mid t \bar{\tau} \mid [\sigma]$	Types
1674 $\sigma ::= \tau \mid \forall \alpha. \sigma$	Type schemes
1675 g	Ground types
1676 s	Ground type schemes
1677 $r ::= \alpha[\phi]$	Ground region
1678 $\mathfrak{G} \subseteq \mathcal{R}$	Sets of ground types
1679 $\mathfrak{R} \subseteq \mathcal{G}$	Sets of ground regions
1680 $C ::= \text{true} \mid \text{false} \mid C_1 \wedge C_2 \mid \exists \alpha. C \mid \forall \alpha. C \mid \tau_1 = \tau_2$	Constraints
1681 $\mid \text{let } x = \lambda \alpha. C_1 \text{ in } C_2 \mid x \tau$	
1682 $\mid \text{match } \tau \text{ with } \bar{\chi}$	
1683 $\mid \bar{\epsilon} \mid \text{let } x \alpha [\bar{\alpha}] = C_1 \text{ in } C_2 \mid \exists i^x. C \mid i[\alpha \rightsquigarrow \tau]$	
1684 $\mid \Omega(\ell/t) \leq \tau_1 \rightarrow \tau_2 \mid \Omega(\ell/t) \leq \tau_1 \rightarrow \tau_2$	
1685 $\mid s \leq \tau \mid \sigma \leq \tau \mid x \leq s \mid x \leq \sigma$	
1686 $\chi ::= \rho \rightarrow C$	Branches
1687 $\rho ::= _ \mid \Pi \alpha_j \mid t _ \mid [s]$	Patterns
1688 $\phi ::= \emptyset \mid \phi[\alpha := g] \mid \phi[x := \mathfrak{G}] \mid \phi[x := \mathfrak{R}] \mid \phi[i := \phi']$	Semantic environment
1689 $\mid \phi[t := t] \mid \phi[s := s]$	
1690 $U ::= \text{true} \mid \text{false} \mid U_1 \wedge U_2 \mid \exists \alpha. U \mid \epsilon$	Unification problems
1691 $\epsilon ::= \emptyset \mid \tau = \epsilon$	Multi-equations
1692 $\mathcal{C} ::= \square \mid \mathcal{C} \wedge C \mid C \wedge \mathcal{C} \mid \exists \alpha. \mathcal{C} \mid \forall \alpha. \mathcal{C}$	Constraint contexts
1693 $\mid \text{let } x = \lambda \alpha. \mathcal{C} \text{ in } C \mid \text{let } x = \lambda \alpha. C \text{ in } \mathcal{C}$	
1694 $\mid \text{let } x \alpha [\bar{\alpha}] = \mathcal{C} \text{ in } C \mid \text{let } x \alpha [\bar{\alpha}] = C \text{ in } \mathcal{C} \mid \exists i^x. \mathcal{C}$	
1695 $\zeta ::= v\bar{y}. \tau$	Shapes
1696 ξ	Canonical principal shapes
1697 $e ::= x \mid () \mid \lambda x. e \mid e_1 e_2 \mid \text{let } x = e_1 \text{ in } e_2 \mid (e : \exists \bar{\alpha}. \tau)$	Terms
1698 $\mid \{l = e\} \mid e.l$	
1699 $\mid (e_1, \dots, e_n) \mid e.j \mid e.j/n$	
1700 $\mid [e] \mid [e : \exists \bar{\alpha}. \sigma] \mid \langle e \rangle \mid \langle e : \exists \bar{\alpha}. \sigma \rangle$	
1701 $\mid \{e\}$	
1702 $l ::= \ell \mid \ell/t \mid \{\ell\} \mid (l : \exists \bar{\alpha}. \tau)$	Labels
1703 $\mathcal{E} ::= \square \mid \mathcal{E} e \mid e \mathcal{E} \mid \text{let } x = \mathcal{E} \text{ in } e \mid \text{let } x = e \text{ in } \mathcal{E} \mid (\mathcal{E} : \exists \bar{\alpha}. \tau)$	Term contexts
1704 $\mid \{l_1 = e_1 \dots l_i = \mathcal{E} \dots l_n = e_n\} \mid \mathcal{E}.l$	
1705 $\mid (e_1, \dots, \mathcal{E}, \dots, e_n) \mid \mathcal{E}.j \mid \mathcal{E}.j/n$	
1706 $\mid [\mathcal{E}] \mid [\mathcal{E} : \exists \bar{\alpha}. \sigma] \mid \langle \mathcal{E} \rangle \mid \langle \mathcal{E} : \exists \bar{\alpha}. \sigma \rangle$	
1707 $\mid \{\mathcal{E}\}$	
1708 $\mathcal{L} ::= \mathcal{E}[e. \square] \mid \mathcal{E}[\{l_1 = e_1; \dots; \square = e_i; \dots; l_n = e_n\}]$	Label contexts
1709 $\Gamma ::= \emptyset \mid \Gamma, x : \sigma$	Typing contexts
1710 $\Omega ::= \emptyset \mid \Omega, \ell : \forall \bar{\alpha}. \tau \rightarrow t \bar{\alpha}$	Label environment
1711 \mid	
1712 \mid	
1713 \mid	
1714 \mid	
1715 \mid	

1716	$\boxed{\phi \vdash C}$	Under the environment ϕ , the constraint C is satisfiable.
1717		
1718	TRUE	$\text{CONJ} \quad \frac{\phi \vdash C_1 \quad \phi \vdash C_2}{\phi \vdash C_1 \wedge C_2}$
1719		$\text{EXISTS} \quad \frac{\phi[\alpha := \mathbf{g}] \vdash C}{\phi \vdash \exists \alpha. C}$
1720		$\text{FORALL} \quad \frac{\forall \mathbf{g}, \phi[\alpha := \mathbf{g}] \vdash C}{\phi \vdash \forall \alpha. C}$
1721		$\text{UNIF} \quad \frac{\phi(\tau_1) = \phi(\tau_2)}{\phi \vdash \tau_1 = \tau_2}$
1722		
1723	LET	$\frac{\phi \vdash \exists \alpha. C_1 \quad \phi[x := \phi(\lambda \alpha. C_1)] \vdash C_2}{\phi \vdash \text{let } x = \lambda \alpha. C_1 \text{ in } C_2}$
1724		$\text{APP} \quad \frac{\phi(\tau) \in \phi(x)}{\phi \vdash x \tau}$
1725		
1726	SUSP-CTX	
1727	$\mathcal{C}[\tau ! \varsigma]$	$\phi \vdash \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]$
1728		$\phi \vdash \mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]$
1729		$\text{MULTI-UNIF} \quad \frac{\forall \tau \in \epsilon, \phi(\tau) = \mathbf{g}}{\phi \vdash \epsilon}$
1730		
1731	LETR	$\frac{\phi \vdash \exists \alpha, \bar{\alpha}. C_1 \quad \phi[x := \phi(\lambda \alpha[\bar{\alpha}]. C_1)] \vdash C_2}{\phi \vdash \text{let } x \alpha [\bar{\alpha}] = C_1 \text{ in } C_2}$
1732		$\text{APP}\bar{R} \quad \frac{\alpha[\phi'] \in \phi(x) \quad \phi(\tau) = \phi'(\alpha)}{\phi \vdash x \tau}$
1733		
1734	EXISTS-INST	
1735	$\alpha[\phi'] \in \phi(x)$	$\phi[i := \phi'] \vdash C$
1736		$\text{PARTIAL-INST} \quad \frac{\phi(i)(\alpha) = \phi(\tau)}{\phi \vdash i[\alpha \rightsquigarrow \tau]}$
1737		$\text{LAB-INST} \quad \frac{\phi \vdash \Omega(\ell/\phi(t)) \leq \tau_1 \rightarrow \tau_2}{\phi \vdash \Omega(\ell/t) \leq \tau_1 \rightarrow \tau_2}$
1738		
1739	SCM-INST	$\frac{\phi \vdash \phi(s) \leq \tau}{\phi \vdash s \leq \tau}$
1740		$\text{ABS-INST} \quad \frac{\phi \vdash x \leq \phi(s)}{\phi \vdash x \leq s}$
1741		
1742	match $\tau := \varsigma$ with $\rho \rightarrow \bar{C}$	$\triangleq \quad \exists \bar{\alpha}. \tau = \varsigma \bar{\alpha} \wedge \theta(C_i) \quad \text{if } \rho_i \text{ matches } \varsigma \bar{\alpha} = \theta$
1743	$\Omega(\ell/t) \leq \tau_1 \rightarrow \tau_2$	$\triangleq \quad \exists \bar{\alpha}. \tau_1 = \tau \wedge \tau_2 = t \bar{\alpha} \quad \text{if } \Omega(\ell/t) = \forall \bar{\alpha}. \tau \rightarrow t \bar{\alpha}$
1744		
1745	$(\forall \bar{\alpha}. \tau') \leq \tau$	$\triangleq \quad \exists \bar{\alpha}. \tau' = \tau$
1746	$x \leq (\forall \bar{\alpha}. \tau)$	$\triangleq \quad \forall \bar{\alpha}. x \tau$
1747		
1748	$\phi(\lambda \alpha[\bar{\alpha}]. C)$	$\triangleq \quad \{\alpha[\phi[\alpha := \mathbf{g}, \bar{\alpha} := \bar{\mathbf{g}}]] \in \mathcal{R} : \phi[\alpha := \mathbf{g}, \bar{\alpha} := \bar{\mathbf{g}}] \vdash C\}$
1749		
1750	$\phi(\lambda \alpha. C)$	$\triangleq \quad \{\mathbf{g} \in \mathcal{G} : \phi[\alpha := \mathbf{g}] \vdash C\}$
1751		
1752	$\mathcal{C}[\tau ! \varsigma]$	$\triangleq \quad \forall \phi, \mathbf{g}. \phi \vdash \lfloor \mathcal{C}[\tau = \mathbf{g}] \rfloor \implies \text{shape}(\mathbf{g}) = \varsigma$
1753	Note:	in most definitions, we ignore the additional OmniML constraints, as they are not particularly interesting.
1754		
1755		
1756	$\boxed{\zeta \preceq \zeta'}$	The shape ζ' is an instance of ζ . Alternatively, ζ' is more general than ζ .
1757		
1758		
1759	INST-SHAPE	
1760		$\frac{\bar{\gamma}_2 \# v\bar{\gamma}_1. \tau}{v\bar{\gamma}_1. \tau \leq v\bar{\gamma}_2. \tau[\bar{\gamma}_1 := \bar{\tau}_1]}$
1761		
1762		
1763	Definition C.1.	A non-trivial shape $\zeta \in \mathcal{S}$ is the principal shape of the type τ iff:
1764		

1765 (1) $\exists \bar{\tau}', \tau = \zeta \bar{\tau}'$

1766 (2) $\forall \zeta' \in \mathcal{S}, \forall \bar{\tau}', \tau = \zeta' \bar{\tau}' \implies \zeta \leq \zeta'$

1767 A principal shape $v\bar{y}.\tau$ is *canonical* if the sequence of its free variables \bar{y} appear in the order in
1768 which the variables occur in τ . shape (τ) is the canonical principal shape of τ .

1769

1770

1771 $\boxed{\rho \text{ matches } \varsigma \bar{\alpha} = \theta}$ The pattern ρ matches the shape ς with components $\bar{\alpha}$ binding
1772 pattern variables in θ .

1773

1774 $\Pi \beta_j \text{ matches } (v\bar{y}.\Pi_{i=1}^n \bar{y}) \bar{\alpha} \triangleq [\beta := \alpha_j] \quad \text{if } n \geq j$

1775

1776 $t_ \text{ matches } (v\bar{y}.t) \bar{\alpha} \triangleq [t := t]$

1777 $[s] \text{ matches } (v\bar{y}.[\sigma]) \bar{\alpha} \triangleq [s := \sigma[\bar{y} := \bar{\alpha}]]$

1778

1779

1780 $\boxed{C \text{ simple}}$ The constraint C is simple.

1781

SIMPLE-TRUE	SIMPLE-FALSE	$\frac{\text{SIMPLE-CONJ}}{C_1 \text{ simple} \quad C_2 \text{ simple}} \quad C_1 \wedge C_2 \text{ simple}$	SIMPLE-EXISTS	SIMPLE-FORALL
$\boxed{\text{true simple}}$	$\boxed{\text{false simple}}$		$\boxed{C \text{ simple}}$	$\boxed{C \text{ simple}}$

1782

SIMPLE-UNIF	$\frac{\text{SIMPLE-LET}}{C_1 \text{ simple} \quad C_2 \text{ simple}}$	$\frac{\text{SIMPLE-APP}}{x \tau \text{ simple}}$	SIMPLE-LET _R	SIMPLE-FORALL
$\boxed{\tau_1 = \tau_2 \text{ simple}}$	$\boxed{\text{let } x = \lambda \alpha. C_1 \text{ in } C_2 \text{ simple}}$		$\boxed{C_1 \text{ simple}}$	$\boxed{C_2 \text{ simple}}$

1783

	$\frac{\text{SIMPLE-EXISTS-INST}}{C \text{ simple}}$	SIMPLE-PARTIAL-INST		
	$\boxed{\exists i^x. C \text{ simple}}$	$\boxed{i[\alpha \rightsquigarrow \tau] \text{ simple}}$		

1784

1785

SIMPLE-CTX-HOLE	$\frac{\text{SIMPLE-CTX-CONJ-LEFT}}{C \text{ simple} \quad C \text{ simple}}$	SIMPLE-CTX-CONJ-RIGHT	
$\boxed{\square \text{ simple}}$	$\boxed{C \wedge C \text{ simple}}$	$\boxed{C \text{ simple} \quad C \text{ simple}}$	$\boxed{C \wedge C \text{ simple}}$

1786

SIMPLE-CTX-EXISTS	$\frac{\text{SIMPLE-CTX-FORALL}}{C \text{ simple}}$	SIMPLE-CTX-LET-ABS	
$\boxed{C \text{ simple}}$		$\boxed{C \text{ simple} \quad C \text{ simple}}$	$\boxed{\text{let } x = \lambda \alpha. C \text{ in } C \text{ simple}}$

1787

	$\frac{\text{SIMPLE-CTX-LET-IN}}{C \text{ simple} \quad C \text{ simple}}$	SIMPLE-CTX-EXISTS-INST	
	$\boxed{\text{let } x = \lambda \alpha. C \text{ in } C \text{ simple}}$	$\boxed{C \text{ simple}}$	$\boxed{\exists i^x. C \text{ simple}}$

1788

1789

1790

1791

1792

1793

1794

1795

1796

1797

1798

1799

1800

1801

1802

1803

1804

1805

1806

1807

1808

1809

1810

1811

1812

1813

$\boxed{[C]}$ The erasure of C .

1814	$\lfloor \text{true} \rfloor \triangleq \text{true}$
1815	$\lfloor \text{false} \rfloor \triangleq \text{false}$
1816	$\lfloor C_1 \wedge C_2 \rfloor \triangleq \lfloor C_1 \rfloor \wedge \lfloor C_2 \rfloor$
1817	$\lfloor \exists \alpha. C \rfloor \triangleq \exists \alpha. \lfloor C \rfloor$
1818	$\lfloor \forall \alpha. C \rfloor \triangleq \forall \alpha. \lfloor C \rfloor$
1819	$\lfloor \tau_1 = \tau_2 \rfloor \triangleq \tau_1 = \tau_2$
1820	$\lfloor \text{let } x = \lambda \alpha. C_1 \text{ in } C_2 \rfloor \triangleq \text{let } x = \lambda \alpha. \lfloor C_1 \rfloor \text{ in } \lfloor C_2 \rfloor$
1821	$\lfloor x \tau \rfloor \triangleq x \tau$
1822	$\lfloor \text{match } \tau \text{ with } \bar{\rho} \rightarrow \bar{C} \rfloor \triangleq \text{true}$
1823	$\lfloor \text{let } x \alpha [\bar{\alpha}] = C_1 \text{ in } C_2 \rfloor \triangleq \text{let } x \alpha [\bar{\alpha}] = \lfloor C_1 \rfloor \text{ in } \lfloor C_2 \rfloor$
1824	$\lfloor \exists i^x. C \rfloor \triangleq \exists i^x. \lfloor C \rfloor$
1825	$\lfloor i[\alpha \rightsquigarrow \tau] \rfloor \triangleq i[\alpha \rightsquigarrow \tau]$

1826

1827

1828

$\boxed{\phi \Vdash C}$ Under the semantic environment ϕ , the constraint C is canonically satisfiable.

1829

1830

1831

1832

1833

1834

1835

1836

1837

1838

$$\frac{\text{CAN-SIMPLE}}{\phi \Vdash_{\text{simple}} C}$$

$$\frac{\text{CAN-SUSP-CTX}}{\mathcal{C}[\tau ! \varsigma] \quad \phi \Vdash \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]} \quad \phi \Vdash \mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]$$

1839

1840

1841

1842

$$\frac{\text{LAB-MAGIC}}{\Gamma \vdash \{\ell\} : \tau' \rightarrow \tau}$$

$$\frac{\text{LAB-ANNOT}}{\Gamma \vdash (l : \exists \bar{\alpha}. \tau) : \tau' \rightarrow \tau[\bar{\alpha} := \bar{\tau}]}$$

$$\frac{\text{LAB-X}}{\Omega(\ell/t) = \forall \bar{\alpha}. \tau \rightarrow t \bar{\alpha}} \quad \frac{}{\Gamma \vdash \ell/t : \bar{\tau}[\bar{\alpha} := \bar{\tau}] \rightarrow t \bar{\tau}}$$

1843

1844

$\boxed{l/t \in \Omega}$ The label l belongs to the nominal record type t in Ω .

1845

1846

$\boxed{\bar{l}/t \in \Omega}$ The set of labels \bar{l} belong to a unique nominal record type t in Ω .

1847

1848

$\boxed{\bar{l}/\tau}$ The set of labels \bar{l} infer the possibly unique type τ .

1849

1850	LAB-INI	$\ell : \forall \bar{\alpha}. \tau \rightarrow t \bar{\alpha} \in \Omega$	LAB-INX	$\ell/t \in \Omega$	LAB-INMAGIC	$\ell/t \in \Omega$	LAB-INANNOT	$l/t \in \Omega$	$(l : \exists \bar{\alpha}. \tau)/t \in \Omega$
1851		$\ell/t \in \Omega$		$(\ell/t)/t \in \Omega$		$\{\ell\}/t \in \Omega$			
1852									
1853									
1854	LAB-U	$\bar{l}/t \in \Omega \quad \forall t', \bar{l}/t' \in \Omega \implies t = t'$	LAB-!	$\bar{l}/t \in \Omega$	LAB-?	$\bar{l}/t \bar{\tau}$			
1855									
1856									
1857									

1858

1859

1860

1861

$\boxed{\Gamma \vdash l = e : \tau}$ Under the typing context Γ , the record assignment $l = e$ has the record type τ .

1862

$\boxed{\Gamma \vdash e : \sigma}$ Under the typing context Γ , the term e is assigned the type σ .

1863	VAR	$\frac{x : \sigma \in \Gamma}{\Gamma \vdash x : \sigma}$	1864	FUN	$\frac{\Gamma, x : \tau_1 \vdash e : \tau_2}{\Gamma \vdash \lambda x. e : \tau_1 \rightarrow \tau_2}$	1865	APP	$\frac{\Gamma \vdash e_1 : \tau_1 \rightarrow \tau_2 \quad \Gamma \vdash e_2 : \tau_1}{\Gamma \vdash e_1 e_2 : \tau_2}$	1866	UNIT	$\frac{}{\Gamma \vdash () : 1}$
1867	ANNOT	$\frac{\Gamma \vdash e : \tau[\bar{\alpha} := \bar{\tau}]}{\Gamma \vdash (e : \exists \bar{\alpha}. \tau) : \tau[\bar{\alpha} := \bar{\tau}]}$	1868	GEN	$\frac{\Gamma \vdash e : \sigma \quad \alpha \# \Gamma}{\Gamma \vdash e : \forall \alpha. \sigma}$	1869	INST	$\frac{\Gamma \vdash e : \forall \alpha. \sigma}{\Gamma \vdash e : \alpha[\alpha := \tau]}$	1870		
1871	LET	$\frac{\Gamma \vdash e_1 : \sigma \quad \Gamma, x : \sigma \vdash e_2 : \tau}{\Gamma \vdash \text{let } x = e_1 \text{ in } e_2 : \tau}$	1872	TUPLE	$\frac{(\Gamma \vdash e_i : \tau_i)_{i=1}^n}{\Gamma \vdash (e_1, \dots, e_n) : \prod_{i=1}^n \tau_i}$	1873	PROJ-X	$\frac{\Gamma \vdash e : \prod_{i=1}^n \tau_i \quad 1 \leq j \leq n}{\Gamma \vdash e.j/n : \tau_j}$	1874		
1875	PROJ-I	$\frac{\mathcal{E}[e \triangleright v\bar{y}. \prod_{i=1}^n \bar{y}_i] \quad \Gamma \vdash \mathcal{E}[e.j/n] : \tau}{\Gamma \vdash \mathcal{E}[e.j] : \tau}$	1876	POLY-X	$\frac{\Gamma \vdash e : \sigma[\bar{\alpha} := \bar{\tau}]}{\Gamma \vdash [e : \exists \bar{\alpha}. \sigma] : [\sigma[\bar{\alpha} := \bar{\tau}]]}$	1877			1878		
1879	POLY-I	$\frac{\mathcal{E}[e \triangleleft v\bar{y}. [\sigma]] \quad \Gamma \vdash \mathcal{E}[[e : \exists \bar{y}. \sigma]] : \tau}{\Gamma \vdash \mathcal{E}[[e]] : \tau}$	1880	USE-X	$\frac{\Gamma \vdash e : [\sigma][\bar{\alpha} := \bar{\tau}]}{\Gamma \vdash \langle e : \exists \bar{\alpha}. \sigma \rangle : \sigma[\bar{\alpha} := \bar{\tau}]}$	1881			1882		
1883	USE-I	$\frac{\mathcal{E}[e \triangleright v\bar{y}. [\sigma]] \quad \Gamma \vdash \mathcal{E}[\langle e : \exists \bar{y}. \sigma \rangle] : \tau}{\Gamma \vdash \mathcal{E}[\langle e \rangle] : \tau}$	1884	RCD-ASSN	$\frac{\Gamma \vdash e : \tau \quad l : \tau \rightarrow \tau'}{\Gamma \vdash l = e : \tau'}$	1885			1886		
1887	RCD	$\frac{(\Gamma \vdash l_i = e_i : \tau)_{i=1}^n \quad \bar{l} ! \tau}{\Gamma \vdash \{l_1 = e_1; \dots; l_n = e_n\} : \tau}$	1888	RCD-PROJ	$\frac{\Gamma \vdash e : \tau' \quad l : \tau \rightarrow \tau' \quad \bar{l} ! \tau}{\Gamma \vdash e.l : \tau}$	1889	LAB-I	$\frac{\mathcal{L}[\ell ! t] \quad \Gamma \vdash \mathcal{L}[\ell / t] : \tau}{\Gamma \vdash \mathcal{L}[\ell] : \tau}$	1890		
1891			1892	MAGIC	$\frac{\Gamma \vdash e : \tau}{\Gamma \vdash \{e\} : \tau'}$	1893			1894		
1895			1896	$E[e \triangleright \varsigma] \triangleq \forall \Gamma, \tau, g, \Gamma \vdash [\mathcal{E}[\{(e : g)\}]] : \tau \implies \text{shape}(g) = \varsigma$		1897			1898		
1899			1900	$E[e \triangleleft \varsigma] \triangleq \forall \Gamma, \tau, g, \Gamma \vdash [\mathcal{E}[\{(e : g)\}]] : \tau \implies \text{shape}(g) = \varsigma$		1901			1902	$\boxed{[\Gamma \vdash e : \tau]}$ $[\Gamma \vdash e : \tau]$ is satisfiable iff e has the expected <i>known</i> type τ under <i>known</i> context Γ .	
1903			1904	$\boxed{[\bar{l} ! \alpha]}$ $[\bar{l} ! \alpha]$ is satisfiable iff \bar{l} has the possibly unique type α .		1905			1906	$\boxed{[l : \alpha \rightarrow \beta]}$ $[l : \alpha \rightarrow \beta]$ is satisfiable iff l has the record type β and field type α .	
1907			1908	$\boxed{[l = e : \alpha]}$ $[l = e : \alpha]$ is satisfiable iff the record assignment $l = e$ has the record type α .		1909			1910	$\boxed{[e : \sigma]}$ $[e : \sigma]$ is satisfiable iff e has the expected <i>known</i> type scheme σ .	
1911			1912	$\boxed{[e : \alpha]}$ $[e : \alpha]$ is satisfiable iff e has the expected type α .							

1912	$\llbracket x : \alpha \rrbracket$	$\triangleq x \alpha$
1913	$\llbracket () : \alpha \rrbracket$	$\triangleq \alpha = 1$
1914	$\llbracket \lambda x. e : \alpha \rrbracket$	$\triangleq \exists \beta, \gamma. \alpha = \beta \rightarrow \gamma \wedge \text{let } x = \lambda \beta'. \beta' = \beta \text{ in } \llbracket e : \gamma \rrbracket$
1915	$\llbracket e_1 e_2 : \alpha \rrbracket$	$\triangleq \exists \beta \gamma. \gamma = \beta \rightarrow \alpha \wedge \llbracket e_1 : \gamma \rrbracket \wedge \llbracket e_2 : \beta \rrbracket$
1916	$\llbracket \text{let } x = e_1 \text{ in } e_2 : \alpha \rrbracket$	$\triangleq \text{let } x = \lambda \beta. \llbracket e_1 : \beta \rrbracket \text{ in } \llbracket e_2 : \alpha \rrbracket$
1917	$\llbracket (e : \exists \bar{\alpha}. \tau) : \alpha \rrbracket$	$\triangleq \exists \bar{\alpha}. \alpha = \tau \wedge \llbracket e : \alpha \rrbracket$
1918	$\llbracket (e_1, \dots, e_n) : \alpha \rrbracket$	$\triangleq \exists \bar{\alpha}. \alpha = \prod_{i=1}^n \bar{\alpha} \wedge \bigwedge_{i=1}^n \llbracket e_i : \alpha_i \rrbracket$
1919	$\llbracket e.j/n : \alpha \rrbracket$	$\triangleq \exists \beta, \bar{\beta}. \llbracket e : \beta \rrbracket \wedge \beta = \prod_{i=1}^n \bar{\beta} \wedge \alpha = \beta_j$
1920	$\llbracket e.j : \alpha \rrbracket$	$\triangleq \exists \beta. \llbracket e : \beta \rrbracket \wedge \text{match } \beta \text{ with } \prod \gamma_j \rightarrow \alpha = \gamma$
1921	$\llbracket [e : \exists \bar{\alpha}. \sigma] : \alpha \rrbracket$	$\triangleq \exists \bar{\alpha}. \llbracket e : \sigma \rrbracket \wedge \alpha = [\sigma]$
1922	$\llbracket \langle e : \exists \bar{\alpha}. \sigma \rangle : \alpha \rrbracket$	$\triangleq \exists \bar{\alpha}, \beta. \llbracket e : \beta \rrbracket \wedge \beta = [\sigma] \wedge \sigma \leq \alpha$
1923	$\llbracket \langle e : \alpha \rangle : \alpha \rrbracket$	$\triangleq \exists \alpha. \llbracket e : \alpha \rrbracket \wedge \text{match } \alpha \text{ with } [s] \rightarrow s \leq \alpha$
1924	$\llbracket [e] : \alpha \rrbracket$	$\triangleq \text{let } x = \lambda \beta. \llbracket e : \beta \rrbracket \text{ in match } \alpha \text{ with } [s] \rightarrow x \leq s$
1925	$\llbracket [e.l] : \alpha \rrbracket$	$\triangleq \exists \beta. \llbracket e : \beta \rrbracket \wedge \llbracket l ! \beta \rrbracket \wedge \llbracket l : \alpha \rightarrow \beta \rrbracket$
1926	$\llbracket \{\overline{l = e}\} : \alpha \rrbracket$	$\triangleq \llbracket \bar{l} ! \alpha \rrbracket \wedge \bigwedge_{i=1}^n \llbracket l_i = e_i : \alpha \rrbracket$
1927	$\llbracket \{e\} : \alpha \rrbracket$	$\triangleq \exists \beta. \llbracket e : \beta \rrbracket$
1928		
1929		
1930		
1931		
1932	$\llbracket e : \tau \rrbracket$	$\triangleq \exists \alpha. \alpha = \tau \wedge \llbracket e : \alpha \rrbracket$
1933	$\llbracket e : \forall \bar{\alpha}. \tau \rrbracket$	$\triangleq \forall \bar{\alpha}. \llbracket e : \tau \rrbracket$
1934		
1935		
1936	$\llbracket l = e : \alpha \rrbracket$	$\triangleq \exists \beta. \llbracket e : \beta \rrbracket \wedge \llbracket l : \beta \rightarrow \alpha \rrbracket$
1937		
1938	$\llbracket \ell : \alpha \rightarrow \beta \rrbracket$	$\triangleq \text{match } \tau_2 \text{ with } t _ \rightarrow \Omega(\ell/t) \leq \alpha \rightarrow \beta$
1939	$\llbracket \ell/t : \alpha \rightarrow \beta \rrbracket$	$\triangleq \Omega(\ell/t) \leq \alpha \rightarrow \beta$
1940	$\llbracket \{\ell\} : \alpha \rightarrow \beta \rrbracket$	$\triangleq \text{true}$
1941	$\llbracket (l : \exists \bar{\alpha}. \tau) : \alpha \rightarrow \beta \rrbracket$	$\triangleq \exists \bar{\alpha}. \beta = \tau \wedge \llbracket l : \alpha \rightarrow \beta \rrbracket$
1942	$\llbracket \bar{l} ! \alpha \rrbracket$	$\triangleq \begin{cases} \exists \bar{\alpha}. \alpha = t \bar{\alpha} & \text{if } \bar{l} ! t \in \Omega \\ \text{true} & \text{otherwise} \end{cases}$
1943		
1944		
1945		
1946		
1947	$\llbracket \emptyset \vdash e : \tau \rrbracket$	$\triangleq \llbracket e : \tau \rrbracket$
1948	$\llbracket x : \sigma, \Gamma \vdash e : \tau \rrbracket$	$\triangleq \text{let } x = \lambda \alpha. \sigma \leq \alpha \text{ in } \llbracket \Gamma \vdash e : \tau \rrbracket$
1949		
1950		
1951		
1952		
1953		
1954		
1955		
1956		
1957		
1958	$e \text{ simple}$	The term e is simple.
1959		
1960		

1961	SIMPLE-VAR	SIMPLE-FUN	SIMPLE-APP	SIMPLE-UNIT
1962	$\frac{}{x \text{ simple}}$	$\frac{e \text{ simple}}{\lambda x. e \text{ simple}}$	$\frac{e_1 \text{ simple} \quad e_2 \text{ simple}}{e_1 e_2 \text{ simple}}$	$\frac{}{() \text{ simple}}$
1963				
1964				
1965	SIMPLE-LET	SIMPLE-ANNOT	SIMPLE-TUPLE	SIMPLE-PROJX
1966	$e_1 \text{ simple} \quad e_2 \text{ simple}$	$\frac{}{e \text{ simple}}$	$\frac{(e_i \text{ simple})_{i=1}^n}{(e_1, \dots, e_n) \text{ simple}}$	$\frac{}{e \text{ simple}}$
1967	$\text{let } x = e_1 \text{ in } e_2 \text{ simple}$	$(e : \exists \bar{\alpha}. \tau) \text{ simple}$		$e.j/n \text{ simple}$
1968				
1969	SIMPLE-POLYX	SIMPLE-USEX	SIMPLE-RCD	SIMPLE-RCD-ASSN
1970	$e \text{ simple}$	$\frac{}{e \text{ simple}}$	$\frac{(l_i = e_i \text{ simple})_{i=1}^n}{\{l_1 = e_1 \dots l_n = e_n\}}$	$\frac{l \text{ simple} \quad e \text{ simple}}{l = e \text{ simple}}$
1971	$[e : \exists \bar{\alpha}. \sigma] \text{ simple}$	$\langle e : \exists \bar{\alpha}. \sigma \rangle \text{ simple}$		
1972				
1973	SIMPLE-RCD-PROJ	SIMPLE-MAGIC	SIMPLE-LAB	SIMPLE-LAB-ANNOT
1974	$e \text{ simple} \quad l \text{ simple}$	$e \text{ simple}$	$\frac{}{\ell/t \text{ simple}}$	$\frac{l \text{ simple}}{(l : \exists \bar{\alpha}. \tau) \text{ simple}}$
1975	$e.l \text{ simple}$	$\{e\} \text{ simple}$	$\frac{}{\{\ell\} \text{ simple}}$	
1976				
1977	$\boxed{[l]}$	The erasure of l .		
1978	$\boxed{[e]}$	The erasure of e .		
1979				
1980				
1981		$[x] \triangleq x$		
1982		$[\lambda x. e] \triangleq \lambda x. [e]$		
1983		$[e_1 e_2] \triangleq [e_1] [e_2]$		
1984		$[() \triangleq ()]$		
1985		$[\text{let } x = e_1 \text{ in } e_2] \triangleq \text{let } x = [e_1] \text{ in } [e_2]$		
1986		$[(e : \exists \bar{\alpha}. \tau)] \triangleq ([e] : \exists \bar{\alpha}. \tau)$		
1987		$[(e_1, \dots, e_n)] \triangleq ([e_1], \dots, [e_n])$		
1988		$[e.j] \triangleq \{ e \}$		
1989		$[e.j/n] \triangleq [e].j/n$		
1990		$[[e : \exists \bar{\alpha}. \sigma]] \triangleq [[e] : \exists \bar{\alpha}. \sigma]$		
1991		$[[e]] \triangleq \{ e \}$		
1992		$[\langle e \rangle] \triangleq \{ e \}$		
1993		$[\langle e : \exists \bar{\alpha}. \sigma \rangle] \triangleq \langle [e] : \exists \bar{\alpha}. \sigma \rangle$		
1994		$[\{l_1 = e_1 \dots l_n = e_n\}] \triangleq \{ l_1 = e_1 \dots l_n = e_n \}$		
1995		$[e.l] \triangleq [e].[l]$		
1996		$[\ell/t] \triangleq \ell/t$		
1997		$[\ell] \triangleq \{\ell\}$		
1998		$[(l : \exists \bar{\alpha}. \tau)] \triangleq ([l] : \exists \bar{\alpha}. \tau)$		
1999		$[\{\ell\}] \triangleq \{\ell\}$		
2000				
2001				
2002	$\boxed{\Gamma \vdash_{\text{simple}}^{\text{sd}} e : \tau}$	Under the typing context Γ , the simple term e has the type τ .		
2003				
2004				
2005	VAR-SD	LET-SD		
2006	$x : \forall \bar{\alpha}. \tau \in \Gamma$	$\Gamma \vdash_{\text{simple}}^{\text{sd}} e_1 : \tau_1 \quad \bar{\alpha} \# \Gamma \quad \Gamma, x : \forall \bar{\alpha}. \tau_1 \vdash_{\text{simple}}^{\text{sd}} e_2 : \tau_2$		
2007				
2008	$\Gamma \vdash_{\text{simple}}^{\text{sd}} x : \tau [\bar{\alpha} := \bar{\tau}]$		$\Gamma \vdash_{\text{simple}}^{\text{sd}} \text{let } x = e_1 \text{ in } e_2 : \tau_2$	
2009				

2010	$\boxed{\vdash e : \tau}$	The term e canonically has the type τ .
2011		
2012	CAN-BASE	
2013	$\emptyset \vdash_{\text{simple}}^{\text{sd}} e : \tau$	CAN-PROJ-I
2014		$\mathcal{E}[e \triangleright v\bar{y}. \Pi_{i=1}^n \bar{y}] \quad \vdash \mathcal{E}[e.j/n] : \tau$
2015	$\vdash e : \tau$	$\vdash \mathcal{E}[e.j] : \tau$
2016		
2017	CAN-USE-I	
2018	$\mathcal{E}[e \triangleright v\bar{y}. [\sigma]] \quad \vdash \mathcal{E}[\langle e : \exists \bar{y}. \sigma \rangle] : \tau$	
2019		$\vdash \mathcal{E}[\langle e \rangle] : \tau$
2020		
2021		
2022	$\boxed{U \longrightarrow U'}$	The unifier rewrites U to U' .
2023		
2024		
2025	U-EXISTS	
2026	$(\exists \alpha. U_1) \wedge U_2 \quad \alpha \# U_2$	U-CYCLE
2027	$\longrightarrow \exists \alpha. U_1 \wedge U_2$	$U \quad \text{cyclic } (U)$
2028		false
2029	U-MERGE	
2030	$\alpha = \epsilon_1 \wedge \alpha = \epsilon_2$	U-STUTTER
2031	$\alpha = \epsilon_1 = \epsilon_2$	$\alpha = \epsilon$
2032		
2033		
2034	U-CLASH	
2035	$\zeta \bar{\alpha} = \zeta' \bar{\beta} = \epsilon \quad \zeta \neq \zeta'$	
2036		false
2037		
2038		
2039	$\boxed{C \longrightarrow C'}$	The constraint solver rewrites C to C' .
2040		
2041	S-UNIF	
2042	$U_1 \quad U_1 \longrightarrow U_2$	S-TRUE
2043	$\longrightarrow U_2$	$C \wedge \text{true}$
2044		
2045		
2046	S-EXISTS-CONJ	
2047	$(\exists \alpha. C_1) \wedge C_2 \quad \alpha \# C_2$	S-LET-EXISTSLEFT
2048	$\longrightarrow \exists \alpha. C_1 \wedge C_2$	$\text{let } x \alpha [\bar{\alpha}] = \exists \beta. C_1 \text{ in } C_2 \quad \beta \# \alpha, \bar{\alpha}, C_2$
2049		$\text{let } x \alpha [\bar{\alpha}, \beta] = C_1 \text{ in } C_2$
2050		
2051	S-LET-EXISTSRIGHT	
2052	$\text{let } x \alpha [\bar{\alpha}] = C_1 \text{ in } \exists \beta. C_2 \quad \beta \# \alpha, \bar{\alpha}, C_1$	S-LET-CONJLEFT
2053	$\longrightarrow \exists \beta. \text{let } x = \lambda \bar{\alpha}. C_1 \text{ in } C_2$	$\text{let } x \alpha [\bar{\alpha}] = C_1 \wedge C_2 \text{ in } C_3 \quad C_1 \# \alpha, \bar{\alpha}$
2054		$\longrightarrow C_1 \wedge \text{let } x \alpha [\bar{\alpha}] = C_2 \text{ in } C_3$
2055	S-LET-CONJRIGHT	
2056	$\text{let } x \alpha [\bar{\alpha}] = C_1 \text{ in } (C_2 \wedge C_3) \quad x \# C_3$	S-MATCH-TYPE
2057	$\longrightarrow C_3 \wedge \text{let } x \alpha = C_1 \text{ in } C_2$	$\text{match } \tau \text{ with } \bar{\chi} \quad \tau \notin \mathcal{V}$
2058		$\longrightarrow \text{match } \tau := \text{shape } (\tau) \text{ with } \bar{\chi}$

<p>2059 S-MATCH-VAR</p> $\frac{\mathcal{C}[\text{match } \alpha \text{ with } \bar{\chi}] \quad \alpha = \tau = \epsilon \in \mathcal{C}}{\mathcal{C}[\text{match } \alpha := \text{shape } (\tau) \text{ with } \bar{\chi}]}$	<p>2060 S-INST-NAME</p> $\frac{i[\alpha \rightsquigarrow \tau] \quad \tau \notin \mathcal{V}}{\exists \gamma. \gamma = \tau \wedge i[\alpha \rightsquigarrow \gamma]}$
<p>2061 S-LET-APPR</p> $\frac{\text{let } x \alpha [\bar{\alpha}] = C \text{ in } \mathcal{C}[x \tau] \quad \gamma \# \tau \quad x \# \text{bv}(\mathcal{C})}{\text{let } x \alpha [\bar{\alpha}] = C \text{ in } \mathcal{C}[\exists \gamma. i^x. \gamma = \tau \wedge i[\alpha \rightsquigarrow \gamma]]}$	
<p>2062 S-INST-COPY</p> $\frac{\begin{array}{l} \text{let } x \alpha [\bar{\alpha}] = C \text{ in } \mathcal{C}[i^x[\alpha' \rightsquigarrow \gamma]] \quad C = C' \wedge \alpha' = \varsigma \bar{\beta} = \epsilon \quad \alpha' \in \alpha, \bar{\alpha} \\ \neg \text{cyclic } (C) \quad \bar{\beta}' \# \alpha', \gamma, \bar{\beta} \quad x \# \text{bv}(\mathcal{C}) \end{array}}{\text{let } x \alpha [\bar{\alpha}] = C \text{ in } \mathcal{C}[\exists \bar{\beta}'. \gamma = \varsigma \bar{\beta}' \wedge i^x[\bar{\beta} \rightsquigarrow \bar{\beta}']]}$	
<p>2063 S-INST-UNIFY</p> $\frac{i[\alpha \rightsquigarrow \gamma_1] \wedge i[\alpha \rightsquigarrow \gamma_2]}{i[\alpha \rightsquigarrow \gamma_1] \wedge \gamma_1 = \gamma_2}$	<p>2064 S-INST-POLY</p> $\frac{\begin{array}{l} \text{let } x \alpha [\bar{\alpha}] = \bar{\epsilon} \wedge C \text{ in } \mathcal{C}[i^x[\alpha' \rightsquigarrow \gamma]] \quad \forall \alpha'. \bar{\epsilon} \equiv \text{true} \\ \alpha' \in \alpha, \bar{\alpha} \quad \alpha' \# C \quad i.\alpha' \# \text{insts}(\mathcal{C}) \quad x \# \text{bv}(\mathcal{C}) \end{array}}{\text{let } x \alpha [\bar{\alpha}] = \bar{\epsilon} \wedge C \text{ in } \mathcal{C}[\text{true}]}$
<p>2065 S-INST-MONO</p> $\frac{\text{let } x \alpha [\bar{\alpha}] = C \text{ in } \mathcal{C}[i^x[\beta \rightsquigarrow \gamma]] \quad \beta \notin \alpha, \bar{\alpha} \quad x, \beta \# \text{bv}(\mathcal{C})}{\text{let } x \alpha [\bar{\alpha}] = C \text{ in } \mathcal{C}[\beta = \gamma]}$	
<p>2066 S-LET-SOLVE</p> $\frac{\text{let } x \alpha [\bar{\alpha}] = \bar{\epsilon} \text{ in } C \quad \exists \alpha, \bar{\alpha}. \bar{\epsilon} \equiv \text{true} \quad x \# C}{C}$	
<p>2067 S-COMPRESS</p> $\frac{\text{let } x \alpha [\bar{\alpha}, \beta] = C_1 \wedge \beta = \gamma = \epsilon \text{ in } C_2 \quad \beta \neq \gamma}{\text{let } x \alpha [\bar{\alpha}] = C_1[\beta := \gamma] \wedge \gamma = \epsilon[\beta := \gamma] \text{ in } C_2[x, \beta := \gamma]}$	
<p>2068 S-Gc</p> $\frac{\text{let } x \alpha [\bar{\alpha}, \beta] = C_1 \wedge \beta = \epsilon \text{ in } C_2 \quad \beta \# C_1, \epsilon, C_2}{\text{let } x \alpha [\bar{\alpha}] = C_1 \wedge \epsilon \text{ in } C_2}$	
<p>2069 S-EXISTS-LOWER</p> $\frac{\text{let } x \alpha [\bar{\alpha}, \bar{\beta}] = C_1 \text{ in } C_2 \quad \exists \alpha, \bar{\alpha}. C_1 \text{ determines } \bar{\beta}}{\exists \bar{\beta}. \text{let } x \alpha [\bar{\alpha}] = C_1 \text{ in } C_2}$	
<p>2070 S-BACKPROP</p> $\frac{\begin{array}{l} \mathcal{C}[\text{let } x \alpha [\bar{\alpha}] = \mathcal{C}_1[\text{match } \alpha' \text{ with } \bar{\chi}] \text{ in } \mathcal{C}_2[i^x[\alpha' \rightsquigarrow \gamma]]] \\ \alpha' \in \alpha, \bar{\alpha} \quad \gamma = \tau = \epsilon \in \mathcal{C}[\mathcal{C}_2] \quad x \# \text{bv}(\mathcal{C}_2) \end{array}}{\mathcal{C}[\text{let } x \alpha [\bar{\alpha}] = \mathcal{C}_1[\text{match } \alpha' := \text{shape } (\tau) \text{ with } \bar{\chi}] \text{ in } \mathcal{C}_2[i^x[\alpha' \rightsquigarrow \gamma]]]}$	<p>2071 S-EXISTS-EXISTS-INST</p> $\frac{\exists i^x. \exists \alpha. C}{\exists \alpha. \exists i^x. C}$
<p>2072 S-EXISTS-INST-CONJ</p> $\frac{\exists i^x. C_1 \wedge C_2 \quad i \# C_1}{C_1 \wedge \exists i^x. C_2}$	<p>2073 S-EXISTS-INST-LET</p> $\frac{\text{let } x \alpha [\bar{\alpha}] = C_1 \text{ in } \exists i^{x'}. C_2 \quad x \neq x'}{\exists i^{x'}. \text{let } x \alpha [\bar{\alpha}] = C_1 \text{ in } C_2}$
<p>2074 S-EXISTS-INST-SOLVE</p> $\frac{\exists i^x. C \quad i \# C}{C}$	

$$\begin{array}{c}
\text{S-ALL-CONJ} \\
\frac{\forall \bar{\alpha}. \exists \bar{\beta}. C_1 \wedge C_2 \quad \bar{\alpha}, \bar{\beta} \# C_1}{C_1 \wedge \forall \bar{\alpha}. \exists \bar{\beta}. C_2} \\
\text{S-EXISTS-ALL} \\
\frac{\forall \bar{\alpha}. \exists \bar{\beta}, \bar{\gamma}. C \quad \exists \bar{\alpha}, \bar{\beta}. C \text{ determines } \bar{\gamma}}{\exists \bar{\gamma}. \forall \bar{\alpha}. \exists \bar{\beta}. C} \\
\text{S-ALL-ESCAPE} \\
\frac{\forall \bar{\alpha}, \alpha. \exists \bar{\beta}. C \wedge \bar{\epsilon} \quad \alpha <_{\bar{\epsilon}}^* \gamma \quad \gamma \# \alpha, \bar{\beta} \quad \alpha \# \bar{\beta}}{\text{false}} \\
\text{S-ALL-RIGID} \\
\frac{\forall \bar{\alpha}, \alpha. \exists \bar{\beta}. C \wedge \alpha = \tau = \epsilon \quad \tau \notin \mathcal{V} \quad \alpha \# \bar{\beta}}{\text{false}} \\
\text{S-ALL-SOLVE} \\
\frac{\forall \bar{\alpha}. \exists \bar{\beta}. \bar{\epsilon} \quad \exists \bar{\beta}. \bar{\epsilon} \equiv \text{true}}{\text{true}}
\end{array}$$

Definition C.2. C determines $\bar{\beta}$ if and only if every ground assignments ϕ and ϕ' that satisfy (the erasure of) C and coincide outside of β coincide on $\bar{\beta}$ as well.

$$C \text{ determines } \beta \triangleq \forall \phi, \phi'. \phi \vdash [C] \wedge \phi' \vdash [C] \wedge \phi =_{\bar{\beta}} \phi' \implies \phi = \phi'$$

Definition C.3. A context \mathcal{C} proves a multi-equation ϵ , written $\epsilon \in \mathcal{C}$, if there exists a decomposition $\mathcal{C} = \mathcal{C}_1[\epsilon \wedge \mathcal{C}_2]$ such that $\text{fv}(\epsilon) \# \text{bv}(\mathcal{C}_2)$

Definition C.4 (Measure). For the relation $\phi \vdash C$, the following measure enables a useful induction principle:

$$\|C\| \triangleq \langle \#\text{match } C, |C| \rangle$$

where $\langle \dots \rangle$ denotes a pair with lexicographic ordering, and:

- (1) $\#\text{match } C$ is the number of match τ with $\bar{\chi}$ constraints in C .
- (2) the last component $|C|$ is a structural measure of constraints *i.e.*, a conjunction $C_1 \wedge C_2$ is larger than the two conjuncts C_1, C_2 .

2157 D Properties of the constraint language

2158 This appendix establishes key properties of the constraint language. The first is the principality of
 2159 shapes **Theorem D.1**: any non-variable type τ admits a non-trivial principal shape ζ .

2160 The second is the canonicalization of satisfiability derivations $\phi \vdash C$, which enables a simple
 2161 induction principle for reasoning about unicity. This canonical form for derivations is a crucial tool
 2162 in our proof of soundness and completeness in §F.

2164 D.1 Principality of shapes

2165 **THEOREM D.1 (PRINCIPAL SHAPES).** *Any non-variable type τ has a non-trivial principal shape ζ .*

2166 **PROOF.** Let us assume τ is a non-variable type.

2167 **Case τ is a type constructor $c \bar{\tau}$.**

2168 c is a top-level type constructor of arity n , which in our setting may be the nullary 1, the binary
 2169 arrow, the n -ary product, or a n -ary nominal type. In all these cases, the shape of τ is $v\bar{y}. c \bar{y}$
 2170 where \bar{y} is a sequence of n distinct type variables. This is clearly principal.

2171 **Case τ is a polytype $[\forall \bar{\alpha}. \tau]$.**

2172 We may assume *w.l.o.g.* that each variable of $\bar{\alpha}$ occurs free in τ . Let $(\pi_i)_{i=1}^n$ be the sequence of
 2173 shortest paths in τ that cannot be extended to reach a (polymorphic) variable in $\bar{\alpha}$, in lexicographic
 2174 order and \bar{y} be a sequence $(\gamma_i)_{i=1}^n$ of distinct variables that do not appear in τ . Let τ_0 be $\tau[\pi_i := \gamma_i]_{i=1}^n$,
 2175 *i.e.*, the term τ where each path π_i has been substituted by the variable γ_i . Let ζ be the
 2176 shape $v\bar{y}. [\forall \bar{\alpha}. \tau_0]$. We claim that ζ is actually the principal shape of $[\forall \bar{\alpha}. \tau]$.

2177 By construction, τ is equal to $\zeta \bar{\tau}$ (1), where $\bar{\tau}$ is the sequence composed of τ_i equal to τ/π_i for
 2178 i ranging from 1 to n . Indeed, by definition, $\zeta \bar{\tau}$ is equal to $(\tau[\pi_i := \gamma_i]_{i=1}^n)[\gamma_i := \tau_i]$ which is
 2179 obviously equal to τ . The remaining of the proof checks that ζ is minimal (2), that is, we assume
 2180 that ζ' is another shape such that $[\forall \bar{\alpha}. \tau]$ is equal to $\zeta' \bar{\tau}'$ for some $\bar{\tau}'$ (3) and show that $\zeta \leq \zeta'$ (4).

2181 It follows from (3) that ζ' must be a polytype shape, *i.e.*, of the form $v\bar{y}'. [\forall \bar{\beta}. \tau']$ and $[\forall \bar{\alpha}. \tau]$
 2182 is equal to $[\forall \bar{\beta}. \tau'][\bar{y}' := \bar{\tau}']$ (5). We may assume *w.l.o.g.* that $\bar{\beta}$ and \bar{y}' are disjoint, that \bar{y}' does
 2183 not contain useless variables, *i.e.*, that they all appear in τ' and that they actually appear in
 2184 lexicographic order. Now that never term contains useless variables, (5) implies that the sequences
 2185 $\bar{\alpha}$ and $\bar{\beta}$ can be put in one-to-one correspondences. Besides, since they all ordered in the order
 2186 of appearance in terms, they the correspondence respects the ordering. Hence, the substitution
 2187 $[\bar{\beta} := \bar{\alpha}]$ is a renaming. Therefore, we can assume *w.l.o.g.* that $\bar{\beta}$ is $\bar{\alpha}$. That is, (5) becomes that
 2188 $[\forall \bar{\alpha}. \tau]$ is equal to $[\forall \bar{\alpha}. \tau'[\bar{y}' := \bar{\tau}']]$, which given that variables $\bar{\alpha}$ appear in the same order in
 2189 both terms, implies that τ is equal to $\tau'[\bar{y}' := \bar{\tau}']$ (6).

2190 Since $\bar{\tau}'$ does not contain any variable in $\bar{\alpha}$, every path π_i is a path in τ' . Thus, we may write
 2191 τ' as $\tau'[\pi_i := \tau''_i]_{i=1}^n$ where τ''_i is τ'/π_i . This is also equal to $(\tau'[\pi_i := \gamma_i]_{i=1}^n)[\gamma_i := \tau''_i]_{i=1}^n$,
 2192 that is $\tau_0[\gamma_i := \tau''_i]_{i=1}^n$. In summary, we have τ' is equal to $\tau_0[\gamma_i := \tau''_i]_{i=1}^n$, which implies that
 2193 $[\forall \bar{\alpha}. \tau']$ is equal to $[\forall \bar{\alpha}. \tau_0[\gamma_i := \tau''_i]_{i=1}^n]$, *i.e.*, $[\forall \bar{\alpha}. \tau_0][\gamma_i := \tau''_i]_{i=1}^n$ (7). By **INST-SHAPE**, we have
 2194 $v\bar{y}. [\forall \bar{\alpha}. \tau_0] \leq v\bar{y}'. [\forall \bar{\alpha}. \tau_0][\gamma_i := \tau''_i]_{i=1}^n$, which, given (7), is exactly (4).

2195 \square

2196 D.2 Canonicalization of satisfiability

2197 They key result in this section is that our semantic derivations $\phi \vdash C$ can always be rewritten to
 2198 only apply the rule **SUSP-CTX** at the very bottom of the derivation, rather than in the middle of
 2199 derivations. This corresponds to explicitating the unique shapes of all suspended constraints (in
 2200 some order that respects the dependency between suspended constraints), and then continuing
 2201 with a syntax-directed proof of a fully-discharged constraint.

2206 We did not impose this ordering in our definition of the semantics to make it more flexible and
 2207 more declarative, but the inversion principle that it provides will be helpful when reasoning about
 2208 the solver in §E.

2209 We define in §C a formal judgment C simple that says that C does not contain any suspended
 2210 match constraint, and extend it trivially to constraint contexts: \mathcal{C} simple. In particular, the erasure
 2211 $[C]$ of a constraint (Definition 3.4) is always simple. We then introduce in §C a “canonical” semantic
 2212 judgment $\phi \Vdash C$ that enforces the structure we mentioned: its derivation starts by discharging
 2213 suspended constraints, until eventually we reach a simple constraint C . Below we prove that any
 2214 semantic derivation $\phi \vdash C$ can be turned into a canonical semantic derivation $\phi \Vdash C$.

2215 We can think of this result as controlling the amount of non-syntax-directedness in our rules:
 2216 we need some of it, but it suffices to have it only at the outside, and it contains a more standard
 2217 derivation that is easy to reason about.

2218 *Inversion.* When C is simple, a derivation of $\phi \vdash C$ does not use the contextual rule (it is a
 2219 derivation in $\phi \vdash_{\text{simple}} C$), so it enjoys the usual inversion principle on syntax-directed judgments;
 2220 for example, if $\phi \vdash_{\text{simple}} C_1 \wedge C_2$ then by inversion $\phi \vdash_{\text{simple}} C_1$ and $\phi \vdash_{\text{simple}} C_2$, etc.

2221 *Congruence.* Congruence does not hold in general in our system due to the contextual rule.
 2222 For example, $C_1 \triangleq (\text{match } \alpha \text{ with } _\rightarrow \text{true})$ is unsatisfiable so we have $C_1 \equiv \text{false}$, but for
 2223 $\mathcal{C} \triangleq (\exists \alpha. \alpha = \text{int} \wedge \square)$ we have $\mathcal{C}[C_1] \equiv \text{true}$ and $\mathcal{C}[\text{false}] \equiv \text{false}$. It holds simply for simple
 2224 constraints.

2225 **LEMMA D.2 (SIMPLE CONGRUENCE).** *Given simple constraints C_1, C_2 and simple context \mathcal{C} . If
 2226 $C_1 \models C_2$, then $\mathcal{C}[C_1] \models \mathcal{C}[C_2]$.*

2227 **PROOF.** Induction on the derivation of \mathcal{C} simple. □

2228 *Composability.* The composability result below is an important test of our definition of the
 2229 unicity condition $\mathcal{C}[\tau ! \varsigma]$, which is in part engineered for this lemma to be simple to prove. In the
 2230 past we used a definition of unicity that also required $\mathcal{C}[\text{true}]$ to be satisfiable, which broke the
 2231 composability property.

2232 **LEMMA D.3 (COMPOSABILITY OF UNICITY).** *If $\mathcal{C}_1[\tau ! \varsigma]$, then $\mathcal{C}_2[\mathcal{C}_1][\tau ! \varsigma]$.*

2233 **PROOF.** Induction on the structure of \mathcal{C}_2 .

2234 **Case** \square immediate.

2235 **Case** $\mathcal{C}_3 \wedge C$.

$\mathcal{C}_1[\tau ! \varsigma]$	Premise
$\mathcal{C}_3[\mathcal{C}_1][\tau ! \varsigma]$	By i.h.
For all ϕ, g	Definition of $(\mathcal{C}_3[\mathcal{C}_1] \wedge C)[\tau ! \varsigma]$
$\phi \vdash [\mathcal{C}_3[\mathcal{C}_1][\tau = g]] \wedge [C]$	$\implies I$
$\phi \vdash [\mathcal{C}_3[\mathcal{C}_1][\tau = g]]$	Simple inversion
shape (g) = ς	$\implies E$ on $\mathcal{C}_3[\mathcal{C}_1][\tau ! \varsigma]$
■ $(\mathcal{C}_3[\mathcal{C}_1] \wedge C)[\tau ! \varsigma]$	Above

2236 **Case** $C \wedge \mathcal{C}_3$.

2237 Similar to the $\mathcal{C}_3 \wedge C$ case.

2238 **Case** $\exists \alpha. \mathcal{C}_3$.

2255	$\mathcal{C}_1[\tau!\varsigma]$	Premise
2256	$\mathcal{C}_3[\mathcal{C}_1][\tau!\varsigma]$	By <i>i.h.</i>
2257	For all ϕ, g	Definition of $(\exists\alpha.\mathcal{C}_3[\mathcal{C}_1])[\tau!\varsigma]$
2258	$\phi \vdash \exists\alpha.[\mathcal{C}_3[\mathcal{C}_1][\tau = g]]$	$\implies I$
2259	$\phi[\alpha := g'] \vdash [\mathcal{C}_3[\mathcal{C}_1][\tau = g]]$	Simple inversion
2260	shape(g) = ς	$\implies E$ on $\mathcal{C}_3[\mathcal{C}_1][\tau!\varsigma]$
2261	$\blacksquare (\exists\alpha.\mathcal{C}_3[\mathcal{C}_1])[\tau!\varsigma]$	Above
2262		
2263	Case $\forall\alpha.\mathcal{C}_3$.	
2264	Similar to $\exists\alpha.\mathcal{C}_3$ case.	
2265	Case $\exists t^x.\mathcal{C}_3$.	
2266	Similar to $\exists\alpha.\mathcal{C}_3$ case.	
2267	Case let $x = \lambda\alpha.\mathcal{C}_3$ in C .	
2268		
2269	$\mathcal{C}_1[\tau!\varsigma]$	Premise
2270	$\mathcal{C}_3[\mathcal{C}_1][\tau!\varsigma]$	By <i>i.h.</i>
2271	For all ϕ, g	Definition of $(\text{let } x \dots)[\tau!\varsigma]$
2272	$\phi \vdash \text{let } x = \lambda\alpha.[\mathcal{C}_3[\mathcal{C}_1][\tau = g]] \text{ in } [C]$	$\implies I$
2273	$\phi \vdash \exists\alpha.[\mathcal{C}_3[\mathcal{C}_1][\tau = g]]$	Simple inversion
2274	$\phi[\alpha := g'] \vdash [\mathcal{C}_3[\mathcal{C}_1][\tau = g]]$	Simple inversion
2275	shape(g) = ς	$\implies E$ on $\mathcal{C}_3[\mathcal{C}_1][\tau!\varsigma]$
2276	$\blacksquare (\text{let } x = \lambda\alpha.\mathcal{C}_3[\mathcal{C}_1] \text{ in } C)[\tau!\varsigma]$	Above
2277		
2278	Case let $x = \lambda\alpha.C$ in \mathcal{C}_3 .	
2279	Similar to let $x = \lambda\alpha.\mathcal{C}_3$ in C case.	
2280	Case let $x \alpha [\bar{\alpha}] = \mathcal{C}_3$ in C .	
2281	Similar to let $x = \lambda\alpha.\mathcal{C}_3$ in C case.	
2282	Case let $x \alpha [\bar{\alpha}] = C$ in \mathcal{C}_3 .	
2283	Similar to let $x = \lambda\alpha.C$ in \mathcal{C}_3 case.	
2284		
2285		\square
2286	LEMMA D.4 (INVERSION OF UNICITY).	
2287	(i) If $(\exists\alpha.\mathcal{C})[\tau!\varsigma]$, then $\mathcal{C}[\tau!\varsigma]$.	
2288	(ii) If $(\forall\alpha.\mathcal{C})[\tau!\varsigma]$, then $\mathcal{C}[\tau!\varsigma]$.	
2289		
2290	PROOF. The definition of $\mathcal{C}[\tau!\varsigma]$ uses simple semantics on the erasure $[\mathcal{C}]$, so these results are	
2291	easily shown by simple inversion.	\square
2292		
2293	LEMMA D.5 (DECANONICALIZATION). If $\phi \Vdash C$, then $\phi \vdash C$.	
2294	PROOF. Induction on the given derivation $\phi \Vdash C$	\square
2295		
2296	THEOREM D.6 (CANONICALIZATION). If $\phi \vdash C$, then $\phi \Vdash C$.	
2297	PROOF. We proceed by induction on $\phi \vdash C$ with the measure $\ C\ $.	
2298	Case	
2299	$\frac{}{\phi \vdash \text{true}}$ TRUE	
2300	$\phi \vdash \text{true}$	
2301	$\blacksquare \phi \Vdash \text{true}$ immediate by CAN-BASE	
2302		
2303		

2304 **Case**

$$\frac{\phi(\tau_1) = \phi(\tau_2)}{\phi \vdash \tau_1 = \tau_2} \text{UNIF}$$

2308 Similar to the **TRUE** case.2309 **Case**

$$\frac{\phi \vdash C_1 \quad \phi \vdash C_2}{\phi \vdash C_1 \wedge C_2} \text{CONJ}$$

2314 $\phi \vdash C_1$ Premise2315 $\phi \vdash C_2$ Premise2316 $\phi \Vdash C_1$ By *i.h.*2317 $\phi \Vdash C_2$ By *i.h.*2318 By cases on $\phi \Vdash C_1, \phi \Vdash C_2$.2319 **Subcase**

$$\frac{\phi \vdash C_1 \quad C_1 \text{ simple}}{\phi \Vdash C_1} \text{CAN-BASE}$$

$$\frac{\phi \vdash C_2 \quad C_2 \text{ simple}}{\phi \Vdash C_2} \text{CAN-BASE}$$

2327 $\blacksquare \phi \Vdash C_1 \wedge C_2$ immediate by **CAN-BASE**2328 **Subcase**

$$\frac{\mathcal{C}[\tau ! \varsigma] \quad \phi \Vdash \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]}{\phi \Vdash \underbrace{\mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]}_{C_1}} \text{CAN-SUSP-CTX}$$

2333 $\phi \Vdash C_2$ 2335 $\phi \Vdash \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]$ Premise2336 $\phi \Vdash \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]$ Lemma D.52337 $\phi \Vdash \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}] \wedge C_2$ By **CONJ**2338 $\phi \Vdash \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}] \wedge C_2$ By *i.h.*2339 $\mathcal{C}[\alpha ! \varsigma]$ Premise2340 $(\mathcal{C} \wedge C_2)[\alpha ! \varsigma]$ Lemma D.32342 $\blacksquare \phi \Vdash \mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]$ By **CAN-SUSP-CTX**2343 **Subcase**2344 $\phi \Vdash C_1$

$$\frac{\mathcal{C}[\tau ! \varsigma] \quad \phi \Vdash \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]}{\phi \Vdash \underbrace{\mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]}_{C_2}} \text{CAN-SUSP-CTX}$$

2350 Symmetric to the above case.

2351

2352

Case

$$\frac{\phi[\alpha := g] \vdash C}{\phi \vdash \exists \alpha. C} \text{ EXISTS}$$

$\phi[\alpha := g] \vdash C$ Premise
 $\phi[\alpha := g] \Vdash C$ By i.h.

By cases on $\phi[\alpha := g] \Vdash C$.

Subcase

$$\frac{\phi[\alpha := g] \vdash C \quad C \text{ simple}}{\phi[\alpha := g] \Vdash C} \text{ CAN-BASE}$$

$\blacksquare \phi \Vdash \exists \alpha. C$ Immediate by CAN-BASE

Subcase

$$\frac{\mathcal{C}[\tau ! \varsigma] \quad \phi[\alpha := g] \Vdash \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}] \quad \phi \Vdash \underbrace{\mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]}_C}_{\phi \Vdash \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]} \text{ CAN-SUSP-CTX}$$

$\phi[\alpha := g] \Vdash \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]$ Premise
 $\phi[\alpha := g] \vdash \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]$ Lemma D.5
 $\phi \vdash \exists \alpha. \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]$ By EXISTS
 $\phi \Vdash \exists \alpha. \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]$ By i.h.
 $\mathcal{C}[\tau ! \varsigma]$ Premise
 $(\exists \alpha. \mathcal{C})[\tau ! \varsigma]$ Lemma D.3
 $\blacksquare \phi \Vdash \exists \alpha. \mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]$ By CAN-SUSP-CTX

Case

$$\frac{\forall g, \phi[\alpha := g] \vdash C}{\phi \vdash \forall \alpha. C} \text{ FORALL}$$

Similar to the EXISTS case.

Case

$$\frac{\phi \vdash \exists \alpha. C_1 \quad \phi[x := \phi(\lambda \alpha. C_1)] \vdash C_2}{\phi \vdash \text{let } x = \lambda \alpha. C_1 \text{ in } C_2} \text{ LET}$$

$\phi \vdash \exists \alpha. C_1$ Premise

$\phi \Vdash \exists \alpha. C_1$ By i.h.

$\phi[x := \phi(\lambda \alpha. C_1)] \vdash C_2$ Premise

$\phi[x := \phi(\lambda \alpha. C_1)] \Vdash C_2$ By i.h.

By cases on $\phi \Vdash \exists \alpha. C_1, \phi[x := \phi(\lambda \alpha. C_1)] \Vdash C_2$.

Subcase

$$\frac{\phi \vdash \exists \alpha. C_1 \quad \exists \alpha. C_1 \text{ simple}}{\phi \Vdash \exists \alpha. C_1} \text{ CAN-BASE}$$

$$\frac{\phi[x := \phi(\lambda \alpha. C_1)] \vdash C_2 \quad C_2 \text{ simple}}{\phi[x := \phi(\lambda \alpha. C_1)] \Vdash C_2} \text{ CAN-BASE}$$

$\blacksquare \phi \Vdash \text{let } x = \lambda \alpha. C_1 \text{ in } C_2$ Immediate by CAN-BASE

Subcase

$$\frac{(\exists \alpha. C_1)[\tau ! \varsigma] \quad \phi \Vdash \exists \alpha. \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]}{\phi \Vdash \exists \alpha. \underbrace{\mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]}_{C_1}} \text{ CAN-SUSP-CTX}$$

$\phi[x := \phi(\lambda \alpha. C_1)] \Vdash C_2$

$(\exists \alpha. \mathcal{C})[\tau ! \varsigma]$ Premise

$\mathcal{C}[\tau ! \varsigma]$ Lemma D.4

$\phi \Vdash \exists \alpha. \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]$ Premise

$\phi \vdash \exists \alpha. \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]$ Lemma D.5

$\phi(\lambda \alpha. C_1) = \phi(\lambda \alpha. \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}])$ Corollary D.8

$\phi \vdash \text{let } x = \lambda \alpha. \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}] \text{ in } C_2$ By LET

$\phi \Vdash \text{let } x = \lambda \alpha. \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}] \text{ in } C_2$ By i.h.

$(\text{let } x = \lambda \alpha. \mathcal{C} \text{ in } C_2)[\tau ! \varsigma]$ Lemma D.3

$\blacksquare \phi \Vdash \text{let } x = \lambda \alpha. \mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}] \text{ in } C_2$ By CAN-SUSP-CTX

Subcase

$\phi \Vdash \exists \alpha. C_1$

$$\frac{\mathcal{C}[\tau ! \varsigma] \quad \phi[x := \phi(\lambda \alpha. C_1)] \Vdash \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]}{\phi[x := \phi(\lambda \alpha. C_1)] \Vdash \underbrace{\mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]}_{C_2}} \text{ CAN-SUSP-CTX}$$

$\mathcal{C}[\tau ! \varsigma]$ Premise

$(\text{let } x = \lambda \alpha. C_1 \text{ in } \mathcal{C})[\tau ! \varsigma]$ Lemma D.3

$\phi[x := \phi(\lambda \alpha. C_1)] \Vdash \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]$ Premise

$\phi[x := \phi(\lambda \alpha. C_1)] \vdash \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]$ Lemma D.5

$\phi \vdash \text{let } x = \lambda \alpha. C_1 \text{ in } \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]$ By LET

$\phi \Vdash \text{let } x = \lambda \alpha. C_1 \text{ in } \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]$ By i.h.

$\blacksquare \phi \Vdash \text{let } x = \lambda \alpha. C_1 \text{ in } \mathcal{C}[\text{match } \tau \text{ with } \varsigma]$ By CAN-SUSP-CTX

Case

$$\frac{\phi(\tau) \in \phi(x)}{\phi \vdash x \tau} \text{ APP}$$

Similar to the TRUE case.

2451 **Case**

$$\frac{\phi \vdash \exists \alpha, \bar{\alpha}. C_1 \quad \phi[x := \phi(\lambda \alpha[\bar{\alpha}]. C_1)] \vdash C_2}{\phi \vdash \text{let } x \alpha [\bar{\alpha}] = C_1 \text{ in } C_2} \text{ LETR}$$

2454

2455 Similar to the **LET** case.2456 **Case**

$$\frac{\alpha[\phi'] \in \phi(x) \quad \phi(\tau) = \phi'(\alpha)}{\phi \vdash x \tau} \text{ APPR}$$

2459

2460 Similar to the **APP** case.2461 **Case**

$$\frac{\alpha[\phi'] \in \phi(x) \quad \phi[i := \phi'] \vdash C}{\phi \vdash \exists i^x. C} \text{ EXISTS-INST}$$

2464

2465 Similar to the **EXISTS** case.2466 **Case**

$$\frac{\forall \tau \in \epsilon, \phi(\tau) = g}{\phi \vdash \epsilon} \text{ MULTI-UNIF}$$

2469

2470 Similar to the **UNIF** case.

2471

2472 **Case**

$$\frac{\phi(i)(\alpha) = \phi(\tau)}{\phi \vdash i[\alpha \rightsquigarrow \tau]} \text{ PARTIAL-INST}$$

2475

2476 Similar to the **APP** case.

2477

□

2478

2479 LEMMA D.7 (INVERSION OF SUSPENSION). If $\phi \vdash \mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]$ and $\mathcal{C}[\tau ! \varsigma]$, then

2480

 $\phi \vdash \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]$.

2481

PROOF. We use canonicalization ([Theorem D.6](#)) to induct on $\phi \Vdash \mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]$ instead of $\phi \vdash \mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]$.

2483

This simplifies the proof, but introduces a circular dependency between [Theorem D.6](#) and [Lemma D.7](#). However, this does not compromise the well-foundedness of induction, as the application of [Lemma D.7](#) (via [Corollary D.8](#)) within the proof of [Theorem D.6](#) is restricted to strictly smaller constraints.

2487

2488 **Case**

$$\frac{\phi \vdash \mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}] \quad \mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}] \text{ simple}}{\phi \Vdash \mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]} \text{ CAN-BASE}$$

2490

The second premise is a contradiction.

2491

2492 **Case**

$$\frac{\mathcal{C}'[\tau' ! \varsigma'] \quad \phi \Vdash \mathcal{C}'[\text{match } \tau' := \varsigma' \text{ with } \bar{\chi}']}{\phi \Vdash \underbrace{\mathcal{C}'[\text{match } \tau' \text{ with } \bar{\chi}']}_{\mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]}} \text{ CAN-SUSP-CTX}$$

2495

By cases on $\mathcal{C} = \mathcal{C}'$.

2496

2497

2498

2499

2500	Subcase $\mathcal{C} = \mathcal{C}'$.		
2501	$\mathcal{C} = \mathcal{C}'$	Premise	
2502	$\tau' = \tau$		
2503	$\varsigma' = \varsigma$		
2504	$\bar{\chi}' = \bar{\chi}$		
2505	$\blacksquare \phi \Vdash \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]$	Premise	
2506	Subcase $\mathcal{C} \neq \mathcal{C}'$.		
2507	$\mathcal{C}_2[\text{match } \tau \text{ with } \bar{\chi}, \text{match } \tau' \text{ with } \bar{\chi}'] = \mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]$	For some 2-hole context \mathcal{C}_2	
2508		$= \mathcal{C}'[\text{match } \tau' \text{ with } \bar{\chi}']$	
2509			
2510			
2511	$\phi \Vdash \mathcal{C}_2[\text{match } \tau \text{ with } \bar{\chi}, \text{match } \tau' := \varsigma' \text{ with } \bar{\chi}']$	Premise	
2512	For all ϕ', g'	Defn. of $\mathcal{C}_2[\square, \text{match } \tau' := \varsigma' \text{ with } \bar{\chi}'][\tau'! \varsigma]$	
2513			
2514			
2515	$\phi' \vdash [\mathcal{C}_2[\tau = g', \text{match } \tau' := \varsigma' \text{ with } \bar{\chi}']]$	$\implies I$	
2516	$\phi' \vdash [\mathcal{C}_2[\tau = g', \text{true}]]$	Lemma D.2	
2517	$[\mathcal{C}_2[\tau = g', \text{true}]] = [\mathcal{C}_2[\tau = g', [\text{match } \tau' \text{ with } \bar{\chi}']]]$	By definition	
2518	$= [\mathcal{C}[\tau = g']]$	By definition	
2519	$\phi' \vdash [\mathcal{C}[\tau = g']]$	Above	
2520	$\text{shape}(g') = \varsigma$	$\implies E \text{ on } \mathcal{C}[\tau'! \varsigma]$	
2521	$\mathcal{C}_2[\square, \text{match } \tau' := \varsigma' \text{ with } \bar{\chi}'][\tau'! \varsigma]$	Above	
2522	$\phi \Vdash \mathcal{C}_2[\text{match } \tau := \varsigma \text{ with } \bar{\chi}, \text{match } \tau' := \varsigma' \text{ with } \bar{\chi}']$	By i.h.	
2523			
2524	For all ϕ', g'	Defn. of $\mathcal{C}_2[\text{match } \tau := \varsigma \text{ with } \bar{\chi}, \square][\tau'! \varsigma']$	
2525			
2526			
2527	$\phi' \vdash [\mathcal{C}_2[\text{match } \tau := \varsigma \text{ with } \bar{\chi}, \tau' = g']]$	$\implies I$	
2528	$\phi' \vdash [\mathcal{C}_2[\text{true}, \tau' = g']]$	Lemma D.2	
2529	$[\mathcal{C}_2[\text{true}, \tau' = g']] = [\mathcal{C}_2[[\text{match } \tau \text{ with } \bar{\chi}], \tau' = g']]$	By definition	
2530	$= [\mathcal{C}'[\tau' = g']]$	By definition	
2531	$\phi' \vdash [\mathcal{C}[\tau = g']]$	Above	
2532	$\mathcal{C}'[\tau'! \varsigma']$	Premise	
2533	$\text{shape}(g') = \varsigma'$	$\implies E \text{ on } \mathcal{C}'[\tau'! \varsigma']$	
2534	$\mathcal{C}_2[\text{match } \tau := \varsigma \text{ with } \bar{\chi}, \square][\tau'! \varsigma']$	Above	
2535	$\blacksquare \phi \Vdash \mathcal{C}_2[\text{match } \tau := \varsigma \text{ with } \bar{\chi}, \text{match } \tau' \text{ with } \bar{\chi}']$	By CON-SUSP-CTX	
2536			
2537			□
2538	COROLLARY D.8. If $\mathcal{C}[\tau'! \varsigma]$, then $\phi(\lambda\alpha. \mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]) = \phi(\lambda\alpha. \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}])$.		
2539	Similarly, $\phi(\lambda\alpha[\bar{\alpha}]. \mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]) = \phi(\lambda\alpha[\bar{\alpha}]. \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}])$.		
2540			
2541	PROOF. It is sufficient to show that $\phi[\alpha := g] \vdash \mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]$ if and only if $\phi \vdash \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]$.		
2542			
2543	Case \implies .		
2544	$\mathcal{C}[\tau'! \varsigma]$	Premise	
2545	$\phi[\alpha := g] \vdash \mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]$	Premise	
2546	$\blacksquare \phi[\alpha := g] \vdash \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]$	Lemma D.7	
2547			
2548			

Case $\Leftarrow .$

2550 $\mathcal{C}[\tau ! \varsigma]$ Premise
2551 $\phi[\alpha := g] \vdash \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]$ Premise
2552 $\neg \phi[\alpha := g] \vdash \mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]$ By **Susp-CTX**
2553 For $\phi(\lambda\alpha[\bar{\alpha}]. \mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]) = \phi(\lambda\alpha[\bar{\alpha}]. \mathcal{C}[\text{match } \tau := \varsigma \text{ with } \bar{\chi}])$, the proof is identical.
2554
2555
2556
2557
2558
2559
2560
2561
2562
2563
2564
2565
2566
2567
2568
2569
2570
2571
2572
2573
2574
2575
2576
2577
2578
2579
2580
2581
2582
2583
2584
2585
2586
2587
2588
2589
2590
2591
2592
2593
2594
2595
2596
2597

2598 **E Properties of the constraint solver**

2599 The primary requirement of our constraint solver is correctness: a constraint C is satisfiable if and
2600 only if the solver terminates with a solution.

2601 This section decomposes this requirement into three properties: preservation, progress, and
2602 termination—and provides proofs for each. Correctness then follows as a corollary of these results.
2603

2604 **E.1 Preservation**

2605 This section details the proof of *preservation* for the solver: if $C_1 \rightarrow C_2$, then $C_1 \equiv C_2$. Since
2606 rewriting may occur under arbitrary contexts, it suffices to check for each rule, that the equivalence
2607 $C_1 \equiv C_2$ holds under all contexts \mathcal{C} .

2608 However, the introduction of suspended match constraints breaks congruence of equivalence.
2609 That is, it is no longer the case that $C_1 \equiv C_2$ implies $\mathcal{C}[C_1] \equiv \mathcal{C}[C_2]$. For instance, we have
2610 match α with $\bar{\chi} \equiv \text{false}$, yet $\mathcal{C}[\text{match } \alpha \text{ with } \bar{\chi}] \not\equiv \mathcal{C}[\text{false}]$ for $\mathcal{C} := \square \wedge \alpha = \text{int}$.

2611 As a result, we must prove *contextual equivalence* for each rewriting rule explicitly. This is both
2612 non-trivial and tedious. To simplify the task, we first present a series of auxiliary lemmas that
2613 recover contextual equivalence for many common cases. Whenever possible, we prefer to work
2614 with equivalences on *simple* constraints, as these retain the desired congruence properties that do
2615 not hold generally in our system.
2616

2617 *Definition E.1 (Contextual equivalence).* Two constraints C_1 and C_2 are contextually equivalence,
2618 written $C_1 \equiv_{\text{ctx}} C_2$, iff:

$$C_1 \equiv_{\text{ctx}} C_2 \triangleq \forall \mathcal{C}. \mathcal{C}[C_1] \equiv \mathcal{C}[C_2]$$

2619 **COROLLARY E.2 (SIMPLE EQUIVALENCE IS CONGRUENT).** Given simple constraints C_1, C_2 and simple
2620 context \mathcal{C} . If $C_1 \equiv C_2$, then $\mathcal{C}[C_1] \equiv \mathcal{C}[C_2]$.

2621 PROOF. Follows from Lemma D.2. □

2622 **LEMMA E.3 (SIMPLE EQUIVALENCE IS CONTEXTUAL).** For simple constraints C_1, C_2 . If $C_1 \equiv C_2$, then
2623 $C_1 \equiv_{\text{ctx}} C_2$.

2624 PROOF. We proceed by induction on the number of suspended match constraints n in \mathcal{C} .

2625 **Case n is 0.** Follows from Corollary E.2.

2626 **Case n is $k + 1$.**

2627 **Subcase \implies .**

2628 $\phi \vdash \mathcal{C}[C_1]$	Premise
2629 $\phi \Vdash \mathcal{C}[C_1]$	Theorem D.6
2630 $\mathcal{C}'[\tau ! \varsigma]$	Inversion of CAN-SUSP-CTX
2631 $\phi \Vdash \mathcal{C}'[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]$	"
2632 $\mathcal{C}[C_1] = \mathcal{C}'[\text{match } \tau := \varsigma \text{ with } \bar{\chi}]$	"
2633 $= \mathcal{C}_2[\text{match } \tau := \varsigma \text{ with } \bar{\chi}, C_1]$	For some two-hole context \mathcal{C}_2
2634 $\phi \vdash \mathcal{C}_2[\text{match } \tau := \varsigma \text{ with } \bar{\chi}, C_2]$	By i.h.
2635 For all ϕ', \mathbf{g}	Defn of $\mathcal{C}'[\tau ! \varsigma]$
2636 $\phi' \vdash [\mathcal{C}_2[\tau := \mathbf{g}, C_2]]$	Premise
2637 $\phi' \vdash [\mathcal{C}_2[\tau := \mathbf{g}, C_1]]$	Corollary E.2
2638 $\phi' \vdash [\mathcal{C}'[\tau := \mathbf{g}]]$	Above
2639 $\text{shape}(\mathbf{g}) = \varsigma$	$\implies \text{E on } \mathcal{C}'[\tau ! \varsigma]$

2647 $\mathcal{C}_2[\square, C_2][\tau ! \varsigma]$ Above
 2648 $\blacksquare \phi \vdash \mathcal{C}_2[\text{match } \tau \text{ with } \bar{\chi}, C_2]$ By SUSP-CTX

2649 **Subcase** \Leftarrow .

2650 Symmetric argument.

2652 \square

2653 LEMMA E.4 (UNIFICATION IS SIMPLE). *For all unification problems U, U simple.*

2655 PROOF. By induction on the structure of U . \square

2656 Definition E.5 (Context equivalence). Two contexts \mathcal{C}_1 and \mathcal{C}_2 are equivalent with guard P , written
 2658 $\mathcal{C}_1 \equiv_{\square}^P \mathcal{C}_2$ iff:

$$2659 \quad \mathcal{C}_1 \equiv_{\square}^P \mathcal{C}_2 \triangleq \forall \bar{C}. P(\bar{C}) \implies \mathcal{C}_1[\bar{C}] \equiv_{\text{ctx}} \mathcal{C}_2[\bar{C}]$$

2660 Definition E.6 (Match-closed). A predicate P on constraints is *match-closed* if, for all constraints
 2661 \bar{C}, \bar{C}' , matches match τ with $\bar{\chi}$ and shapes ς ,

$$2663 \quad P(\bar{C}, \text{match } \tau \text{ with } \bar{\chi}, \bar{C}') \implies P(\bar{C}, \text{match } \tau := \varsigma \text{ with } \bar{\chi}, \bar{C}')$$

2664 LEMMA E.7 (DETERMINES IS MATCH-CLOSED). C determines $\bar{\beta}$ is *match-closed*.

2666 PROOF. Follows from the definition of C determines $\bar{\beta}$ and Lemma D.2. \square

2668 LEMMA E.8 (SIMPLE CONTEXT EQUIVALENCE). *For any two simple contexts $\mathcal{C}_1, \mathcal{C}_2$ and a match-
 2669 closed guard P . If the two contexts \mathcal{C}_1 and \mathcal{C}_2 are equivalent under any simple constraints satisfying
 2670 P , then $\mathcal{C}_1 \equiv_{\square}^P \mathcal{C}_2$.*

2671 PROOF. Let us assume that (\dagger) holds:

$$2673 \quad \forall \mathcal{C}, \bar{C} \text{ simple. } P(\bar{C}) \implies \mathcal{C}[\mathcal{C}_1[\bar{C}]] \equiv \mathcal{C}[\mathcal{C}_2[\bar{C}]]$$

2674 We proceed by induction on the number of suspended match constraints n with the statement
 2675 $Q(n) := \forall \bar{C}, \mathcal{C}. \# \text{match } \mathcal{C} + \# \text{match } \bar{C} = n \implies P(\bar{C}) \implies \mathcal{C}[\mathcal{C}_1[\bar{C}]] \equiv \mathcal{C}[\mathcal{C}_2[\bar{C}]]$.

2676 **Case** n is 0.

$$2677 \quad \begin{array}{ll} \mathcal{C}, \bar{C} \text{ simple} & \text{Premise } (n \text{ is } 0) \\ \blacksquare P(\bar{C}) \implies \mathcal{C}[\mathcal{C}_1][\bar{C}] \equiv \mathcal{C}[\mathcal{C}_2][\bar{C}] & \dagger \end{array}$$

2680 **Case** n is $k+1$.

2681 **Subcase** \implies .

$$\begin{array}{ll} 2682 \quad P(\bar{C}) & \text{Premise} \\ 2683 \quad \phi \vdash \mathcal{C}[\mathcal{C}_1][\bar{C}] & \text{Premise} \\ 2684 \quad \phi \Vdash \mathcal{C}[\mathcal{C}_1][\bar{C}] & \text{Theorem D.6} \\ 2685 \quad \phi \Vdash \mathcal{C}'[\text{match } \tau := \varsigma \text{ with } \bar{\chi}] & \text{Inversion of CAN-SUSP-CTX} \\ 2686 \quad \mathcal{C}'[\tau ! \varsigma] & " \\ 2687 \quad \mathcal{C}[\mathcal{C}_1][\bar{C}] = \mathcal{C}'[\text{match } \tau \text{ with } \bar{\chi}] & " \\ 2688 \quad \text{Cases on } \mathcal{C}, \bar{C}. & \end{array}$$

2689 **Subsubcase** \mathcal{C} contains \mathcal{C}' 's hole.

$$\begin{array}{ll} 2691 \quad \mathcal{C}[\mathcal{C}_1][\bar{C}] = \mathcal{C}_3[\text{match } \tau \text{ with } \bar{\chi}, \mathcal{C}_1[\bar{C}]] & \text{For some 2-hole context } \mathcal{C}_3 \\ 2692 \quad \phi \Vdash \mathcal{C}_3[\text{match } \tau := \varsigma \text{ with } \bar{\chi}, \mathcal{C}_1[\bar{C}]] & \\ 2693 \quad k = \# \text{match } \mathcal{C}_3[\text{match } \tau := \varsigma \text{ with } \bar{\chi}, \mathcal{C}_1[\bar{C}]] & \\ 2694 \quad \phi \vdash \mathcal{C}_3[\text{match } \tau := \varsigma \text{ with } \bar{\chi}, \mathcal{C}_2[\bar{C}]] & \text{By i.h.} \\ 2695 & \end{array}$$

2696 For all ϕ', \mathbf{g}
 2697 $\phi' \vdash [\mathcal{C}_3[\tau = \mathbf{g}, \mathcal{C}_2[\bar{C}]]]$ Premise
 2698 $\phi' \vdash [\mathcal{C}_3[\tau = \mathbf{g}, \mathcal{C}_1[\bar{C}]]]$ †
 2699 shape (\mathbf{g}) = ς $\implies \text{E on } \mathcal{C}'[\tau ! \varsigma]$
 2700 $\mathcal{C}_3[\square, \mathcal{C}_2[\bar{C}]][\tau ! \varsigma]$ Above
 2701 $\blacksquare \quad \phi \vdash \mathcal{C}_3[\text{match } \tau \text{ with } \bar{\lambda}, \mathcal{C}_2[\bar{C}]]$ By [SUSP-CTX](#)

2702 **Subsubcase** C_i contains \mathcal{C}' 's hole.

2703 Similar argument to the above case, but relies on the match-closure of P .

2704 **Subcase** \Leftarrow .

2705 Symmetric argument.

2706 \square

2707 **LEMMA E.9 (SIMPLE LET EQUIVALENCE).** *Given simple constraints C_1, C_2 and a simple context \mathcal{C} . Suppose that*

$$2708 \quad \forall \phi, \phi', \bar{C} \text{ simple. } \phi'(x) = \phi(\lambda \alpha[\bar{\alpha}]. \mathcal{C}[\bar{C}]) \implies \phi' \vdash C_1 \iff \phi' \vdash C_2$$

2709 Then, for any context \mathcal{C}' that does not re-bind x , we have:

$$2710 \quad \text{let } x \alpha [\bar{\alpha}] = \mathcal{C}[\bar{\square}] \text{ in } \mathcal{C}'[C_1] \equiv_{\square}^P \text{let } x \alpha [\bar{\alpha}] = \mathcal{C}[\bar{\square}] \text{ in } \mathcal{C}'[C_2]$$

2711 for any match-closed guard P on the holes.

2712 **PROOF.** Let us assume (\dagger):

$$2713 \quad \forall \phi, \phi', \bar{C}. \phi'(x) = \phi(\lambda \alpha[\bar{\alpha}]. \mathcal{C}[\bar{C}]) \implies \phi' \vdash C_1 \iff \phi' \vdash C_2$$

2714 We proceed by induction on the number of suspended match constraints in $\mathcal{C}'', \mathcal{C}', \bar{C}$ with the
 2715 statement $P(n) := \forall \mathcal{C}'', \mathcal{C}', \bar{C}. \# \text{match } \mathcal{C}'', \mathcal{C}', \bar{C} = n \implies \mathcal{C}''[\text{let } x \alpha [\bar{\alpha}] = \mathcal{C}[\bar{C}] \text{ in } \mathcal{C}'[C_1]] \equiv$
 $\mathcal{C}''[\text{let } x \alpha [\bar{\alpha}] = \mathcal{C}[\bar{C}] \text{ in } \mathcal{C}'[C_2]]$.

2716 **Case** n is 0.

2717 Thus $\mathcal{C}'', \mathcal{C}', \bar{C}$ are simple. It suffices to show the equivalence on the let-constraint directly and
 2718 use congruence of equivalence for simple constraints ([Lemma E.3](#)) to establish the result.

2719 We proceed by induction on the structure of \mathcal{C}' with the statement (\ddagger):

$$2720 \quad \forall \phi, \phi'. \phi'(x) = \phi(\lambda \alpha[\bar{\alpha}]. \mathcal{C}[\bar{C}]) \implies \phi' \vdash \mathcal{C}'[C_1] \iff \phi' \vdash \mathcal{C}'[C_2]$$

2721 This holds due to the compositionality of simple equivalence using \dagger as a base case.

2722 **Subcase** \implies .

2723 $\phi \vdash \text{let } x \alpha [\bar{\alpha}] = \mathcal{C}[\bar{C}] \text{ in } \mathcal{C}'[C_1]$ Premise
 2724 $\phi \vdash \exists \alpha, \bar{\alpha}. \mathcal{C}[\bar{C}]$ Simple inversion
 2725 $\phi[x := \phi(\lambda \alpha[\bar{\alpha}]. \mathcal{C}[\bar{C}])] \vdash \mathcal{C}'[C_1]$ "
 2726 $\phi[x := \phi(\lambda \alpha[\bar{\alpha}]. \mathcal{C}[\bar{C}])] \vdash \mathcal{C}'[C_2]$ \ddagger
 2727 $\phi \vdash \text{let } x \alpha [\bar{\alpha}] = \mathcal{C}[\bar{C}] \text{ in } \mathcal{C}'[C_2]$ By [LET R](#)

2728 **Subcase** \Leftarrow .

2729 Symmetric argument.

2730 **Case** n is $k + 1$.

2731 Analogous to the inductive step in [Lemma E.8](#).

2732 \square

2733 **LEMMA E.10.** *If $\alpha = \tau = \epsilon \in \mathcal{C}$ and $\tau \notin \mathcal{V}$, then $\mathcal{C}[\alpha ! \text{shape}(\tau)]$.*

2734

2745 PROOF.

2746	$\alpha = \tau = \epsilon \in \mathcal{C}$	Premise
2747	$\tau \notin \mathcal{V}$	Premise
2748	$\tau = \text{shape}(\tau) \bar{\tau}$	For some $\bar{\tau}$
2749	$\mathcal{C} = \mathcal{C}_1[\alpha = \tau = \epsilon \wedge \mathcal{C}_2]$	By definition
2750	$\text{fv}(\alpha, \tau, \epsilon) \# \text{bv}(\mathcal{C}_2)$	"
2751	For all ϕ, g	Defn. of $\mathcal{C}[\alpha ! \text{shape}(\tau)]$
2752	$\phi \vdash [\mathcal{C}_1[\alpha = \text{shape}(\tau) \bar{\tau} = \epsilon \wedge \mathcal{C}_2[\alpha = g]]]$	Premise
2753	$\phi_1 \vdash \alpha = \text{shape}(\tau) \bar{\tau} = \epsilon$	Inversion of \mathcal{C}_1
2754	$\phi_2 \vdash \alpha = g$	Inversion of \mathcal{C}_2
2755	$g = \phi_2(\alpha)$	Simple inversion
2756	$= \phi_1(\alpha)$	$\alpha \# \text{bv}(\mathcal{C}_2)$
2757	$= \text{shape}(\tau) \phi_1(\bar{\tau})$	Simple inversion
2758	$\blacksquare \text{shape}(g) = \text{shape}(\tau)$	Applying shape to both sides
2759		□
2760		

2761 LEMMA E.11. If $\gamma = \tau = \epsilon \in \mathcal{C}[\mathcal{C}_2]$ and $\tau \notin \mathcal{V}$, then

$$\mathcal{C}[\text{let } x \alpha [\bar{\alpha}] = \mathcal{C}_1[\square] \text{ in } \mathcal{C}_2[i^x[\alpha' \rightsquigarrow \gamma]]][\alpha' ! \text{shape}(\tau)]$$

2762 PROOF. Similar proof to Lemma E.10. □

2763 LEMMA E.12 (UNIFICATION PRESERVATION). If $U_1 \longrightarrow U_2$, then $U_1 \equiv U_2$

2764 PROOF. By induction on the given derivation $U_1 \longrightarrow U_2$. See Pottier and Rémy [2005] for more
2765 details. □

2766 THEOREM E.13 (PRESERVATION). If $C_1 \longrightarrow C_2$, then $C_1 \equiv C_2$.

2767 PROOF. We proceed by induction on the given derivation. It suffices to show that for each
2768 individual rule R ($C_1 \longrightarrow_R C_2$), that $C_1 \equiv_{\text{ctx}} C_2$.

2769 Case

$$\frac{U_1 \quad U_1 \longrightarrow U_2}{U_2} \text{S-UNIF}$$

2770 $U_1 \longrightarrow U_2$ Premise

2771 $U_1 \equiv U_2$ Lemma E.12

2772 U_1, U_2 simple Lemma E.4

2773 $\blacksquare U_1 \equiv_{\text{ctx}} U_2$ Lemma E.3

2774 Case

$$\frac{(\exists \alpha. C_1) \wedge C_2 \quad \alpha \# C_2}{\exists \alpha. C_1 \wedge C_2} \text{S-EXISTS-CONJ}$$

2775 $\alpha \# C_2$ Premise

2776 Sufficient to show equivalence for simple constraints. Lemma E.8

2777 Suppose C_1, C_2 simple. Premise

2778

2779

2794 **Subcase** \implies .

2795 For all ϕ

2796 $\phi \vdash (\exists \alpha. C_1) \wedge C_2$ Premise

2797 $\phi[\alpha := g] \vdash C_1$ Simple inversion

2798 $\phi \vdash C_2$ Simple inversion

2799 $\phi[\alpha := g] \vdash C_2$ $\alpha \# C_2$

2800 $\phi[\alpha := g] \vdash C_1 \wedge C_2$ By CONJ

2801 $\phi \vdash \exists \alpha. C_1 \wedge C_2$ By EXISTS

2802 **Subcase** \Leftarrow .

2803 Symmetric argument.

2804 **Case** S-LET, S-TRUE, S-FALSE, S-LET-EXISTSLEFT, S-LET-EXISTS-INSTLEFT, S-LET-EXISTSRIGHT, S-LET-EXISTS-INSTRIGHT,
S-LET-CONJLEFT, S-LET-CONJRIGHT, S-INST-NAME, S-EXISTS-EXISTS-INST, S-EXISTS-INST-CONJ, S-EXISTS-INST-LET,
S-EXISTS-INST-SOLVE, S-ALL-CONJ.

2805 Similar argument to the S-EXISTS-CONJ case.

2806 **Case**

$$\frac{\text{match } \tau \text{ with } \bar{\chi} \quad \tau \notin \mathcal{V}}{\text{match } \tau := \text{shape}(\tau) \text{ with } \bar{\chi}} \rightarrow \text{S-MATCH-TYPE}$$

2807 $\tau \notin \mathcal{V}$ Premise
 2808 $\square[\tau ! \text{shape}(\tau)]$ By definition

2809 Sufficient to show equivalences between constraints. Lemma D.3

2810 **Subcase** \implies .

2811 For all ϕ

2812 $\phi \vdash \text{match } \tau \text{ with } \bar{\chi}$ Premise

2813 $\phi \vdash \text{match } \tau := \text{shape}(\tau) \text{ with } \bar{\chi}$ Lemma D.7

2814 **Subcase** \Leftarrow .

2815 For all ϕ

2816 $\phi \vdash \text{match } \tau := \text{shape}(\tau) \text{ with } \bar{\chi}$ Premise

2817 $\phi \vdash \text{match } \tau \text{ with } \bar{\chi}$ By SUSP-CTX

2818 **Case**

$$\frac{\mathcal{C}[\text{match } \alpha \text{ with } \bar{\chi}] \quad \alpha = \tau = \epsilon \in \mathcal{C}}{\mathcal{C}[\text{match } \alpha := \text{shape}(\tau) \text{ with } \bar{\chi}]} \rightarrow \text{S-MATCH-VAR}$$

2819 $\alpha = \tau = \epsilon \in \mathcal{C}$ Premise
 2820 $\mathcal{C}[\alpha ! \text{shape}(\tau)]$ Lemma E.10

2821 Sufficient to show equivalences between constraints. Lemma D.3

2822 **Subcase** \implies .

2823 For all ϕ

2824 $\phi \vdash \mathcal{C}[\text{match } \alpha \text{ with } \bar{\chi}]$ Premise

2825 $\phi \vdash \mathcal{C}[\text{match } \alpha := \text{shape}(\tau) \text{ with } \bar{\chi}]$ Lemma D.7

2826 **Subcase** \Leftarrow .

2827 For all ϕ

2828 $\phi \vdash \mathcal{C}[\text{match } \alpha \text{ with } \bar{\chi}]$ Premise

2829 $\phi \vdash \mathcal{C}[\text{match } \alpha := \text{shape}(\tau) \text{ with } \bar{\chi}]$ By SUSP-CTX

2843 **Case**

$$\frac{\text{let } x \alpha [\bar{\alpha}] = C_1 \text{ in } \mathcal{C}[x \tau] \quad \gamma \# \tau \quad x \# \text{bv}(\mathcal{C})}{\text{let } x \alpha [\bar{\alpha}] = C_1 \text{ in } \mathcal{C}[\exists y, i^x. y = \tau \wedge i[\alpha \rightsquigarrow \gamma]]} \text{S-LET-APP R}$$

$$\begin{array}{ll} \gamma \# \tau & \text{Premise} \\ x \# \text{bv}(\mathcal{C}) & \text{Premise} \end{array}$$

Sufficient to show equivalence between $x \tau$ and $\exists y, i^x. y = \tau \wedge i[\alpha \rightsquigarrow \gamma]$. Lemma E.9
 Suppose $\phi'(x) = \phi(\lambda \alpha[\bar{\alpha}]. C_1)$. Premise

2852 **Subcase** \implies .

$$\begin{array}{lll} \phi' \vdash x \tau & & \text{Premise} \\ \alpha[\phi_1] \in \phi(x) & & \text{Simple inversion} \\ \phi_1(\alpha) = \phi'(\tau) & & " \\ \phi'[\gamma := \phi'(\tau), i := \phi_1] \vdash i[\alpha \rightsquigarrow \gamma] & & \text{By PARTIAL-INST} \\ \phi'[\gamma := \phi'(\tau), i := \phi_1] \vdash \gamma = \tau & & \text{By UNIF} \\ \blacksquare \quad \phi' \vdash \exists y, i^x. y = \tau \wedge i[\alpha \rightsquigarrow \gamma] & & \text{By EXISTS, EXISTS-INST and CONJ} \end{array}$$

2860 **Subcase** \Leftarrow .

2861 Symmetric argument.

2862 **Case**

$$\frac{\begin{array}{c} \text{let } x \alpha [\bar{\alpha}] = C \text{ in } \mathcal{C}[i^x[\alpha' \rightsquigarrow \gamma]] \\ C = C' \wedge \alpha' = \varsigma \bar{\beta} = \epsilon \quad \alpha' \in \alpha, \bar{\alpha} \quad \neg\text{cyclic}(C) \quad \bar{\beta}' \# \alpha', \gamma, \bar{\beta} \quad x \# \text{bv}(\mathcal{C}) \end{array}}{\text{let } x \alpha [\bar{\alpha}] = C \text{ in } \mathcal{C}[\exists \bar{\beta}'. \gamma = \varsigma \bar{\beta}' \wedge i^x[\bar{\beta} \rightsquigarrow \bar{\beta}']] \text{S-INST-COPY}}$$

$$x \# \text{bv}(\mathcal{C}) \quad \text{Premise}$$

$$\bar{\beta}' \# \alpha', \gamma, \bar{\beta} \quad \text{Premise}$$

Sufficient to show equivalence between $i^x[\alpha' \rightsquigarrow \gamma]$ and $\exists \bar{\beta}'. \gamma = \varsigma \bar{\beta}' \wedge i^x[\bar{\beta} \rightsquigarrow \bar{\beta}']$. Lemma E.9

Suppose $\phi'(x) = \phi(\lambda \alpha[\bar{\alpha}]. C)$. Premise

2873 **Subcase** \implies .

$$\begin{array}{lll} \phi' \vdash i^x[\alpha' \rightsquigarrow \gamma] & & \text{Premise} \\ \alpha[\phi_1] \in \phi(x) & & \exists i^x. \in \mathcal{C} \\ \phi'_i(i) = \phi_1 & & " \\ \phi'(y) = \phi(i)(\alpha') & & \text{Simple inversion} \\ & & \text{Above} \\ \phi_1 \vdash C' \wedge \alpha' = \varsigma \bar{\beta} = \epsilon & & \text{Above} \\ \phi_1 \vdash \alpha' = \varsigma \bar{\beta} = \epsilon & & \text{Simple inversion} \\ \phi_1(\alpha') = \varsigma \phi_1(\bar{\beta}) & & " \\ \phi'(y) = \varsigma \phi_1(\bar{\beta}) & & \text{Above} \\ \phi'[\bar{\beta}' := \phi_1(\bar{\beta})] \vdash \gamma = \varsigma \bar{\beta}' & & \text{By UNIF} \\ \phi'[\bar{\beta}' := \phi_1(\bar{\beta})] \vdash i^x[\bar{\beta} \rightsquigarrow \bar{\beta}'] & & \text{By PARTIAL-INST} \\ \blacksquare \quad \phi' \vdash \exists \bar{\beta}'. \gamma = \varsigma \bar{\beta}' \wedge i^x[\bar{\beta} \rightsquigarrow \bar{\beta}'] & & \text{By EXISTS and CONJ} \end{array}$$

2888 **Subcase** \Leftarrow .

2889 Symmetric argument.

2890

2892 **Case**

$$\frac{i[\alpha \rightsquigarrow \gamma_1] \wedge i[\alpha \rightsquigarrow \gamma_2]}{i[\alpha \rightsquigarrow \gamma_1] \wedge \gamma_1 = \gamma_2} \text{S-INST-UNIF}$$

2896 Sufficient to show equivalence between $i[\alpha \rightsquigarrow \gamma_1] \wedge i[\alpha \rightsquigarrow \gamma_2]$ and $i[\alpha \rightsquigarrow \gamma_1] \wedge \gamma_1 = \gamma_2$. Lemma E.8
 2897

2898 **Subcase** \implies .

$$\begin{array}{lll} 2900 & \phi \vdash i[\alpha \rightsquigarrow \gamma_1] \wedge i[\alpha \rightsquigarrow \gamma_2] & \text{Premise} \\ 2901 & \phi \vdash i[\alpha \rightsquigarrow \gamma_1] & \text{Simple inversion} \\ 2902 & \phi \vdash i[\alpha \rightsquigarrow \gamma_2] & " \\ 2903 & \phi(\gamma_1) = \phi(i)(\alpha) & " \\ 2904 & \phi(\gamma_2) = \phi(i)(\alpha) & " \\ 2905 & \phi(\gamma_1) = \phi(\gamma_2) & \text{Above} \\ 2906 & \phi \vdash \gamma_1 = \gamma_2 & \text{By UNIF} \\ 2907 & \blacksquare \quad \phi \vdash i[\alpha \rightsquigarrow \gamma_1] \wedge \gamma_1 = \gamma_2 & \text{By CONJ} \\ 2908 \end{array}$$

2909 **Subcase** \Leftarrow .

2910 Symmetric argument.
 2911

2912 **Case**

$$\frac{\begin{array}{c} \text{let } x \alpha [\bar{\alpha}] = \bar{\epsilon} \wedge C \text{ in } \mathcal{C}[i^x[\alpha' \rightsquigarrow \gamma]] \\ \forall \alpha'. \exists \alpha, \bar{\alpha}. \bar{\epsilon} \equiv \text{true} \quad \alpha' \in \alpha, \bar{\alpha} \quad \alpha' \# C \quad i.\alpha' \# \text{insts}(\mathcal{C}) \quad x \# \text{bv}(\mathcal{C}) \end{array}}{\text{let } x \alpha [\bar{\alpha}] = \bar{\epsilon} \wedge C \text{ in } \mathcal{C}[\text{true}]} \text{S-INST-POLY}$$

$$\begin{array}{ll} 2917 & \forall \alpha'. \exists \alpha, \bar{\alpha}. \bar{\epsilon} \equiv \text{true} & \text{Premise} \\ 2918 & \alpha' \# C & \text{Premise} \\ 2919 & i.\alpha' \# \text{insts}(\mathcal{C}) & \text{Premise} \\ 2920 & x \# \text{bv}(\mathcal{C}) & \text{Premise} \\ 2921 \end{array}$$

2922 Sufficient to show equivalence between $i^x[\alpha' \rightsquigarrow \gamma]$ and true. Lemma E.9
 2923 Suppose $\phi'(x) = \phi(\lambda \alpha[\bar{\alpha}, \alpha']. \bar{\epsilon} \wedge C)$.
 2924 Premise

2925 **Subcase** \implies .

$$\begin{array}{ll} 2926 & \phi' \vdash i^x[\alpha' \rightsquigarrow \gamma] & \text{Premise} \\ 2927 & \blacksquare \quad \phi' \vdash \text{true} & \text{By TRUE} \\ 2928 \end{array}$$

2929 **Subcase** \Leftarrow .

$$\begin{array}{ll} 2930 & \phi' \vdash \text{true} & \text{Premise} \\ 2931 & \alpha[\phi_1] \in \phi'(x) & \mathcal{C} = \mathcal{C}_1[\exists i^x. \mathcal{C}_2] \\ 2932 & \phi'(i) = \phi_1 & " \\ 2933 & \text{By cases on } \phi_1(\alpha'). \\ 2934 \end{array}$$

2935 **Subsubcase** $\phi_1(\alpha') = \phi'(\gamma)$.

$$\begin{array}{ll} 2936 & \phi_1(\alpha') = \phi'(\gamma) & \text{Premise} \\ 2937 & \blacksquare \quad \phi' \vdash i^x[\alpha' \rightsquigarrow \gamma] & \text{By PARTIAL-INST} \\ 2938 \end{array}$$

2939 **Subsubcase** $\phi_1(\alpha') \neq \phi'(\gamma)$.

2990	For all ϕ	
2991	$\phi \vdash \text{let } x \alpha [\bar{\alpha}] = \bar{\epsilon} \text{ in } C$	Premise
2992	$\phi \vdash \exists \alpha, \bar{\alpha}. \bar{\epsilon}$	Simple inversion
2993	$\phi[x := \phi(\lambda \alpha[\bar{\alpha}]. \bar{\epsilon})] \vdash C$	"
2994	$\blacksquare \phi \vdash C$	$x \# C$
2995	Subcase \Leftarrow .	
2996	For all ϕ	
2997	$\phi \vdash C$	Premise
2998	$\phi[x := \phi(\lambda \alpha[\bar{\alpha}]. \bar{\epsilon})] \vdash C$	$x \# C$
2999	$\phi \vdash \exists \alpha, \bar{\alpha}. \bar{\epsilon}$	
3000	$\blacksquare \phi \vdash \text{let } x \alpha [\bar{\alpha}] = \bar{\epsilon} \text{ in } C$	By LETR
3001	Case	
3002	$\text{let } x \alpha [\bar{\alpha}, \bar{\beta}] = C_1 \text{ in } C_2 \quad \exists \alpha, \bar{\alpha}. C_1 \text{ determines } \bar{\beta}$	$\xrightarrow{\text{S-EXISTS-LOWER}}$
3003	$\exists \bar{\beta}. \text{let } x \alpha [\bar{\alpha}] = C_1 \text{ in } C_2$	
3004	$\exists \alpha, \bar{\alpha}. C_1 \text{ determines } \bar{\beta}$	Premise
3005	Sufficient to show equivalence for simple constraints.	Lemma E.8 and Lemma E.7
3006	Suppose C_1, C_2 simple.	Premise
3007	Subcase \implies .	
3008	$\phi \vdash \text{let } x \alpha [\bar{\alpha}, \bar{\beta}] = C_1 \text{ in } C_2$	Premise
3009	$\phi \vdash \exists \alpha, \bar{\alpha}, \bar{\beta}. C_1$	Simple inversion
3010	$\phi[x := \phi(\lambda \alpha[\bar{\alpha}, \bar{\beta}]. C_1)] \vdash C_2$	"
3011	$\phi[\alpha := g, \bar{\alpha} := \bar{g}, \bar{\beta} := \bar{g}'] \vdash C_1$	"
3012	$\phi[\bar{\beta} := \bar{g}'] \vdash \exists \alpha, \bar{\alpha}. C_1$	By EXISTS
3013	Sufficient to show $\phi[x := \phi(\lambda \alpha[\bar{\alpha}, \bar{\beta}]. C_1)] = \phi[\bar{\beta} := \bar{g}'](\lambda \alpha[\bar{\alpha}]. C_1)$.	
3014	Subsubcase \implies .	
3015	$\phi[\alpha := g_1, \bar{\alpha} := \bar{g}_1, \bar{\beta} := \bar{g}_2] \vdash C_1$	Premise
3016	$\phi[\bar{\beta} := \bar{g}_2] \vdash \exists \alpha, \bar{\alpha}. C_1$	By EXISTS
3017	$\bar{g}_2 = \bar{g}'$	By definition of determines
3018	$\blacksquare \phi[\bar{\beta} := \bar{g}', \alpha := g_1, \bar{\alpha} := \bar{g}_1] \vdash C_1$	Above
3019	Subsubcase \Leftarrow .	
3020	Symmetric argument.	
3021	Subcase \Leftarrow .	
3022	Symmetric argument.	
3023	Case	
3024	$\mathcal{C}[\text{let } x \alpha [\bar{\alpha}] = \mathcal{C}_1[\text{match } \alpha' \text{ with } \bar{\chi}] \text{ in } \mathcal{C}_2[i^x[\alpha' \rightsquigarrow \gamma]]]$	
3025	$\alpha' \in \alpha, \bar{\alpha} \quad \gamma = \tau = \epsilon \in \mathcal{C}[\mathcal{C}_2] \quad x \# \text{bv}(\mathcal{C}_2)$	
3026	$\mathcal{C}[\text{let } x \alpha [\bar{\alpha}] = \mathcal{C}_1[\text{match } \alpha' := \text{shape } (\tau) \text{ with } \bar{\chi}] \text{ in } \mathcal{C}_2[i^x[\alpha' \rightsquigarrow \gamma]]]$	$\xrightarrow{\text{S-BACKPROP}}$
3027	Similar argument to S-MATCH-VAR , using Lemma E.11 .	
3028	Case S-COMPRESS , S-GC , S-EXISTS-ALL , S-ALL-ESCAPE , S-ALL-RIGID , S-ALL-SOLVE .	
3029	Similar argument. Use Lemma E.8 . The simple equivalences are standard, see Pottier and Rémy [2005].	

3039

3040

3041

E.2 Progress

3042 LEMMA E.14 (UNIFICATION PROGRESS). *If unification problem U cannot take a step $U \longrightarrow U'$, then
3043 either:*

- 3044 (i)
- U
- is solved.
-
- 3045 (ii)
- U
- is false.
-
- 3046

3047 PROOF. This is a standard result. See Pottier and Rémy [2005]. □

3048

3049 THEOREM E.15 (PROGRESS). *If constraint C cannot take a step $C \longrightarrow C'$, then either:*

- 3050 (i)
- C
- is solved.
-
- 3051 (ii)
- C
- is stuck, it is either: (a) false; (b)
- $\hat{\mathcal{C}}[x \tau]$
- where
- $x \# \hat{\mathcal{C}}$
- ; (c)
- $\hat{\mathcal{C}}[i^x[\alpha \rightsquigarrow \gamma]]$
- where
- $x \# \hat{\mathcal{C}}$
- and
-
- 3052
- $i.\alpha \# \text{insts}(\hat{\mathcal{C}})$
- ; (d) for every match constraint
- $\hat{\mathcal{C}}[\text{match } \alpha \text{ with } \bar{\chi}]$
- in
- C
- ,
- $\hat{\mathcal{C}}[\alpha ! \varsigma]$
- does not
-
- 3053 hold for any
- ς
- . Here,
- $\hat{\mathcal{C}}$
- is a normal context i.e., such that no other rewrites can be applied.
-
- 3054

3055 PROOF. We proceed by induction on the structure of C . We focus on suspended match constraints,
3056 conjunctions, and let rules.3057 Case match τ with $\bar{\chi}$. We have two cases:3058 Subcase τ is a non-variable type. Apply S-MATCH-TYPE.3059 Subcase τ is a type variable α .3060 We have $\square[\alpha \mathbb{X}]$. It suffices that every match constraint in a context-reachable position
3061 $\hat{\mathcal{C}}[\text{match } \alpha' \text{ with } \bar{\chi}]$ satisfies $\hat{\mathcal{C}}[\alpha' \mathbb{X}]$. By the definition of constraint contexts, there is
3062 only one such $\hat{\mathcal{C}}$, namely \square , for which we already have $\square[\alpha \mathbb{X}]$. Hence match τ with $\bar{\chi}$ is
3063 stuck.
30643065 Case $C_1 \wedge C_2$. We begin by inducting on C_1 and C_2 . Then we consider cases:3066 Subcase C_1 (or C_2) take a step. Apply congruence rewriting rule.3067 Subcase C_1 (or C_2) is true. Apply S-TRUE.3068 Subcase C_1 (or C_2) is false. Apply S-FALSE.3069 Subcase C_1 (or C_2) begins with \exists . Apply S-EXISTS-CONJ.3070 Subcase C_1, C_2 are solved.3071 We either apply the above \exists case, or both C_1 and C_2 are solved multi-equations $\bar{\epsilon}_1, \bar{\epsilon}_2$. We
3072 perform cases on this:3073 Subsubcase $\bar{\epsilon}_1$ and $\bar{\epsilon}_2$ are mergeable. Apply U-MERGE.3074 Subsubcase cyclic ($\bar{\epsilon}_1, \bar{\epsilon}_2$). Apply U-CYCLE.3075 Subsubcase Otherwise. The conjunction $\bar{\epsilon}_1 \wedge \bar{\epsilon}_2$ is solved.3076 Subcase C_1 and C_2 are stuck (and not false).3077 w.l.o.g., consider cases C_1 .3078 Subsubcase $\hat{\mathcal{C}}_1[x \tau]$. We have $x \# \text{bv}(\hat{\mathcal{C}}_1)$.3079 $\hat{\mathcal{C}}_1[x \tau] \wedge C_2$ is stuck as we do not bind x in $\hat{\mathcal{C}}_1 \wedge C_2$.3080 Subsubcase $\hat{\mathcal{C}}_1[i^x[\alpha \rightsquigarrow \gamma]]$. We have $x \# \text{bv}(\hat{\mathcal{C}}_1)$ and $i.\alpha \# \text{insts}(\hat{\mathcal{C}}_1)$.3081 If $i.\alpha \in \text{insts}(C_2)$ and $i \# \text{bv}(\hat{\mathcal{C}}_1)$, then apply S-INST-UNIFY. It must be the case that we can ap-
3082 ply S-INST-UNIFY, otherwise, we could lift these instantiation constraints using S-EXISTS-LOWER
3083 and S-LET-CONJLEFT, contradicting that $\hat{\mathcal{C}}_1$ is stuck.3084 Otherwise, $x \# \text{bv}(\hat{\mathcal{C}}_1 \wedge C_2)$, thus $\hat{\mathcal{C}}_1[i^x[\alpha \rightsquigarrow \gamma]]$ is stuck.3085 Subsubcase $\hat{\mathcal{C}}_1[\text{match } \alpha' \text{ with } \bar{\chi}]$. We have $\mathcal{C}_1[\alpha' \mathbb{X}]$.3086 Consider a match constraint match α' with $\bar{\chi}$ in C_1 .

3087

If $\alpha' = \tau = \epsilon \in C_2$ and $\tau \notin \mathcal{V}$. By the above logic, it must be at the root (otherwise C_2 is not stuck). So we have $\alpha' = \tau = \epsilon \in \hat{\mathcal{C}}_1 \wedge C_2$. Thus we can apply S-MATCH-TYPE.

If $\gamma = \tau = \epsilon \in C_2$, $\tau \notin \mathcal{V}$, and $\hat{\mathcal{C}}_1$ contains let $x \alpha [\bar{\alpha}] = \hat{\mathcal{C}}_3$ [match α' with $\bar{\chi}$] in $\hat{\mathcal{C}}_4[i^x[\alpha' \rightsquigarrow \gamma]]$. Apply S-BACKPROP.

Otherwise, we are stuck and $(\mathcal{C}_1 \wedge C_2)[\alpha' \times]$.

Case let $x \alpha [\bar{\alpha}] = C_1$ in C_2 . We begin by inducting on C_1 and C_2 . Then we consider cases:

Subcase C_1 (or C_2) take a step. Apply congruence rewriting rule.

Subcase C_1 (or C_2) is false. Apply S-FALSE.

Subcase C_1 begins with \exists . Apply S-LET-EXISTSLEFT

Subcase C_2 begins with \exists . Apply S-LET-EXISTSRIGHT

Subcase C_2 begins with \wedge with $x \#$ from conjunct. Apply S-LET-CONJRIGHT.

Subcase C_1 begins with \wedge with $\alpha, \bar{\alpha} \#$ from conjunct . Try apply S-LET-CONJLEFT

Subcase C_2 begins with $\exists i^x . , x \neq x'$. Apply S-EXISTS-INST-LET

Subcase $\alpha' \in \bar{\alpha}$ is determined by C_1 . Apply S-EXISTS-LOWER

Subcase C_2 is solved.

Thus C_2 must be true (due to above cases).

Subsubcase C_1 is solved. Thus C_1 must be $\bar{\epsilon}$.

There are two cases:

- $\exists \alpha, \bar{\alpha}, \bar{\epsilon} \equiv \text{true}$. Apply S-LET-SOLVE.

- $\exists \alpha, \bar{\alpha}, \bar{\epsilon} \not\equiv \text{true}$. It must be the case there is some β that dominates a α' in $\alpha, \bar{\alpha}$ in $\bar{\epsilon}$. Hence $\exists \alpha, \bar{\alpha} \setminus \alpha'. \bar{\epsilon}$ determines α' . So we can apply S-EXISTS-LOWER.

Subsubcase C_1 is stuck.

The constraint let $x \alpha [\bar{\alpha}] = C_1$ in C_2 remains stuck, since no additional term variable bindings occur for the scope of C_1 , ruling out the instantiation cases. Additionally, we cannot apply backpropagation since C_2 is true.

Subcase C_2 is stuck.

Subsubcase $\hat{\mathcal{C}}[x \tau]$. We have $x \# \text{bv}(\hat{\mathcal{C}})$.

Apply S-LET-APPB.

Subsubcase $\hat{\mathcal{C}}[i^x[\alpha' \rightsquigarrow \gamma]]$. We have $x \# \text{bv}(\hat{\mathcal{C}})$ or $i.\alpha' \# \text{insts}(\hat{\mathcal{C}})$.

- $\alpha' \in \alpha, \bar{\alpha}$.

We can either apply S-INST-COPY or S-COMPRESS if a multi-equation involving α' occurs in C_1 .

Otherwise, we consider cases where C_1 is solved or stuck.

If C_1 is solved, then it must be of the form $\bar{\epsilon}$. There are two cases:

- $\exists \alpha, \bar{\alpha}, \bar{\epsilon} \equiv \text{true}$. As α' does not appear in the head position of any multi-equation in $\bar{\epsilon}$, it must be polymorphic. Thus $\forall \alpha'. \exists \alpha, \bar{\alpha} \setminus \alpha'. \bar{\epsilon} \equiv \text{true}$. So we can apply S-INST-POLY.
- $\exists \alpha, \bar{\alpha}, \bar{\epsilon} \not\equiv \text{true}$. Apply S-LOWER-EXISTS (using the same logic as above).

If C_1 is stuck, then neither case regarding instantiations in C_1 is fixed, so in these cases the constraint remains stuck. If C_1 is stuck with $\hat{\mathcal{C}}'$ [match β with $\bar{\chi}'$]. Then either backpropagation (S-BACKPROP) applies with an equation in $\hat{\mathcal{C}}$, or the entire constraint is stuck.

- $\alpha' \notin \alpha, \bar{\alpha}$. Apply S-INST-MONO.

Subsubcase For any $\hat{\mathcal{C}}[\text{match } \alpha' \text{ with } \bar{\chi}]$. We have $\hat{\mathcal{C}}[\alpha' \times]$.

Either let $x \alpha [\bar{\alpha}] = C_1$ in C_2 can progress with an instantiation constraint (in the above case) to discharge the match constraint or let $x \alpha [\bar{\alpha}] = C_1$ in C_2 is stuck.

□

3137 **E.3 Termination**

3138 This section presents a proof of termination for our solver. Most rewrite rules, in both unification
 3139 and constraint solving, are *destructive*—that is, they eliminate or modify the structure of a constraint
 3140 in a way that prevents the rule from being applied again. Consequently, to establish termination, it
 3141 suffices to consider only those rules that are not inherently destructive.

3142

3143 **LEMMA E.16 (UNIFICATION TERMINATION).** *The unifier terminates on all inputs.*

3144 PROOF. Let every shape ς have an integer *weight* defined by $\text{sw}(\varsigma) \triangleq 4 + 2 \times |\varsigma|$, where $|\varsigma|$ is
 3145 the arity of the shape ς . The weight of a type $\text{tw}(\tau)$ is defined by:

3146

$$\begin{aligned}\text{tw}(\alpha) &\triangleq 1 \\ \text{tw}(\varsigma \bar{\tau}) &\triangleq \text{iw}(\varsigma \bar{\tau}) - 2 \\ \text{iw}(\alpha) &\triangleq 0 \\ \text{iw}(\varsigma \bar{\tau}) &\triangleq \text{sw}(\varsigma) + \text{iw}(\bar{\tau}) \\ \text{iw}(\bar{\tau}) &\triangleq \sum_{i=1}^n \text{iw}(\tau_i)\end{aligned}$$

3152 The helper $\text{iw}(\tau)$ computes the “internal” weight of τ ; in the common case of shallow types it is
 3153 just the weight of its head shape.

3154 We define the weight of a multi-equation as the sum of the weights of its members. The weight
 3155 of a unification problem $\text{uw}(U)$ is defined as the sum of the weights of its multi-equations.

3156 In $U \rightarrow U'$, the rules **U-DECOMP** and **U-NAME** are not obviously destructive, as they may introduce
 3157 new constraints that are structurally larger than the constraint being rewritten.

3158 However, we show that this is not problematic: in both cases, the unification weight $\text{uw}(U)$
 3159 strictly decreases. The remaining rules are obviously destructive and either maintain or decrease
 3160 the unification weight.

3161 **Case**

3162

$$\frac{\varsigma \bar{\alpha} = \varsigma \bar{\beta} = \epsilon}{\varsigma \bar{\alpha} = \epsilon \wedge \bar{\alpha} = \bar{\beta}} \xrightarrow{\text{U-DECOMP}}$$

3163 We have:

$$\begin{aligned}(+) \quad \text{uw}(\varsigma \bar{\alpha} = \varsigma \bar{\beta} = \epsilon) &= \text{tw}(\varsigma \bar{\alpha}) + \text{tw}(\varsigma \bar{\beta}) + \text{tw}(\epsilon) \\ (-) \quad \text{uw}(\varsigma \bar{\alpha} = \epsilon \wedge \bar{\alpha} = \bar{\beta}) &= \text{tw}(\varsigma \bar{\alpha}) + \text{tw}(\epsilon) + \text{tw}(\bar{\alpha}) + \text{tw}(\bar{\beta}) \\ &= \text{tw}(\varsigma \bar{\beta}) - \text{tw}(\bar{\alpha}) - \text{tw}(\bar{\beta}) \\ &= (\text{sw}(\varsigma) + 0 - 2) - 2|\varsigma| \\ &= (2 + 2|\varsigma|) - 2|\varsigma| = 2\end{aligned}$$

3164 Hence $\text{uw}(\varsigma \bar{\alpha} = \varsigma \bar{\beta} = \epsilon) > \text{uw}(\varsigma \bar{\alpha} = \epsilon \wedge \bar{\alpha} = \bar{\beta})$.

3165 **Case**

3166

$$\frac{\varsigma(\bar{\tau}, \tau_i, \bar{\tau}') = \epsilon \quad \alpha \# \bar{\tau}, \bar{\tau}', \epsilon \quad \tau_i \notin \mathcal{V}}{\exists \alpha. \varsigma(\bar{\tau}, \alpha, \bar{\tau}') = \epsilon \wedge \alpha = \tau_i} \xrightarrow{\text{U-NAME}}$$

3167 Given $\tau_i \notin \mathcal{V}$, by **Theorem D.1**, $\tau_i = \varsigma' \bar{\tau}''$ for some shape ς' and types $\bar{\tau}''$. So we have:

3168

$$\begin{aligned}(+) \quad \text{uw}(\varsigma(\bar{\tau}, \tau_i, \bar{\tau}') = \epsilon) &= \text{sw}(\varsigma) + \text{iw}(\bar{\tau}) + \text{iw}(\tau_i) + \text{iw}(\bar{\tau}') - 2 + \text{uw}(\epsilon) \\ (-) \quad \text{uw}(\exists \alpha. \alpha = \tau_i \wedge \varsigma(\bar{\tau}, \alpha, \bar{\tau}') = \epsilon) &= \text{sw}(\varsigma) + \text{iw}(\bar{\tau}) + 0 + \text{iw}(\bar{\tau}') - 2 + \text{uw}(\epsilon) + 1 + \text{tw}(\tau_i) \\ &= \text{iw}(\tau_i) - \text{iw}(\alpha) - \text{tw}(\tau_i) - 1 \\ &= \text{iw}(\tau_i) - 0 - (\text{iw}(\tau_i) - 2) - 1 \\ &= 1\end{aligned}$$

3169 Hence $\text{uw}(\varsigma(\bar{\tau}, \tau_i, \bar{\tau}') = \epsilon) > \text{uw}(\exists \alpha. \varsigma(\bar{\tau}, \alpha, \bar{\tau}') = \epsilon \wedge \alpha = \tau_i)$.

3170

□

3186
 3187
3188 THEOREM E.17 (TERMINATION). *The constraint solver terminates on all inputs.*
 3189

3190 PROOF. The difficulty for termination comes from the “discharge” rules **S-MATCH-TYPE**, **S-MATCH-VAR**
 3191 which can make arbitrary sub-constraints appear in the non-suspended part of the constraint; and
 3192 from the instantiation rules that copy/duplicate existing structure in another part of the constraint,
 3193 increasing its total size.

3194 As we argued before, the other rewrite rules are *destructive*, they strictly simplify the constraint
 3195 towards a normal form and can only be applied finitely many times when taken together. The
 3196 fragment without discharge rules and incremental instantiation is also extremely similar to the
 3197 constraint language of Pottier and Rémy [2005], so their termination proof applies directly.

3198 *Discharge rules.* The discharge rules strictly decrease the number of occurrences of suspended
 3199 match constraint (if we also count nested suspended constraints), and no rewriting rule introduces
 3200 new suspended match constraints. So these discharge rules can only be applied finitely many times.
 3201 To prove termination of constraint solving, it thus suffices to prove that rewriting sequences that
 3202 do not contain one of the discharge rules (those that occur in-between two discharge rules) are
 3203 always finite.

3204
 3205 *Starting instantiations.* By a similar argument, the number of non-partial instantiations $x \tau$
 3206 decreases strictly on **S-LET-APP** when a partial instantiation starts, and is preserved by other non-
 3207 discharge rules. The rule **S-LET-APP** can thus only occur finitely many times in non-discharging
 3208 sequences, and it suffices to prove that all rewriting sequences that are non-discharging and do not
 3209 contain **S-LET-APP** are finite.

3210
 3211 *Other instantiation rules.* Among other instantiation rules, the rule of concern is **S-INST-COPY**,
 3212 which is not destructive: it introduces new instantiation constraints and structurally increases the
 3213 size of the constraint.

3214 Intuitively, **S-INST-COPY** should not endanger termination because the amount of copying it can
 3215 perform for a given instantiation is bounded by the size of the types in the constraint C it is copying
 3216 from. (C could have cyclic equations with infinite unfoldings, but **S-INST-COPY** forbids copying in
 3217 that case.) The difficulty is that rewrites to C can be interleaved with instantiation rules, so that the
 3218 equations that are being copied can grow strictly during instantiation.

3219 To control this, we perform a structural induction: to prove that (let $x \alpha [\bar{\alpha}] = C_1$ in C_2) does not
 3220 contain infinite non-discharging non-instance-starting rewrite rules, we can assume that the result
 3221 holds for the strictly smaller constraint C_1 , and then prove termination of the partial instantiations
 3222 of x in C_2 . (The notion of structural size used here is preserved by non-discharging rewrite rules,
 3223 as they do not affect the let-structure of the constraint.)

3224 Assuming that C_1 has no infinite rewriting sequence, it suffices to prove that only finitely many
 3225 rewrites in the rest of the constraint (namely C_2) can occur between each rewrite of C_1 .

3226 We define a weight that captures the contribution of types within C_1 to the partial instances
 3227 in C_2 :

$$\begin{aligned} \text{tw}(\varsigma \bar{\tau}) &\triangleq 2 \times \text{sw}(\varsigma) + \sum_{i=1}^n \text{tw}(\tau_i) \\ \text{tw}(\alpha) &\triangleq \begin{cases} \sup \{\text{tw}(\tau) : \alpha = \tau \in C_1\} & \text{if } C_1 \text{ is acyclic} \\ 0 & \text{otherwise} \end{cases} \end{aligned}$$

3228 The weight of a partial instantiation $\text{cw}(i^x[\alpha \rightsquigarrow \tau])$ is defined as the sum of $\text{tw}(\tau)$ and $\text{tw}(\alpha)$. The
 3229 weight of other constraints is given using the measure uw defined in the the proof of Lemma E.16.

3235 **Case**

$$\frac{\begin{array}{c} \text{let } x \alpha [\bar{\alpha}] = C \text{ in } \mathcal{C}[i^x[\alpha' \rightsquigarrow \gamma]] \\ C = C' \wedge \alpha' = \varsigma \bar{\beta} = \epsilon \quad \alpha' \in \alpha, \bar{\alpha} \quad \neg\text{cyclic}(C) \quad \bar{\beta}' \# \alpha', \gamma, \bar{\beta} \quad x \# \text{bv}(\mathcal{C}) \end{array}}{\text{let } x \alpha [\bar{\alpha}] = C \text{ in } \mathcal{C}[\exists \bar{\beta}'. \gamma = \varsigma \bar{\beta}' \wedge i^x[\bar{\beta} \rightsquigarrow \bar{\beta}']]}$$

S-INST-COPY

3239 We aim to show that the weight of the rewritten constraint $\exists \bar{\beta}'. \gamma = \varsigma \bar{\beta}' \wedge i^x[\bar{\beta} \rightsquigarrow \bar{\beta}']$ is strictly less than the original $i^x[\alpha' \rightsquigarrow \gamma]$.

$$\begin{aligned} \text{cw}(i^x[\alpha' \rightsquigarrow \gamma]) &= 1 + \text{tw}(\alpha) \\ &\geq 1 + 2 \times \text{sw}(\varsigma) + \sum_{i=1}^n \text{tw}(\beta_i) \\ \text{cw}(\exists \bar{\beta}'. \gamma = \varsigma \bar{\beta}' \wedge i^x[\bar{\beta} \rightsquigarrow \bar{\beta}']) &= 1 + \text{sw}(\varsigma) + \sum_{i=1}^n \text{tw}(\beta_i) + |\bar{\beta}'| \end{aligned}$$

3247 To ensure a strict decrease, it suffices to show that $\text{sw}(\varsigma) > |\bar{\beta}'|$. Given that $|\bar{\beta}'| = |\varsigma|$, and by
3248 the definition of $\text{sw}(\varsigma)$, this inequality holds. Therefore, the weight strictly decreases under
3249 S-INST-COPY.

3251 Thus the constraint solver terminates. □

3253 E.4 Correctness

3254 LEMMA E.18. *Given non-simple C constraint. If every match constraint $\mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}] = C$
3255 satisfies $\mathcal{C}[\tau \mathbb{X}]$, then C is unsatisfiable.*

3256 PROOF. By contradiction, inverting on the canonical derivation of C. □

3259 COROLLARY E.19. *For the closed-term-variable constraint C, C is satisfiable if and only if $C \longrightarrow^* \hat{C}$
3260 and \hat{C} is a solved form equivalence to C.*

3262 PROOF. We show each direction individually:

3263 **Case \implies .**

3264 By transfinite induction on the well-ordering of constraints whose existence is shown in [Theorem E.17](#).

3265 We have C is satisfiable. By [Theorem E.15](#), we have three cases:

3266 **Subcase C is solved.** We have $C \longrightarrow^* C$ and $C \equiv C$ by reflexivity. So we are done.

3267 **Subcase C is stuck.** Given C is a closed-term-variable constraint, it must be the case that either
3268 C is false or $\mathcal{C}[\text{match } \tau \text{ with } \bar{\chi}]$ and $\mathcal{C}[\tau \mathbb{X}] \varsigma$ for any shape ς .

3269 If C is false, this contradicts our assumption that C is satisfiable. Similarly, by [Lemma E.18](#), if
3270 C is $\hat{\mathcal{C}}[\text{match } \tau \text{ with } \bar{\chi}]$, then this also contradicts the satisfiability of C.

3271 **Subcase $C \longrightarrow C'$.**

3272 By [Theorem E.13](#), we have $C \equiv C'$, thus C' is satisfiable. So by induction, we have $C' \longrightarrow^* \hat{C}$
3273 and \hat{C} is a solved form equivalent to C' . By transitivity of equivalence, we therefore have
3274 $\hat{C} \equiv C$, as required.

3275 **Case \Leftarrow .**

3276 By induction on the rewriting $C \longrightarrow^* \hat{C}$.

3277 **Subcase**

$$\frac{}{\hat{C} \longrightarrow^* \hat{C}} \text{ZERO-STEP}$$

3278 We have $C = \hat{C}$ by inversion. All solved forms are satisfiable, thus C is satisfiable.

3279

Subcase

$$\frac{C \longrightarrow C' \quad C' \longrightarrow^* \hat{C}}{C \longrightarrow^* \hat{C}}$$

ONE-STEP

By induction, we have C' is satisfiable. By [Theorem E.13](#), $C \equiv C'$, hence C is satisfiable.

□

3289

3290

3291

3292

3293

3294

3295

3296

3297

3298

3299

3300

3301

3302

3303

3304

3305

3306

3307

3308

3309

3310

3311

3312

3313

3314

3315

3316

3317

3318

3319

3320

3321

3322

3323

3324

3325

3326

3327

3328

3329

3330

3331

3332

3333 **F Properties of OmniML**

3334 This section states and proves the two central metatheoretic properties of OmniML. The first is the
 3335 *soundness and completeness* of the constraint generator $\llbracket e : \alpha \rrbracket$ with respect to the OmniML typing
 3336 rules. The second is the existence of *principal types*, which follows as a consequence of soundness
 3337 and completeness: every closed well-typed term e admits a most general type.

3338 Throughout this section, we restrict our attention to *closed terms*. This is because the typing
 3339 context Γ can contain bindings to terms whose type is “guessed”. When we generate constraints for
 3340 a term e under a context Γ , we encode the type schemes in Γ as part of the constraint itself using
 3341 let-constraints. However, these schemes are treated as known within the constraint! As a result,
 3342 we assume terms are closed from the outside to avoid Γ leaking any guessed type information.
 3343

3344 **F.1 Simple syntax-directed system**

3345 As a first step towards proving soundness and completeness of constraint generation, we first
 3346 present a variant of the OmniML type system for *simple terms*. For this system, the syntax tree
 3347 completely determines the derivation tree.

3348 We use the standard technique of removing the `INST` and `GEN` rules, and always apply instantiations
 3349 in `VAR` (`VAR-SD`) and always generalize at let-bindings (`LET-SD`). We can show that this system is
 3350 sound and complete with respect to the declarative rules.

3351 **THEOREM F.1 (SOUNDNESS OF THE SYNTAX DIRECTED RULES).** *Given the simple term e . If $\Gamma \vdash_{\text{simple}}^{\text{sd}}$
 3352 $e : \tau$ then we also have $\Gamma \vdash_{\text{simple}} e : \tau$*

3354 PROOF. Induction on the given derivation. □

3356 **THEOREM F.2 (COMPLETENESS OF THE SYNTAX DIRECTED RULES).** *Given the simple term e . If
 3357 $\Gamma \vdash_{\text{simple}} e : \sigma$, then $\Gamma \vdash_{\text{simple}} e : \tau$ for any instance τ of σ .*

3359 PROOF. Induction on the given derivation. □

3361 The simple syntax-directed system has an inversion lemma:

3362 **LEMMA F.3 (SIMPLE INVERSION).**

- (i) If $\Gamma \vdash_{\text{simple}}^{\text{sd}} x : \tau$, then $x : \forall \bar{x}. \tau' \in \Gamma$ and $\tau = \tau'[\bar{x} := \bar{\tau}]$.
- (ii) If $\Gamma \vdash_{\text{simple}}^{\text{sd}} \lambda x. e : \tau$, then $\Gamma, x : \tau_1 \vdash_{\text{simple}}^{\text{sd}} e : \tau_2$ and $\tau = \tau_1 \rightarrow \tau_2$.
- (iii) If $\Gamma \vdash_{\text{simple}}^{\text{sd}} e_1 e_2 : \tau$, then $\Gamma \vdash_{\text{simple}}^{\text{sd}} e_1 : \tau' \rightarrow \tau$ and $\Gamma \vdash_{\text{simple}}^{\text{sd}} e_2 : \tau'$.
- (iv) If $\Gamma \vdash_{\text{simple}}^{\text{sd}} () : \tau$, then $\tau = 1$.
- (v) If $\Gamma \vdash_{\text{simple}}^{\text{sd}} \text{let } x = e_1 \text{ in } e_2 : \tau$, then $\Gamma \vdash_{\text{simple}}^{\text{sd}} e_1 : \tau'$, $\bar{x} \# \Gamma$, and $\Gamma, x : \forall \bar{x}. \tau' \vdash_{\text{simple}}^{\text{sd}} e_2 : \tau$.
- (vi) If $\Gamma \vdash_{\text{simple}}^{\text{sd}} (e : \exists \bar{x}. \tau') : \tau$, then $\Gamma \vdash_{\text{simple}}^{\text{sd}} e : \tau'[\bar{x} := \bar{\tau}]$ and $\tau = \tau'[\bar{x} := \bar{\tau}]$.
- (vii) If $\Gamma \vdash_{\text{simple}}^{\text{sd}} \{\bar{l} = \bar{e}\} : \tau$, then $\Gamma \vdash_{\text{simple}}^{\text{sd}} l_i = e_i : \tau$ for $1 \leq i \leq n$ and $\bar{l} ! \tau$.
- (viii) If $\Gamma \vdash_{\text{simple}}^{\text{sd}} l = e : \tau$, then $\Gamma \vdash_{\text{simple}}^{\text{sd}} e : \tau'$ and $l : \tau' \rightarrow \tau$.
- (ix) If $\Gamma \vdash_{\text{simple}}^{\text{sd}} e.l : \tau$, then $\Gamma \vdash_{\text{simple}}^{\text{sd}} e : \tau'$, $l : \tau' \rightarrow \tau$ and $l ! \tau$.
- (x) If $\Gamma \vdash_{\text{simple}}^{\text{sd}} (e_1, \dots, e_n) : \tau$, then $\Gamma \vdash_{\text{simple}}^{\text{sd}} e_i : \tau_i$ for all $1 \leq i \leq n$ and $\tau = \prod_{i=1}^n \tau_i$.
- (xi) If $\Gamma \vdash_{\text{simple}}^{\text{sd}} e.j/n : \tau$, then $\Gamma \vdash_{\text{simple}}^{\text{sd}} e : \prod_{i=1}^n \tau_i$ and $\tau = \tau_j$, with $n \geq j$.
- (xii) If $\Gamma \vdash_{\text{simple}}^{\text{sd}} [e : \exists \bar{x}. \forall \bar{\beta}. \tau'] : \tau$, then $\Gamma \vdash_{\text{simple}}^{\text{sd}} e : \tau[\bar{x} := \bar{\tau}], \bar{\beta} \# \Gamma$ and $\tau = [\forall \bar{\beta}. \tau'][\bar{x} := \bar{\tau}]$.
- (xiii) If $\Gamma \vdash_{\text{simple}}^{\text{sd}} \langle e : \exists \bar{x}. \sigma \rangle : \tau$, then $\Gamma \vdash_{\text{simple}}^{\text{sd}} e : [\sigma][\bar{x} := \bar{\tau}]$ and $\sigma \leq \tau$.
- (xiv) If $\Gamma \vdash_{\text{simple}}^{\text{sd}} \{e\} : \tau$, then $\Gamma \vdash_{\text{simple}}^{\text{sd}} e : \tau'$.

3382 **F.2 Canonicalization of typability**

3383 Our system satisfies a similar canonicalization theorem to constraint satisfiability.

3384 **LEMMA F.4 (COMPOSABILITY OF UNICITY).**

- 3385 (i) If $\mathcal{E}_1[e \triangleleft \varsigma]$, then $\mathcal{E}_2[\mathcal{E}_1][e \triangleleft \varsigma]$.
 3386 (ii) If $\mathcal{E}_1[e \triangleright \varsigma]$, then $\mathcal{E}_2[\mathcal{E}_1][e \triangleright \varsigma]$.
 3387 (iii) If $\mathcal{L}[\ell ! t]$, then $\mathcal{E}_2[\mathcal{L}][\ell ! t]$.

3388 **PROOF.** By induction on \mathcal{E}_2 . □

3389 **LEMMA F.5 (DECANONICALIZATION).** If $\Vdash e : \tau$, then $\emptyset \vdash e : \tau$.

3390 **PROOF.** By induction on the given derivation $\Vdash e : \tau$. □

3391 **THEOREM F.6 (CANONICALIZATION).** If $\vdash e : \sigma$, then $\Vdash e : \tau$ for any instance τ of σ .

3392 **PROOF.** By induction on the following measure of e :

$$\|e\| \triangleq \langle \#\text{implicit } e, |e| \rangle$$

3393 where $\langle \dots \rangle$ denotes a lexicographically ordered pair, and

- 3394 (1) $\#\text{implicit } e$ is the number of implicit constructs in e i.e., overloaded tuple projections $e.j$,
 3395 field projections $e.\ell$, records $\{\overline{\ell = e}\}$, polytype instantiations $\langle e \rangle$ and polytype boxing $[e]$.
 3396 (2) the last component $|e|$ is a structural measure of terms i.e., a application $e_1 e_2$ is larger than
 3397 the two terms e_1, e_2 .

3398 This measure is analogous to the measure $\|C\|$ for constraints. □

3399 **F.3 Unifiers**

3400 A substitution ϑ is an idempotent function from type variables to types. The (finite) domain of ϑ is
 3401 the set of type variables such that $\vartheta(\alpha) \neq \alpha$ for any $\alpha \in \text{dom } \vartheta$, while the codomain consists of the
 3402 free type variables of its range. We use the notation $[\bar{\alpha} := \bar{\tau}]$ for the substitution ϑ with domain $\bar{\alpha}$
 3403 and $\vartheta(\bar{\alpha}) = \bar{\tau}$.

3404 The constraint induced by a substitution ϑ , written $\exists \vartheta$, is $\exists \bar{\beta}. \bar{\alpha} = \bar{\tau}$ where $\bar{\beta} = \text{rng } \vartheta$, $\bar{\alpha} = \text{dom } \vartheta$
 3405 and $\vartheta(\bar{\alpha}) = \bar{\tau}$.

3406 **Definition F.7 (Unifier).** A substitution ϑ is a unifier of C if $\exists \vartheta$ entails C . A unifier ϑ of C is *most*
 3407 *general* when $\exists \vartheta$ is equivalent to C .

3408 **LEMMA F.8 (SIMPLE INVERSION OF UNIFIERS).**

- 3409 • If ϑ is a unifier of $\tau_1 = \tau_2$, then $\vartheta(\tau_1) = \vartheta(\tau_2)$.
- 3410 • For simple C_1, C_2 , if ϑ is a unifier of $C_1 \wedge C_2$, then ϑ is a unifier of C_1 and C_2 .
- 3411 • For simple C , if ϑ is a unifier of $\exists \alpha. C$, then $\vartheta[\alpha := \tau]$ is a unifier of C for some τ .
- 3412 • For simple C , if ϑ is a unifier of $\forall \alpha. C$, then ϑ is a unifier of C .

3413 **PROOF.** Follows by simple inversion. □

3414 **LEMMA F.9.** If ϑ unifies $\exists \alpha. C$, then there exists a unifier ϑ' that extends ϑ with α , where ϑ' is most
 3415 general unifier of $\exists \vartheta \wedge C$.

3416 Then $\lambda \alpha. C$ is equivalent to $\lambda \alpha. \sigma \leq \alpha$ under ϑ , where $\sigma = \forall \bar{\beta}. \vartheta'(\alpha)$ and $\bar{\beta} = \text{fv}(\vartheta'(\alpha)) \setminus \text{rng } \vartheta$. We
 3417 write this equivalent constraint abstraction as $\llbracket \lambda \alpha. C \rrbracket_{\vartheta}$.

3418 **PROOF.** See Pottier and Rémy [2005]. □

3431 LEMMA F.10 (LET INVERSION OF UNIFIERS). *For simple C_1, C_2 . If ϑ unifies let $x = \lambda\alpha. C_1$ in C_2 , then
3432 ϑ unifies $\exists\alpha. C_1$ and ϑ unifies let $x = \llbracket \lambda\alpha. C_1 \rrbracket_\vartheta$ in C_2*

3433 PROOF. Follows from Lemma F.9 and simple inversion. \square

3435 LEMMA F.11. *For two substitutions ϑ, ϑ' . If $\exists\vartheta \models \exists\vartheta'$, there exists ϑ'' such that $\vartheta = \vartheta'' \circ \vartheta'$.*

3436 PROOF. Standard result, follows from definition of $\exists\vartheta$. \square

F.4 Soundness and completeness of constraint generation

3439 LEMMA F.12. *For any term context \mathcal{E} , term e , $\llbracket \mathcal{E}[\square : \alpha] : \beta \rrbracket [\llbracket e : \alpha \rrbracket] = \llbracket \mathcal{E}[e] : \beta \rrbracket$.*

3440 PROOF. By induction on the structure of \mathcal{E} . \square

3441 LEMMA F.13. *For any term e , $\lfloor \llbracket e : \alpha \rrbracket \rfloor = \llbracket \lfloor e \rfloor : \alpha \rrbracket$.*

3442 PROOF. By induction on e . \square

3443 LEMMA F.14 (SIMPLE SOUNDNESS AND COMPLETENESS). *For simple terms e . $\vartheta(\Gamma) \vdash_{\text{simple}}^{\text{sd}} e : \vartheta(\tau)$ if
3444 and only if ϑ is a unifier of $\llbracket \Gamma \vdash e : \tau \rrbracket$.*

3445 PROOF. By induction on e simple. \square

3446 THEOREM F.15 (SOUNDNESS AND COMPLETENESS). $\Vdash e : \vartheta(\alpha)$ if and only if ϑ is a unifier of $\llbracket e : \alpha \rrbracket$

3447 PROOF. By induction on the number n of implicit terms in e .

3448 Case n is 0.

$$\begin{array}{c} e \text{ simple} & \text{Premise} \\ \emptyset \vdash_{\text{simple}}^{\text{sd}} e : \vartheta(\alpha) \iff \vartheta \text{ unifies } \llbracket e : \alpha \rrbracket & \text{Lemma F.14} \\ \emptyset \vdash_{\text{simple}}^{\text{sd}} e : \vartheta(\alpha) \iff \Vdash e : \vartheta(\alpha) & \text{When } e \text{ simple} \\ \Leftrightarrow \Vdash e : \vartheta(\alpha) \iff \vartheta \text{ unifies } \llbracket e : \alpha \rrbracket & \text{Above} \end{array}$$

3449 Case n is $k + 1$.

3450 Subcase \Rightarrow .

3451 Subsubcase

$$\frac{\mathcal{E}[e \triangleright v\bar{\gamma}. \prod_{i=1}^n \bar{\gamma}] \quad \vartheta(\Gamma) \Vdash \mathcal{E}[e.j/n] : \vartheta(\alpha)}{\Vdash \mathcal{E}[e.j] : \vartheta(\alpha)} \text{ CAN-PROJ-I}$$

$$\begin{aligned} & \vartheta(\Gamma) \Vdash \mathcal{E}[e.j/n] : \vartheta(\alpha) && \text{Premise} \\ & \vartheta \text{ unifies } \llbracket \Gamma \vdash \mathcal{E}[e.j/n] : \alpha \rrbracket && \text{By i.h.} \\ & \llbracket \Gamma \vdash \mathcal{E}[e.j/n] : \alpha \rrbracket = \text{let } \Gamma \text{ in } \llbracket \mathcal{E}[e.j/n] : \alpha \rrbracket && \text{By definition} \\ & & = \text{let } \Gamma \text{ in } \llbracket \mathcal{E}[\square : \beta] : \alpha \rrbracket [\llbracket e.j/n : \beta \rrbracket] & \text{Lemma F.12} \\ & \llbracket e.j/n : \beta \rrbracket \equiv \exists \alpha_1 \bar{\gamma}. \llbracket e : \alpha_1 \rrbracket \wedge \alpha_1 = \prod_{i=1}^n \bar{\gamma} \wedge \beta = \gamma_j && \text{By definition} \\ & & \equiv \exists \alpha_1. \llbracket e : \alpha_1 \rrbracket \wedge \text{match } \alpha_1 := v\bar{\gamma}. \prod_{i=1}^n \bar{\gamma} \text{ with } \Pi \gamma_j \rightarrow \beta = \gamma & " \\ & \vartheta \text{ unifies let } \Gamma \text{ in } \llbracket \mathcal{E}[\square : \beta] : \alpha \rrbracket [\exists \alpha_1. \llbracket e : \alpha_1 \rrbracket \wedge \dots] && \text{Above} \\ & & \mathcal{E}[e \triangleright v\bar{\gamma}. \prod_{i=1}^n \bar{\gamma}] & \text{Premise} \\ & \text{Let } \mathcal{C} = \text{let } \Gamma \text{ in } \llbracket \mathcal{E}[\square : \beta] : \alpha \rrbracket [\exists \alpha_1. \llbracket e : \alpha_1 \rrbracket \wedge \square]. && \text{Premise} \\ & \phi \vdash \lfloor \mathcal{C}[\alpha_1 = g] \rfloor && \text{Premise} \\ & \exists \alpha_1. \llbracket e : \alpha_1 \rrbracket \wedge \alpha_1 = g = \exists \alpha_1. \llbracket (e : g) : \alpha_1 \rrbracket && \text{By definition} \end{aligned}$$

3480	$= \llbracket \{(e : g)\} : \beta \rrbracket$	"
3481	$\llbracket \mathcal{C}[\alpha_1 = g] \rrbracket = \llbracket \text{let } \Gamma \text{ in } \llbracket \mathcal{E}[\square : \beta] : \alpha \rrbracket [\llbracket \{(e : g)\} : \beta \rrbracket] \rrbracket$	"
3482	$= \llbracket \text{let } \Gamma \text{ in } \llbracket \mathcal{E}[\{(e : g)\}] : \alpha \rrbracket \rrbracket$	Lemma F.12
3483	$= \text{let } \Gamma \text{ in } \llbracket \llbracket \mathcal{E}[\{(e : g)\}] : \alpha \rrbracket \rrbracket$	By definition
3484	$= \text{let } \Gamma \text{ in } \llbracket \llbracket \mathcal{E}[\{(e : g)\}] \rrbracket : \alpha \rrbracket$	Lemma F.13
3485	$\phi \text{ unifies let } \Gamma \text{ in } \llbracket \llbracket \mathcal{E}[\{(e : g)\}] \rrbracket : \alpha \rrbracket$	Above
3486	$\vdash \llbracket \mathcal{E}[\{(e : g)\}] \rrbracket : \phi(\alpha)$	By i.h.
3487	$\emptyset \vdash \llbracket \mathcal{E}[\{(e : g)\}] \rrbracket : \phi(\alpha)$	Lemma F.5
3488	$\text{shape } (g) = v\bar{\gamma}. \Pi_{i=1}^n \bar{\gamma}$	$\implies E$
3489	$\mathcal{C}[\alpha_1 ! v\bar{\gamma}. \Pi_{i=1}^n \bar{\gamma}]$	Above
3490	$\vartheta \text{ unifies } \mathcal{C}[\text{match } \alpha_1 \text{ with } \Pi \gamma_j \rightarrow \beta = \gamma]$	By Susp-CTX
3491	$\llbracket e.j : \beta \rrbracket = \exists \alpha_1. \llbracket e : \alpha_1 \rrbracket \wedge \text{match } \alpha_1 \text{ with } \dots$	By definition
3492	$\mathcal{C}[\text{match } \alpha_1 \text{ with } \dots] = \text{let } \Gamma \text{ in } \llbracket \mathcal{E}[\square : \beta] : \alpha \rrbracket [\exists \alpha_1. \llbracket e : \alpha_1 \rrbracket \wedge \dots]$	"
3493	$= \text{let } \Gamma \text{ in } \llbracket \mathcal{E}[\square : \beta] : \alpha \rrbracket [\llbracket e.j : \beta \rrbracket]$	Above
3494	$= \text{let } \Gamma \text{ in } \llbracket \mathcal{E}[e.j] : \alpha \rrbracket$	Lemma F.12
3495	$= \llbracket \mathcal{E}[e.j] : \alpha \rrbracket$	
3496	$\vartheta \text{ unifies } \llbracket \mathcal{E}[e.j] : \alpha \rrbracket$	
3497	\blacksquare	
3498		

Subsubcase *CAN-POLY-I*, *CAN-USE-I*, *CAN-LAB-I*.

Similar arguments.

Subcase \Leftarrow .

Subsubcase

$$\frac{\mathcal{C}[\alpha_1 ! v\bar{\gamma}. \Pi_{i=1}^n \bar{\gamma}] \quad \vartheta \text{ unifies } \mathcal{C}[\text{match } \alpha_1 := v\bar{\gamma}. \Pi_{i=1}^n \bar{\gamma} \text{ with } \dots]}{\vartheta \text{ unifies } \underbrace{\mathcal{C}[\text{match } \alpha_1 \text{ with } \Pi \gamma_j \rightarrow \beta = \gamma]}_{\llbracket e : \alpha \rrbracket}} \text{ CAN-SUSP-CTX}$$

3504	$\llbracket e : \tau \rrbracket = \text{let } \Gamma \text{ in } \llbracket \mathcal{E}[e.j] : \alpha \rrbracket$	Premise
3505	$\mathcal{C} = \text{let } \Gamma \text{ in } \llbracket \mathcal{E}[\square : \beta] : \alpha \rrbracket [\exists \alpha. \llbracket e : \alpha \rrbracket \wedge \square]$	Premise
3506	$\vartheta \text{ unifies } \mathcal{C}[\text{match } \alpha_1 := v\bar{\gamma}. \Pi_{i=1}^n \bar{\gamma} \text{ with } \dots]$	Premise
3507	$\vartheta \text{ unifies } \llbracket \mathcal{E}[e.j/n] : \alpha \rrbracket$	Above (See \implies direction)
3508	$\llbracket e : \tau \rrbracket = \text{let } \Gamma \text{ in } \llbracket \mathcal{E}[e.j] : \alpha \rrbracket$	By i.h.
3509	$\mathcal{C} = \text{let } \Gamma \text{ in } \llbracket \mathcal{E}[\square : \beta] : \alpha \rrbracket [\exists \alpha. \llbracket e : \alpha \rrbracket \wedge \square]$	Premise
3510	$\vartheta \text{ unifies } \mathcal{C}[\text{match } \alpha_1 := v\bar{\gamma}. \Pi_{i=1}^n \bar{\gamma} \text{ with } \dots]$	Premise
3511	$\vartheta \text{ unifies } \llbracket \mathcal{E}[e.j/n] : \alpha \rrbracket$	By <i>i.h.</i>
3512	$\vdash \mathcal{E}[e.j/n] : \vartheta(\alpha)$	Premise
3513	$\Gamma' \vdash \mathcal{E}[\{(e : g)\}] : \tau'$	Premise
3514	$\Gamma' = \emptyset$	$\mathcal{E}[\{(e : g)\}]$ is closed
3515		Lemma F.5
3516	$\vdash \mathcal{E}[\{(e : g)\}] : \tau'$	
3517	$[\alpha := \tau'] \text{ unifies } \llbracket \mathcal{E}[\{(e : g)\}] : \alpha \rrbracket$	By i.h.
3518	$\phi[\alpha := \phi(\tau')] \vdash \llbracket \mathcal{E}[\{(e : g)\}] : \alpha \rrbracket$	By definition
3519	$\mathcal{C}[\alpha_1 ! v\bar{\gamma}. \Pi_{i=1}^n \bar{\gamma}]$	Premise
3520	$\text{shape } (g) = v\bar{\gamma}. \Pi_{i=1}^n \bar{\gamma}$	$\implies E$
3521	$\mathcal{E}[e \triangleright v\bar{\gamma}. \Pi_{i=1}^n \bar{\gamma}]$	Above
3522	$\vdash \mathcal{E}[e.j] : \vartheta(\alpha)$	By CAN-PROJ-I
3523		

Subsubcase $[e], \langle e \rangle, \ell$.

Similar arguments.

□

3529 **F.5 Principal types**

3530 THEOREM F.16 (PRINCIPAL TYPES). *For any well-typed closed term e , there exists a type τ , which
3531 we call principal, such that: (i) $\vdash e : \tau$. (ii) For any other typing $\vdash e : \tau'$, then $\tau' = \theta(\tau)$ for some
3532 substitution θ .*

3533 PROOF. Let e be an arbitrary closed well-typed term; that is, there exists a type τ such that $\vdash e : \tau$.
3534 By [Theorem F.15](#), the constraint $\llbracket e : \alpha \rrbracket$ is satisfiable (specifically under the unifier $\alpha = \tau$). By
3535 [Corollary E.19](#), there exists a solved constraint \hat{C} such that $\hat{C} \equiv \llbracket e : \alpha \rrbracket$. From \hat{C} , we extract a unifier
3536 ϑ . Since $\hat{C} \equiv \exists \vartheta$, it follows that ϑ is *most general*.

3537 We claim that $\vartheta(\alpha)$ is the principal type of e . This amounts to showing:

- 3538 (i) $\vdash e : \vartheta(\alpha)$
3539 (ii) For any other typing $\vdash e : \tau'$, then $\tau' = \theta(\vartheta(\alpha))$ for some θ .

3540 Since ϑ is a unifier of $\llbracket e : \alpha \rrbracket$, it follows immediately from [Theorem F.15](#) that $\vdash e : \vartheta(\alpha)$, proving (i).
3541 For (ii), suppose $\vdash e : \tau'$ for some τ' . Then by [Theorem F.15](#) again, there exists a unifier ϑ' of $\llbracket e : \alpha \rrbracket$
3542 such that $\vartheta'(\alpha) = \tau'$. Since ϑ is most general, we have $\exists \vartheta' \models \exists \vartheta$, and by [Lemma F.11](#), this implies
3543 the existence of a substitution ϑ'' such that $\vartheta' = \vartheta'' \circ \vartheta$. Hence, $\tau' = \vartheta'(\alpha) = \vartheta''(\vartheta(\alpha))$, witnessing
3544 that τ' is an instance of $\vartheta(\alpha)$, as required (ii). \square

3545

3546

3547

3548

3549

3550

3551

3552

3553

3554

3555

3556

3557

3558

3559

3560

3561

3562

3563

3564

3565

3566

3567

3568

3569

3570

3571

3572

3573

3574

3575

3576

3577

CONTENTS

3578		
3579	Abstract	1
3580	1 Introduction	1
3581	2 Suspended constraints: an overview	5
3582	3 Semantics of constraints	10
3583	4 The OmniML calculus	15
3584	5 Solving constraints	17
3585	6 Implementation	22
3586	7 Related work	23
3587	8 Conclusions and future work	24
3588	References	26
3589	A The OmniML calculus: typing rules and constraint generation	27
3590	B Unification	34
3591	C Full technical reference	35
3592	D Properties of the constraint language	45
3593	E Properties of the constraint solver	54
3594	F Properties of OmniML	69
3595	Contents	74
3596		
3597		
3598		
3599		
3600		
3601		
3602		
3603		
3604		
3605		
3606		
3607		
3608		
3609		
3610		
3611		
3612		
3613		
3614		
3615		
3616		
3617		
3618		
3619		
3620		
3621		
3622		
3623		
3624		
3625		
3626		